



DeepLearning.AI

CNNs: Feature Maps and Receptive Fields

Specialized Approaches to Vision

Module 2: Specialized Approaches to Vision

- **CNNs:** Feature Maps and Receptive Fields
- **Interpretability:** Saliency & Class Activation Maps
- **Diffusion Models:** Image Generation

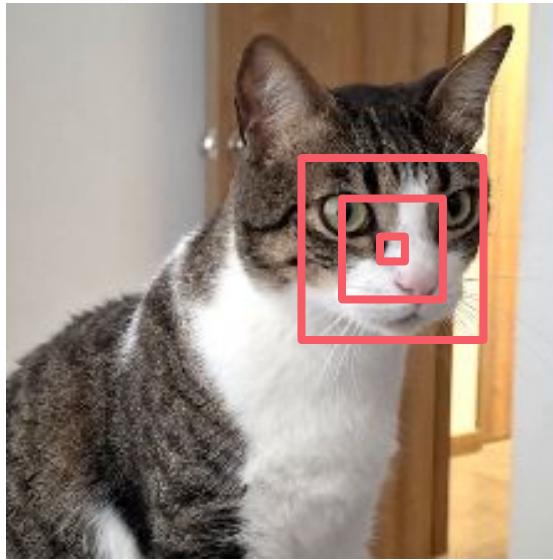
CNN Filters

-1	0	1
-2	0	2
-1	0	1



Detection Scale

-1	0	1
-2	0	2
-1	0	1

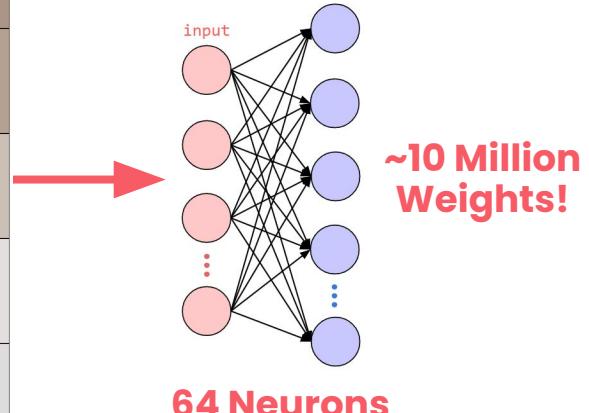


Convolutions vs. Linear Layers



A 5x5 grid of colored squares representing pixel values from a 224x224 image. Each square contains three numerical values: Red (R), Green (G), and Blue (B). The values range from approximately 170 to 252. The grid is labeled "150,000+ RGB Values" in red text at the bottom.

R: 188 G: 179 B: 172	R: 192 G: 178 B: 166	R: 200 G: 180 B: 164	R: 187 G: 164 B: 149	R: 170 G: 149 B: 134
R: 208 G: 201 B: 195	R: 215 G: 197 B: 184	R: 211 G: 191 B: 176	R: 206 G: 185 B: 170	R: 181 G: 161 B: 145
R: 232 G: 228 B: 225	R: 229 G: 216 B: 207	R: 224 G: 207 B: 196	R: 216 G: 200 B: 189	R: 204 G: 192 B: 182
R: 252 G: 249 B: 250	R: 240 G: 234 B: 231	R: 229 G: 224 B: 219	R: 228 G: 223 B: 220	R: 226 G: 223 B: 222
R: 249 G: 247 B: 250	R: 232 G: 228 B: 230	R: 232 G: 230 B: 231	R: 231 G: 229 B: 230	R: 223 G: 220 B: 221



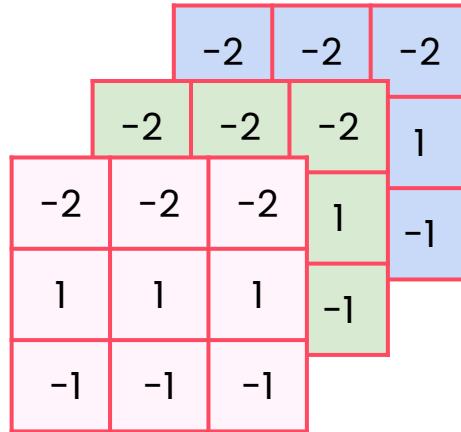
Convolutions vs. Linear Layers



224x224 image

R: 188 G: 179 B: 172	R: 192 G: 178 B: 166	R: 200 G: 180 B: 164	R: 187 G: 164 B: 149	R: 170 G: 149 B: 134
R: 208 G: 201 B: 195	R: 215 G: 197 B: 184	R: 211 G: 191 B: 176	R: 206 G: 185 B: 170	R: 181 G: 161 B: 145
R: 232 G: 228 B: 225	R: 229 G: 216 B: 207	R: 224 G: 207 B: 196	R: 216 G: 200 B: 189	R: 204 G: 192 B: 182
R: 252 G: 249 B: 250	R: 240 G: 234 B: 231	R: 229 G: 224 B: 219	R: 228 G: 223 B: 220	R: 226 G: 223 B: 222
R: 249 G: 247 B: 250	R: 232 G: 228 B: 230	R: 232 G: 230 B: 231	R: 231 G: 229 B: 230	R: 223 G: 220 B: 221

150,000+ RGB Values



$$3 \times 3 \times 3 = 27 \text{ weights} + 1 \text{ bias}$$

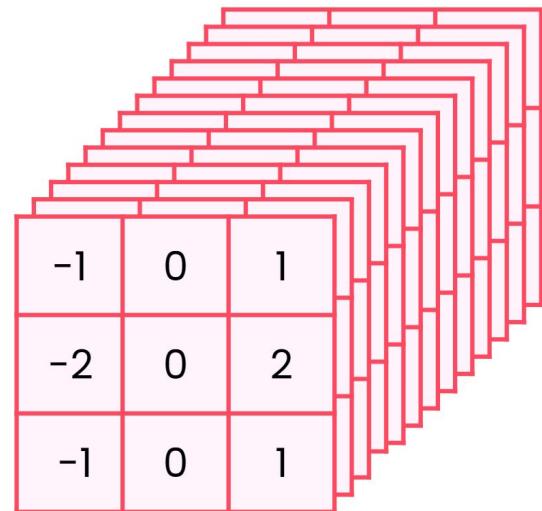
Convolutions vs. Linear Layers



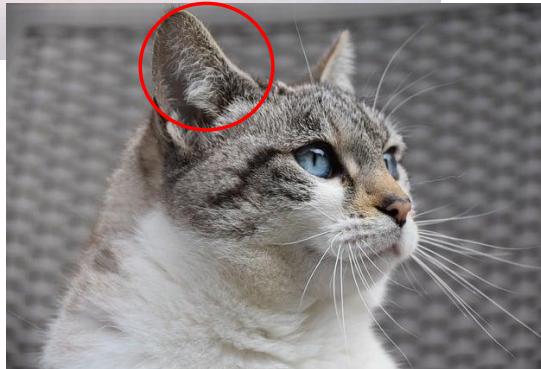
224x224 image

R: 188 G: 179 B: 172	R: 192 G: 178 B: 166	R: 200 G: 180 B: 164	R: 187 G: 164 B: 149	R: 170 G: 149 B: 134
R: 208 G: 201 B: 195	R: 215 G: 197 B: 184	R: 211 G: 191 B: 175	R: 206 G: 185 B: 170	R: 181 G: 161 B: 145
R: 232 G: 228 B: 225	R: 239 G: 226 B: 207		R: 235 G: 200 B: 190	R: 204 G: 192 B: 182
R: 252 G: 249 B: 250	R: 240 G: 234 B: 231	R: 229 G: 224 B: 219	R: 228 G: 223 B: 220	R: 226 G: 223 B: 222
R: 249 G: 247 B: 250	R: 232 G: 228 B: 230	R: 232 G: 230 B: 231	R: 231 G: 229 B: 230	R: 223 G: 220 B: 221

150,000+ RGB Values



Convolutions vs. Linear Layers



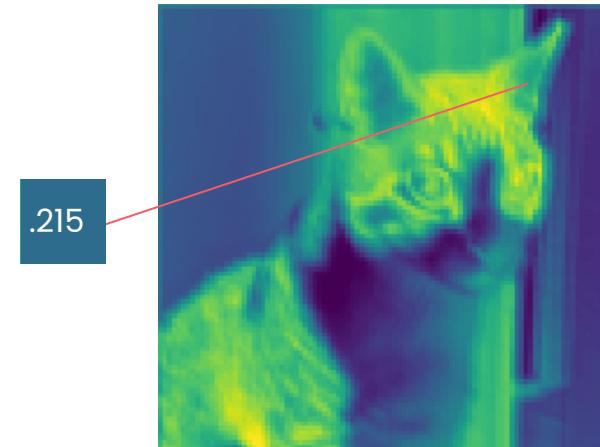
Receptive Fields



96x96 image

R: 229 G: 220 B: 224	R: 234 G: 228 B: 233	R: 229 G: 224 B: 219
R: 231 G: 223 B: 227	R: 234 G: 227 B: 233	R: 240 G: 234 B: 239
R: 224 G: 216 B: 219	R: 240 G: 232 B: 237	R: 245 G: 239 B: 243

3x3 filter = 9 Pixels



Feature Map
96x96 activations

Receptive Fields



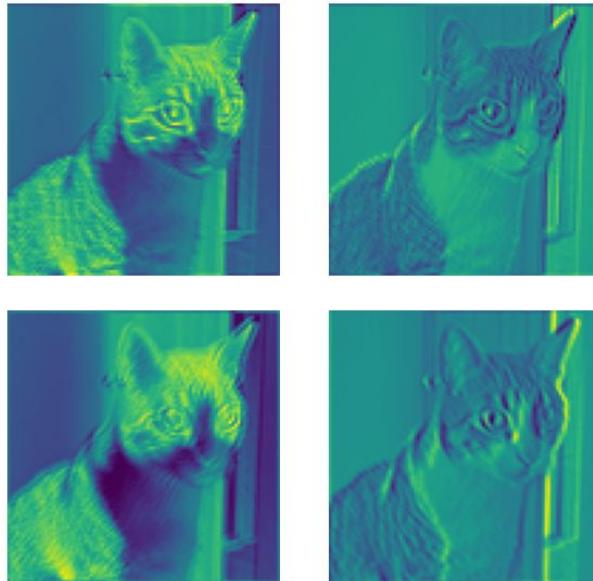
Receptive Field

96x96 image

Receptive Fields

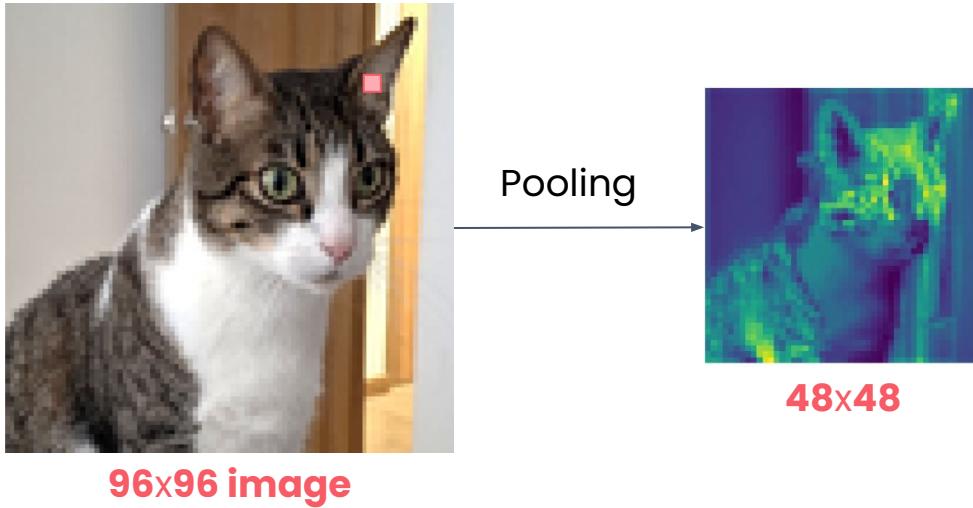


96x96 image

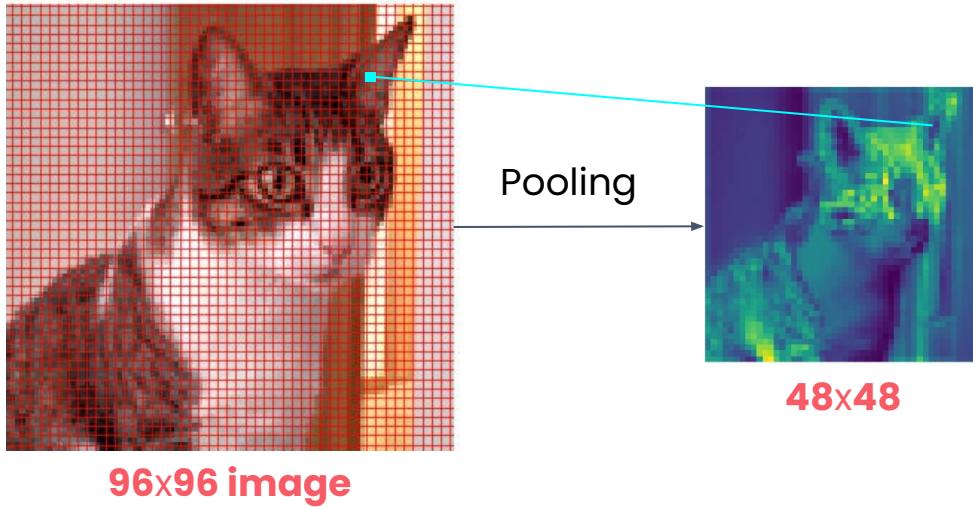


Feature Maps

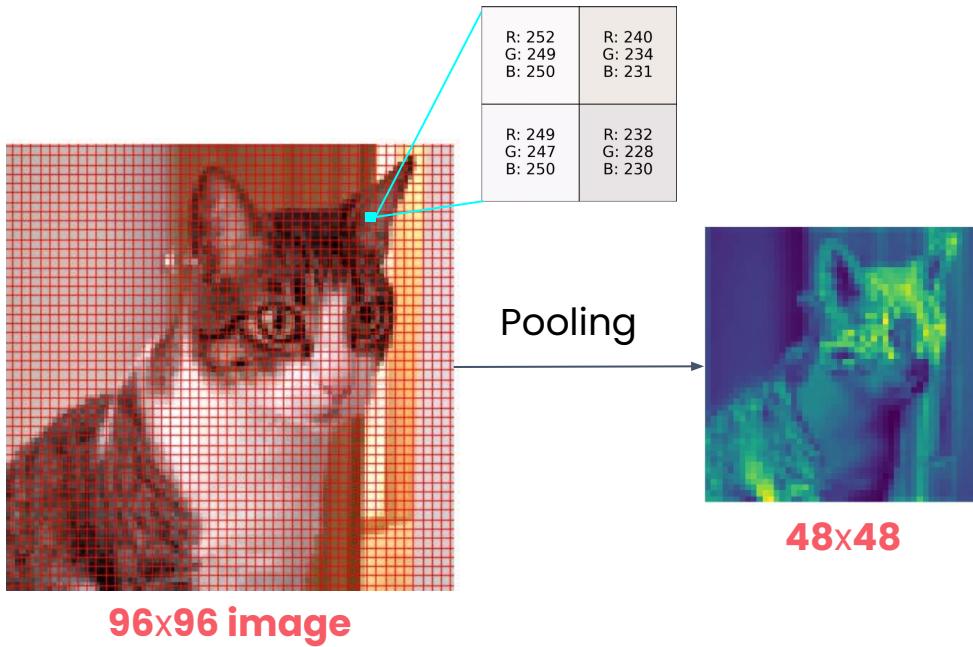
Receptive Fields



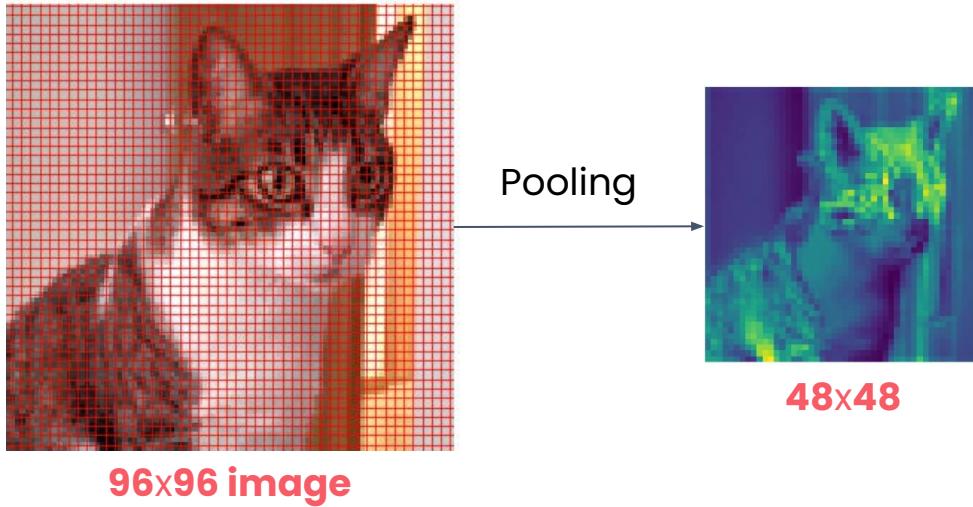
Receptive Fields



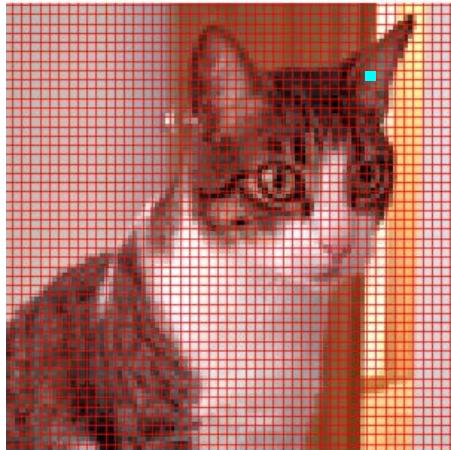
Receptive Fields



Receptive Fields



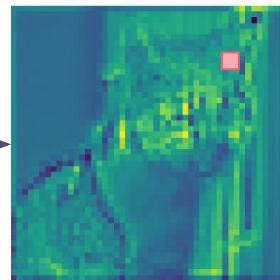
Receptive Fields



96x96 image

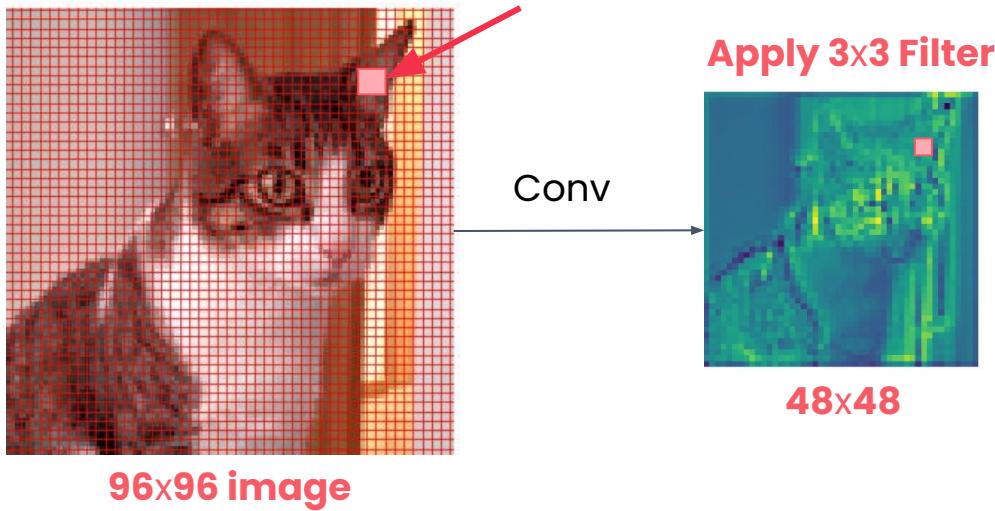
Conv

Apply 3x3 Filter

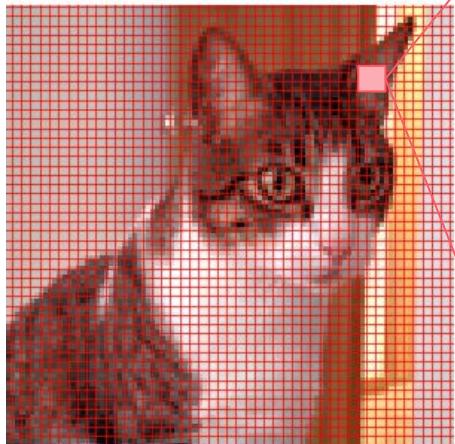


48x48

Receptive Fields



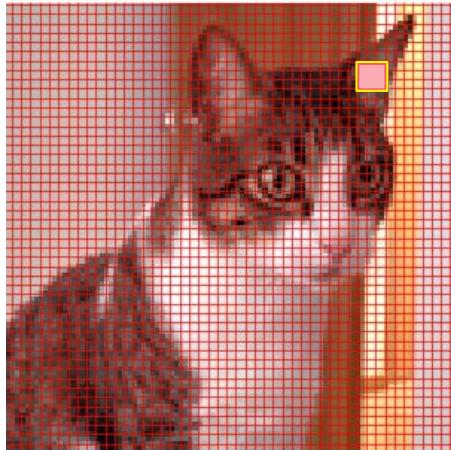
Receptive Fields



R: 229 G: 220 B: 224	R: 234 G: 228 B: 233	R: 245 G: 243 B: 247	R: 240 G: 239 B: 243	R: 235 G: 235 B: 239	R: 156 G: 153 B: 158
R: 231 G: 223 B: 227	R: 234 G: 227 B: 233	R: 240 G: 234 B: 239	R: 243 G: 241 B: 247	R: 234 G: 234 B: 241	R: 208 G: 207 B: 212
R: 224 G: 216 B: 219	R: 240 G: 232 B: 237	R: 245 G: 239 B: 243	R: 231 G: 229 B: 234	R: 204 G: 203 B: 210	R: 193 G: 193 B: 198
R: 219 G: 206 B: 210	R: 237 G: 228 B: 231	R: 244 G: 237 B: 240	R: 231 G: 226 B: 231	R: 202 G: 199 B: 205	R: 206 G: 206 B: 209
R: 212 G: 194 B: 198	R: 231 G: 215 B: 220	R: 237 G: 222 B: 224	R: 223 G: 211 B: 215	R: 222 G: 219 B: 223	R: 232 G: 232 B: 235
R: 201 G: 178 B: 178	R: 226 G: 202 B: 206	R: 230 G: 207 B: 209	R: 222 G: 204 B: 205	R: 232 G: 227 B: 232	R: 242 G: 242 B: 245

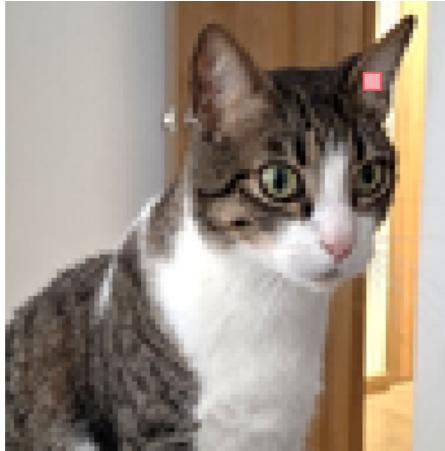
3x3 filter = 36 Pixels

Receptive Fields



96x96 image

Receptive Fields

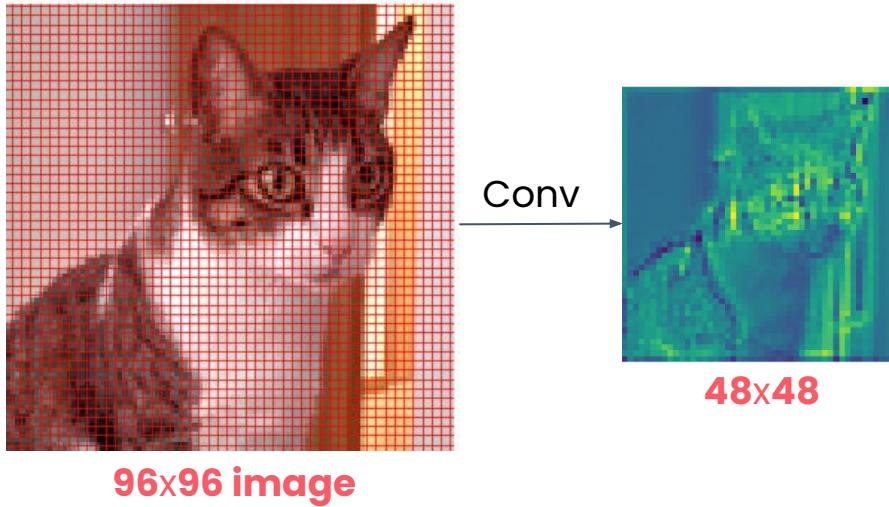


Receptive Field
Before 1st MaxPool: 9 pixels

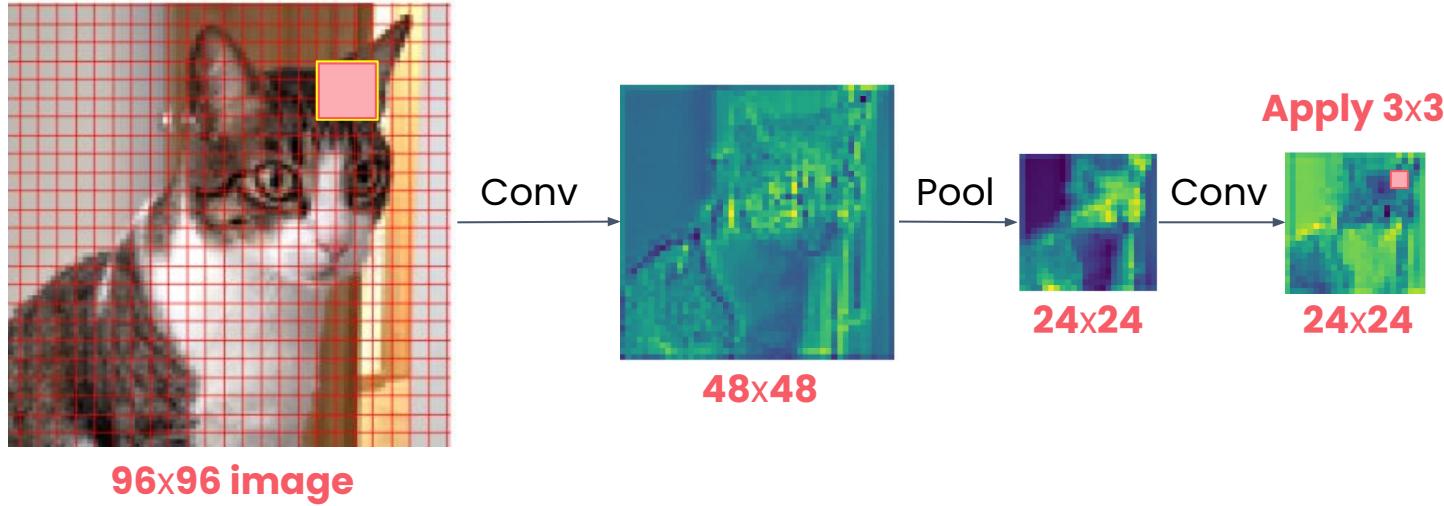


Receptive Field
After 1st MaxPool: 36 pixels

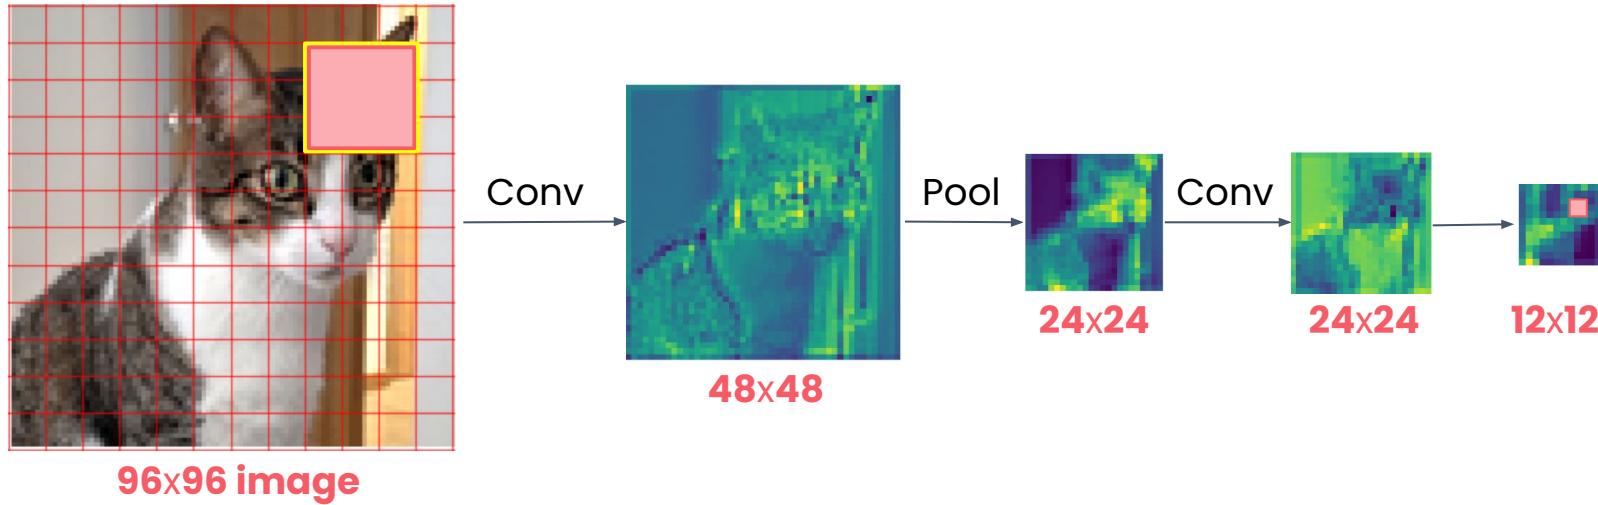
Receptive Fields



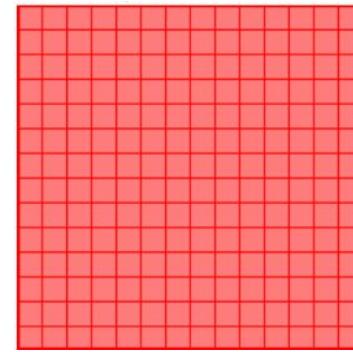
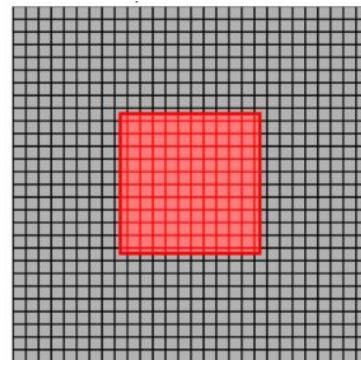
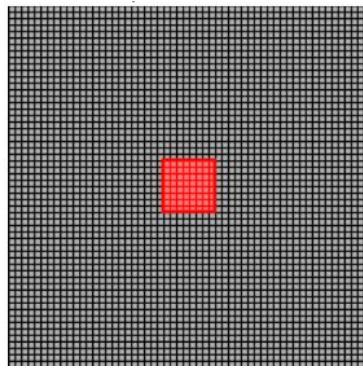
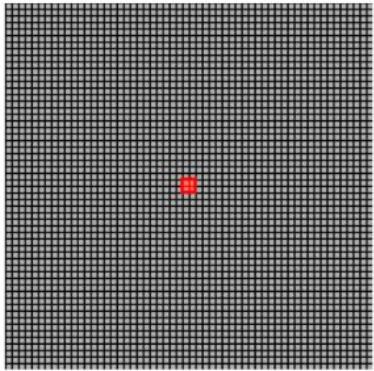
Receptive Fields



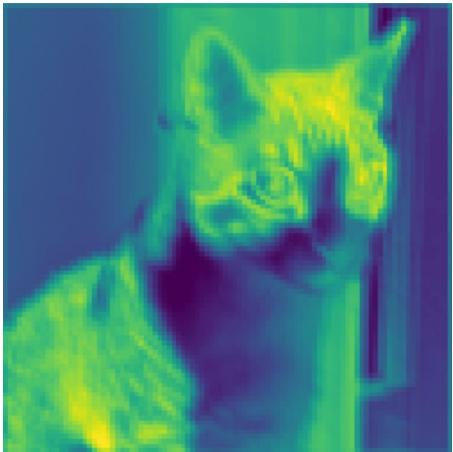
Receptive Fields



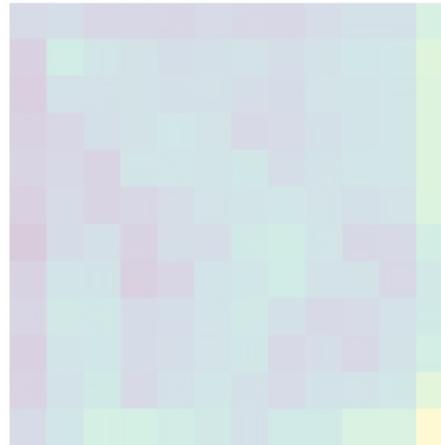
Receptive Fields



Receptive Fields



96x96

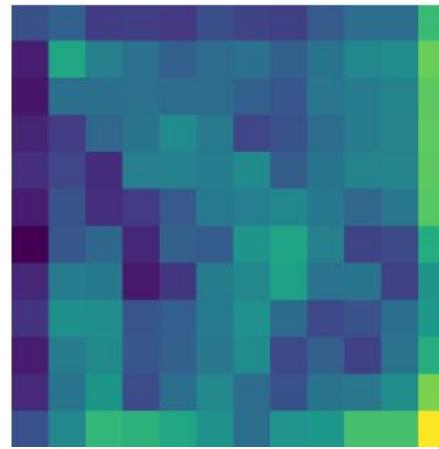


12x12

Receptive Fields

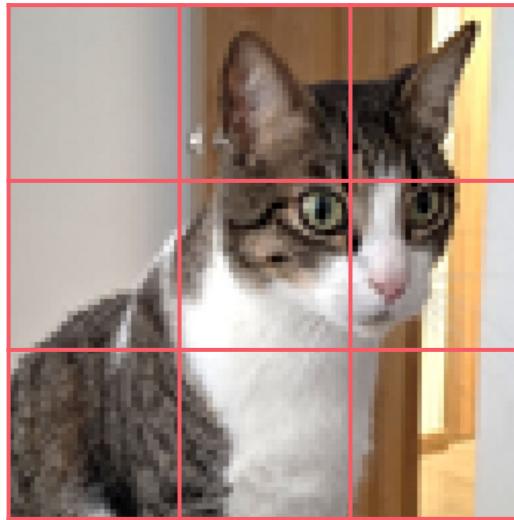


96x96

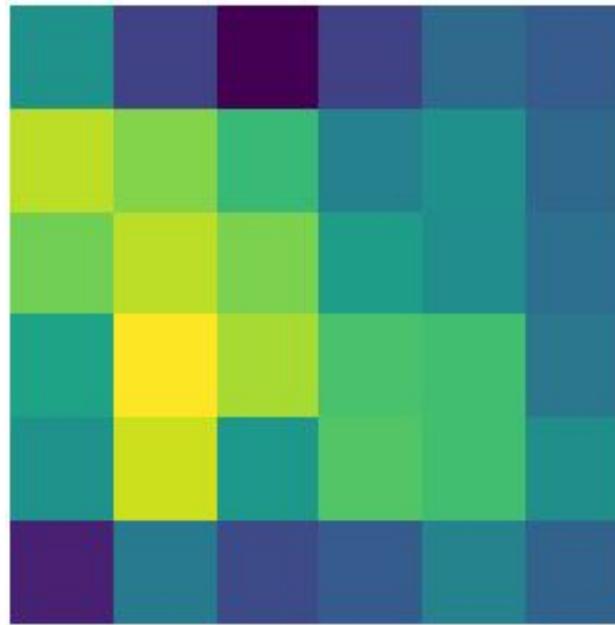


12x12

Receptive Fields



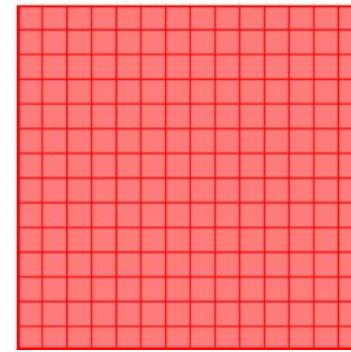
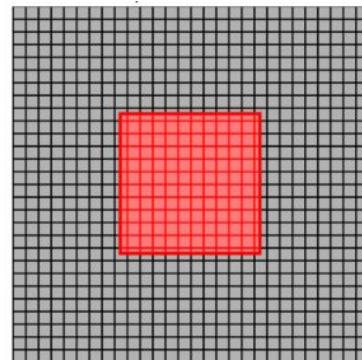
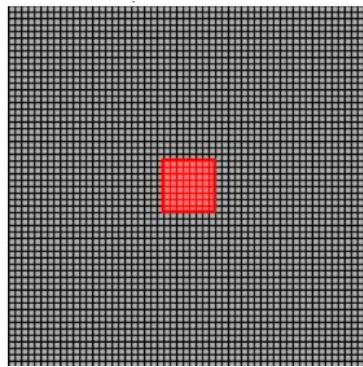
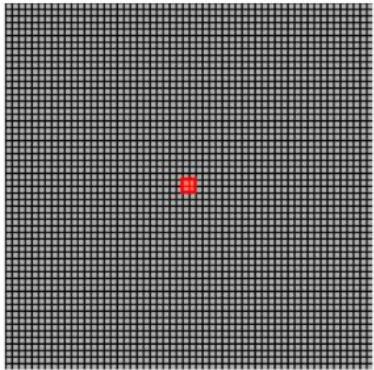
Receptive Fields



This is a cat!

6x6

Receptive Fields



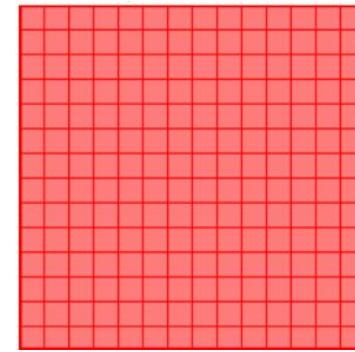
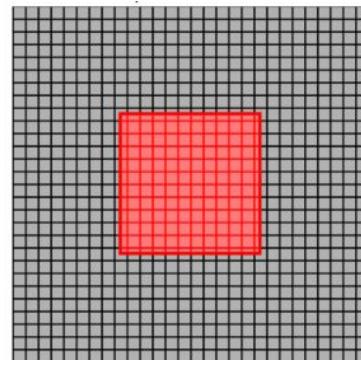
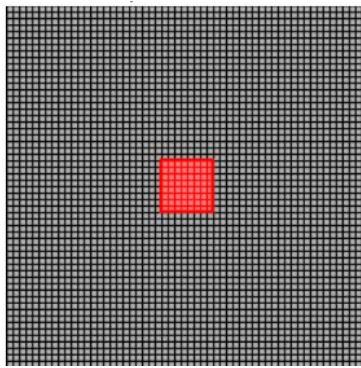
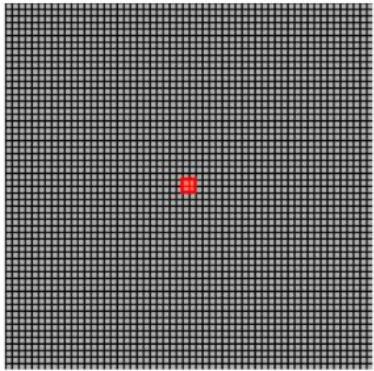


DeepLearning.AI

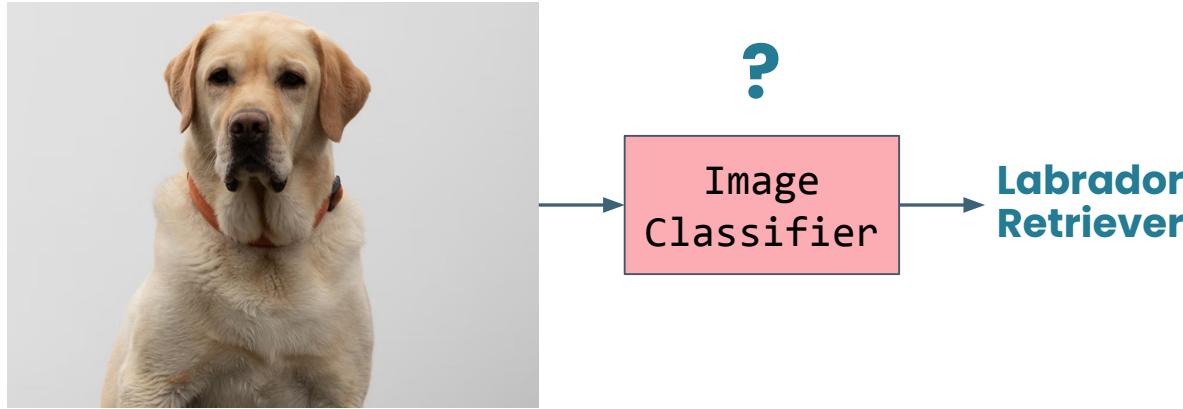
Saliency Maps

Specialized Approaches to Vision

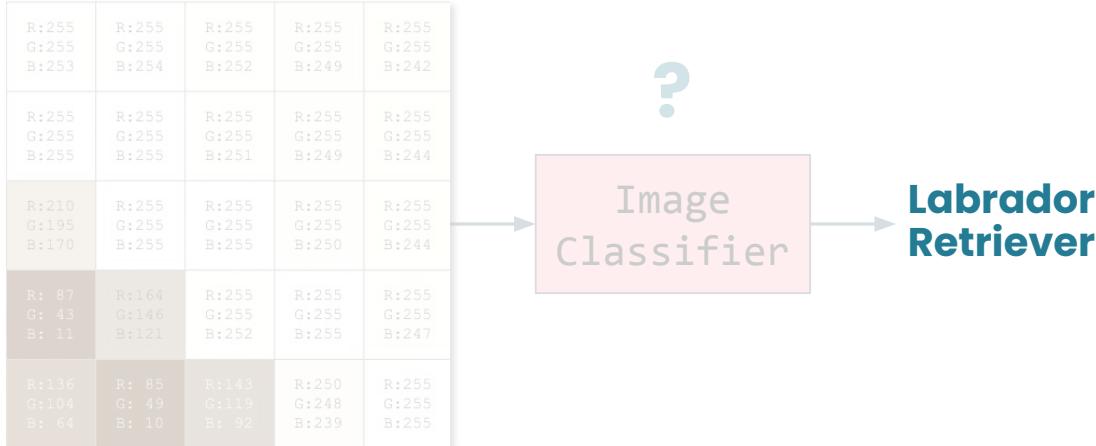
Receptive Fields



Classification



Classification



Interpretability Techniques

Reveal which parts of an input most affect a models' prediction



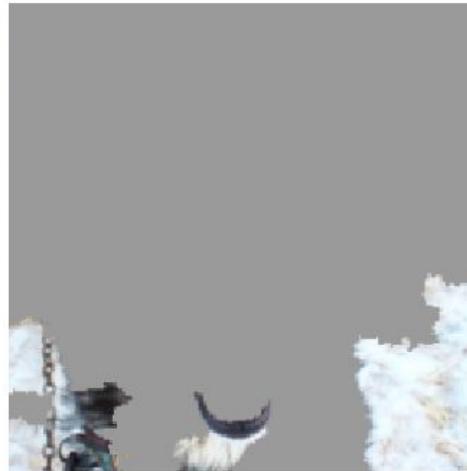
Why Interpretability Matters

- Trust



Why Interpretability Matters

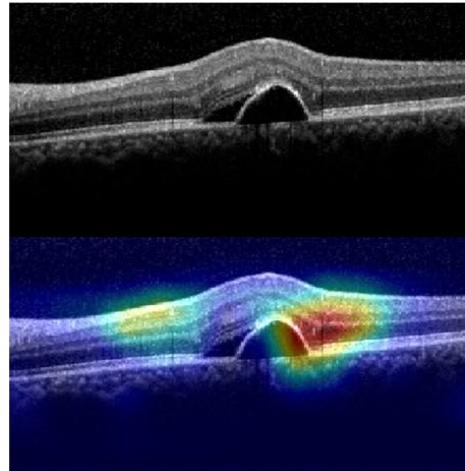
- Trust
- Debugging



Ribeiro, Singh & Guestrin (2016)

Why Interpretability Matters

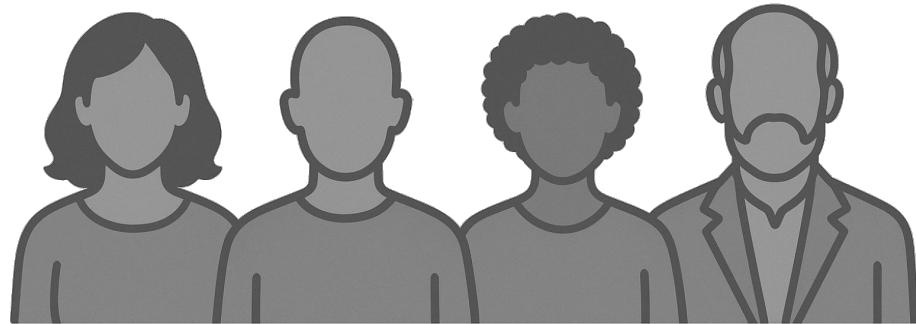
- Trust
- Debugging
- Discovery



Yoon et al. (2020)

Why Interpretability Matters

- Trust
- Debugging
- Discovery
- Ethical Concerns



Saliency Maps

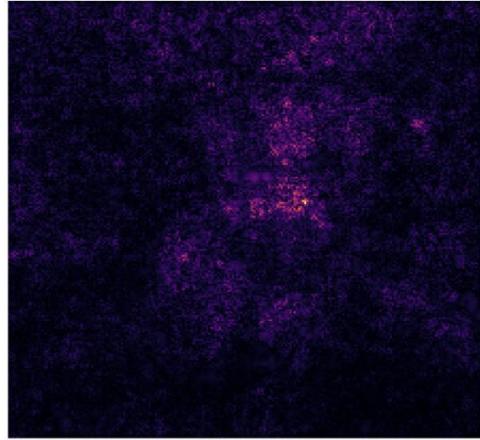


Image
Classifier

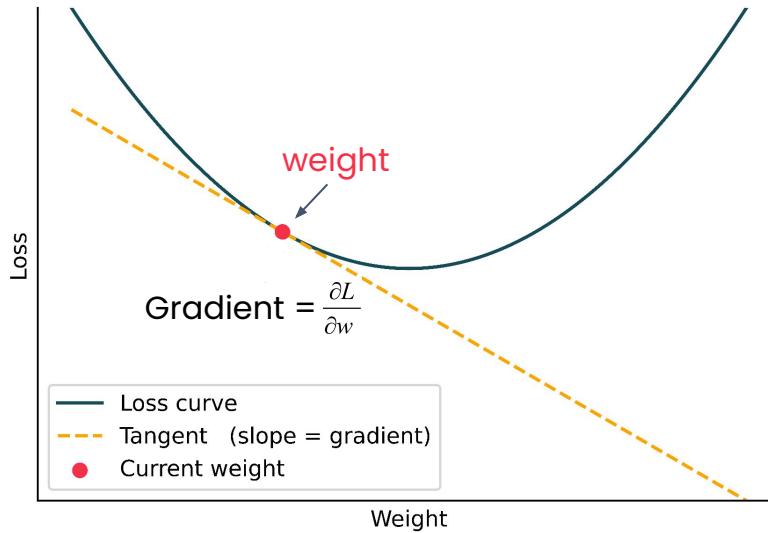
Dog

R:255 G:255 B:253	R:255 G:255 B:254	R:255 G:255 B:252	R:255 G:255 B:249	R:255 G:255 B:242
R:255 G:255 B:255	R:255 G:255 B:255	R:255 G:255 B:251	R:255 G:255 B:249	R:255 G:255 B:244
R:210 G:195 B:170	R:255 G:255 B:255	R:250 G:248 B:239	R:255 G:255 B:250	R:255 G:255 B:244
R: 87 G: 43 B: 11	R:164 G:146 B:121	R:255 G:255 B:252	R:255 G:255 B:255	R:255 G:255 B:247
R:136 G:104 B: 64	R: 85 G: 49 B: 10	R:143 G:119 B: 92	R:250 G:248 B:239	R:255 G:255 B:255

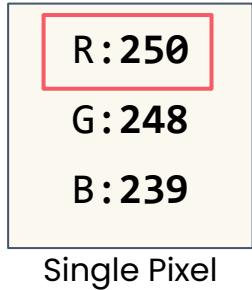
Saliency Maps



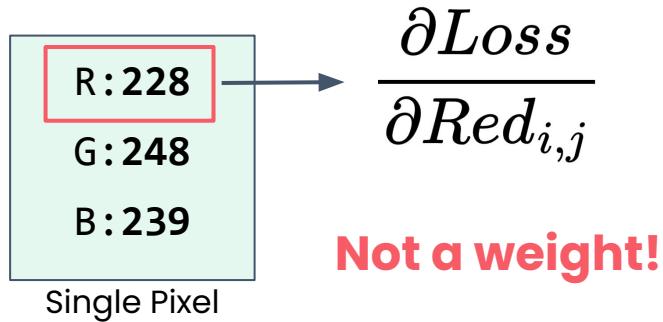
Gradients



Gradients on Pixels



Gradients on Pixels



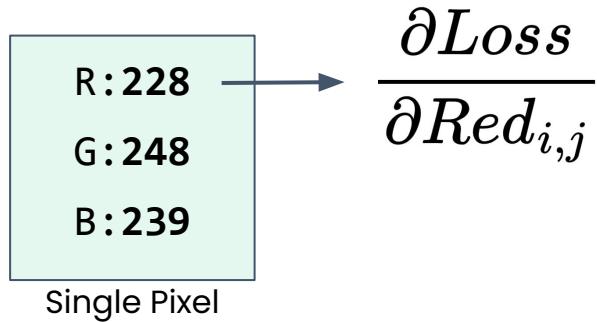
Gradients on Pixels

```
input_image.requires_grad_()
```

Gradients on Pixels

```
input_image.requires_grad_()  
  
output = model(input_image)  
output[0, target_class].backward()
```

Gradients on Pixels



Saliency Maps



Saliency Maps

```
input_image = input_image.clone().detach()
input_image.requires_grad_()

output = model(input_image)
model.zero_grad()
output[0, target_class].backward()

gradients = input_image.grad.data[0]

saliency_map = torch.abs(gradients).sum(dim=0).cpu().numpy()

saliency_map = (saliency_map - saliency_map.min()) /
    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
input_image = input_image.clone().detach()  
input_image.requires_grad_()  
  
output = model(input_image)  
model.zero_grad()  
output[0, target_class].backward()  
  
gradients = input_image.grad.data[0]  
  
saliency_map = torch.abs(gradients).sum(dim=0).vcpu().numpy()  
  
saliency_map = (saliency_map - saliency_map.min()) /  
    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
input_image = input_image.clone().detach()  
input_image.requires_grad_()  
  
output = model(input_image)  
model.zero_grad()  
output[0, target_class].backward()  
  
gradients = input_image.grad.data[0]  
  
saliency_map = torch.abs(gradients).sum(dim=0).vcpu().numpy()  
  
saliency_map = (saliency_map - saliency_map.min()) /  
    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
input_image = input_image.clone().detach()  
input_image.requires_grad_()  
  
output = model(input_image)  
model.zero_grad()  
output[0, target_class].backward()  
  
gradients = input_image.grad.data[0]  
  
saliency_map = torch.abs(gradients).sum(dim=0).cpu().numpy()  
  
saliency_map = (saliency_map - saliency_map.min()) /  
    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
input_image = input_image.clone().detach()  
input_image.requires_grad_()  
  
output = model(input_image)  
model.zero_grad()  
output[0, target_class].backward()  
  
gradients = input_image.grad.data[0]  
  
saliency_map = torch.abs(gradients).sum(dim=0).vcpu().numpy()  
  
saliency_map = (saliency_map - saliency_map.min()) /  
    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
# Load pre-trained ResNet50 model
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

Saliency Maps

```
# Load pre-trained ResNet50 model
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

Saliency Maps

```
# Load pre-trained ResNet50 model
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

Saliency Maps

```
# Load pre-trained ResNet50 model
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

Saliency Maps

```
img_path = 'images/dog.jpg'

transform = transforms.Compose([
    transforms.Resize(224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))
])

if not os.path.exists(img_path):
    raise FileNotFoundError("Please provide a valid path to an image.")

img_pil = load_image(img_path)
img_tensor = transform(img_pil).unsqueeze(0).to(device)
```

Saliency Maps

```
img_path = 'images/dog.jpg'

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])

if not os.path.exists(img_path):
    raise FileNotFoundError("Please provide a valid path to an image.")

img_pil = load_image(img_path)
img_tensor = transform(img_pil).unsqueeze(0).to(device)
```

Saliency Maps

```
img_path = 'images/dog.jpg'

transform = transforms.Compose([
    transforms.Resize(224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))
])

if not os.path.exists(img_path):
    raise FileNotFoundError("Please provide a valid path to an image.")

img_pil = load_image(img_path)
img_tensor = transform(img_pil).unsqueeze(0).to(device)
```

Saliency Maps

```
img_path = 'images/dog.jpg'

transform = transforms.Compose([
    transforms.Resize(224, 224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])

if not os.path.exists(img_path):
    raise FileNotFoundError("Please provide a valid path to an image.")

img_pil = load_image(img_path)
img_tensor = transform(img_pil).unsqueeze(0).to(device)
```

Saliency Maps

```
img_path = 'images/dog.jpg'

transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]))
])

if not os.path.exists(img_path):
    raise FileNotFoundError("Please provide a valid path to an image.")

img_pil = load_image(img_path)
img_tensor = transform(img_pil).unsqueeze(0).to(device)
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    ...
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    ...
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    ...
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    ...
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    input_image = input_image.clone().detach() leaf tensor  
    input_image.requires_grad_()
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()  
  
    output = model(input_image)
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()  
  
    output = model(input_image)  
    model.zero_grad()  
    output[0, target_class].backward()
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()  
  
    output = model(input_image)  
    model.zero_grad()  
    output[0, target_class].backward()
```

Saliency Maps

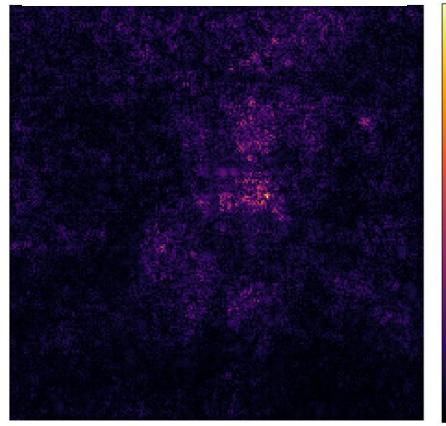
```
def compute_saliency_map(model, input_image, target_class=None):  
  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()  
  
    output = model(input_image)  
    model.zero_grad()  
    output[0, target_class].backward()  
  
    gradients = input_image.grad.data[0]
```

Saliency Maps

```
def compute_saliency_map(model, input_image, target_class=None):  
  
    input_image = input_image.clone().detach()  
    input_image.requires_grad_()  
  
    output = model(input_image)  
    model.zero_grad()  
    output[0, target_class].backward()  
  
    gradients = input_image.grad.data[0]  
  
    saliency_map = torch.abs(gradients).sum(dim=0).cpu().numpy()  
  
    saliency_map = (saliency_map - saliency_map.min()) /  
                    (saliency_map.max() - saliency_map.min() + 1e-8)
```

Saliency Maps

```
def visualize_saliency(img_display, saliency_map pred_class, pred_score, title):  
    . . .
```

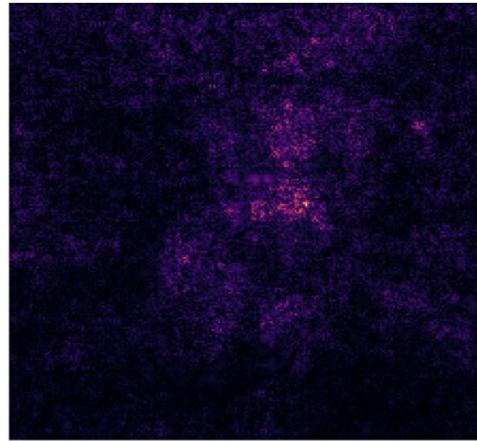


Gradients on Pixels

R:228
G:248
B:239

Single Pixel

$$\frac{\partial Loss}{\partial Red_{i,j}}$$



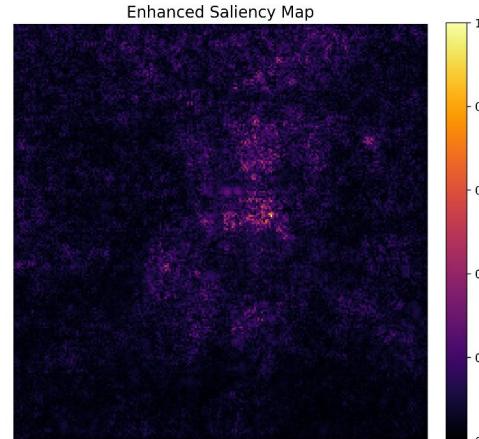


DeepLearning.AI

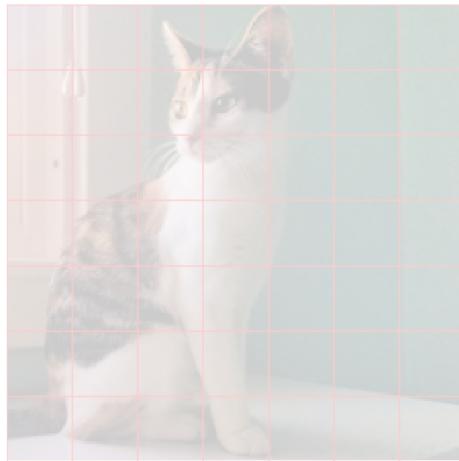
Class Activation Maps

Specialized Approaches to Vision

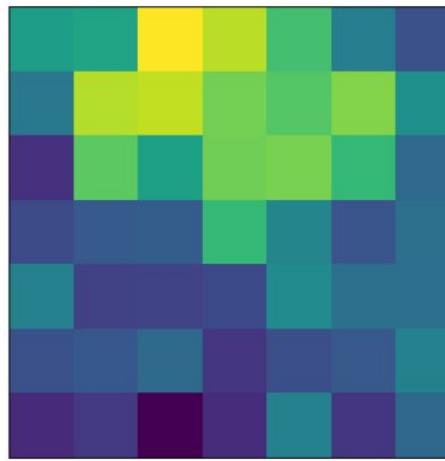
Saliency Maps



Feature Maps

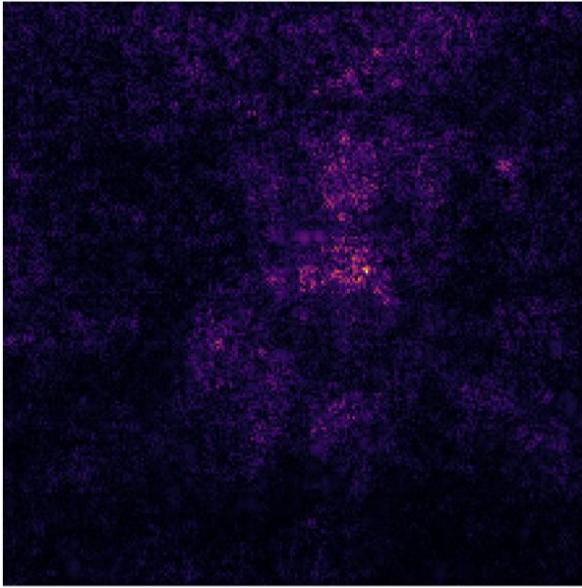


224x224 image

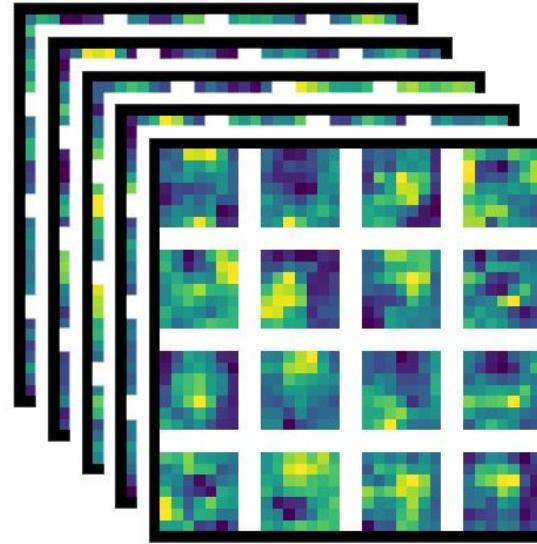


7x7 Feature Map

CAMs

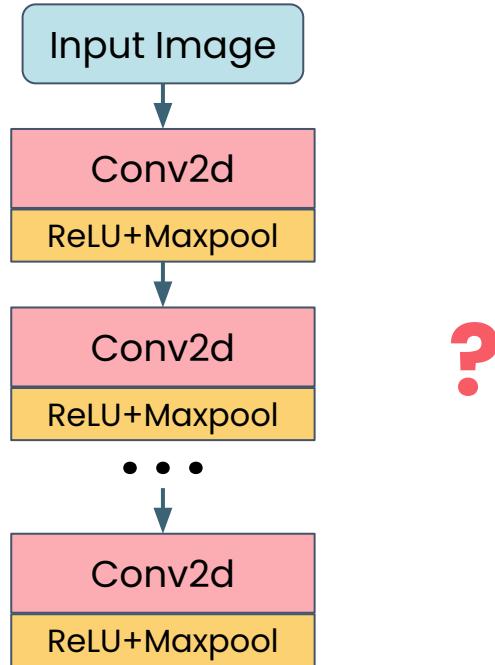


Saliency Maps:
Which pixels matter most?

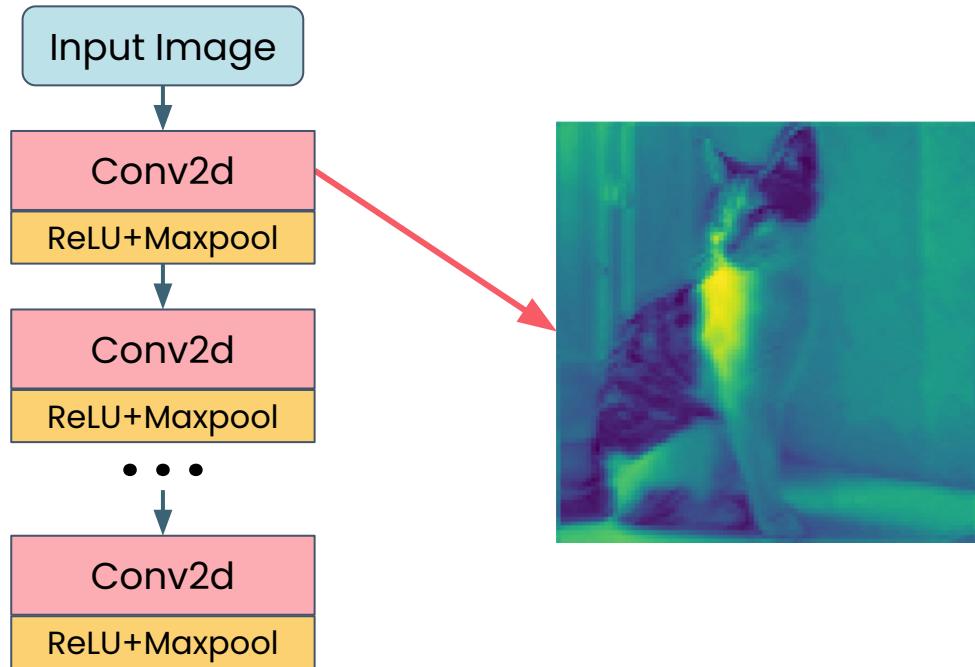


Class Activation Maps (CAMs):
Which feature maps matter most?

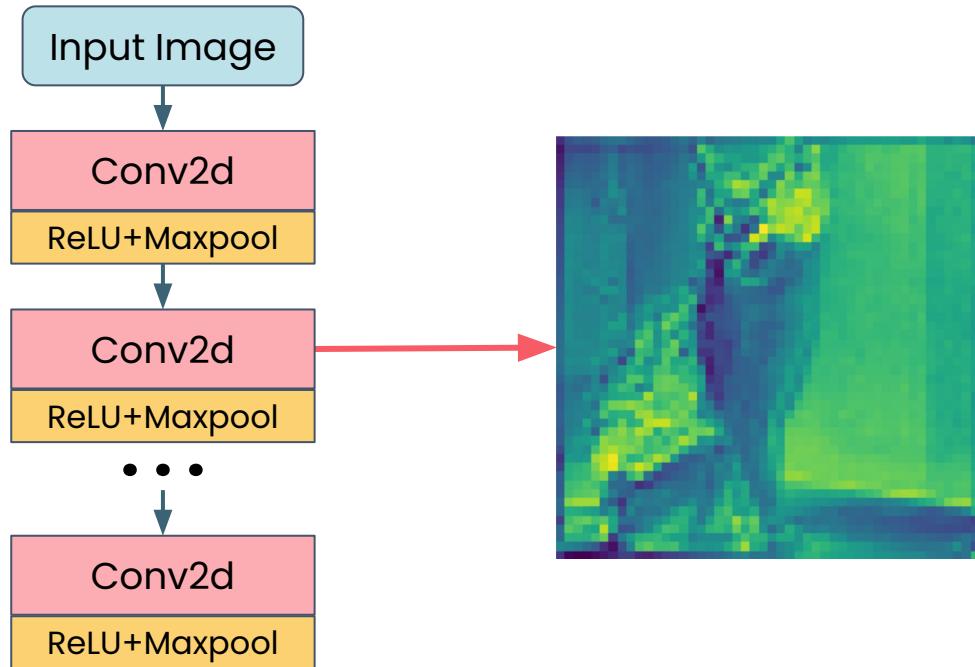
Feature Maps



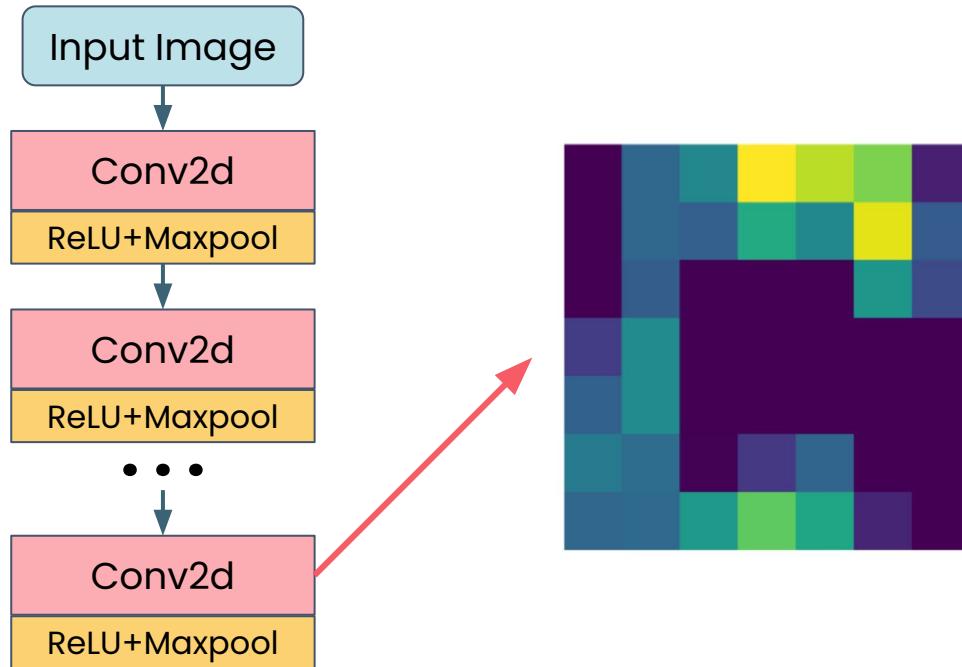
Feature Maps



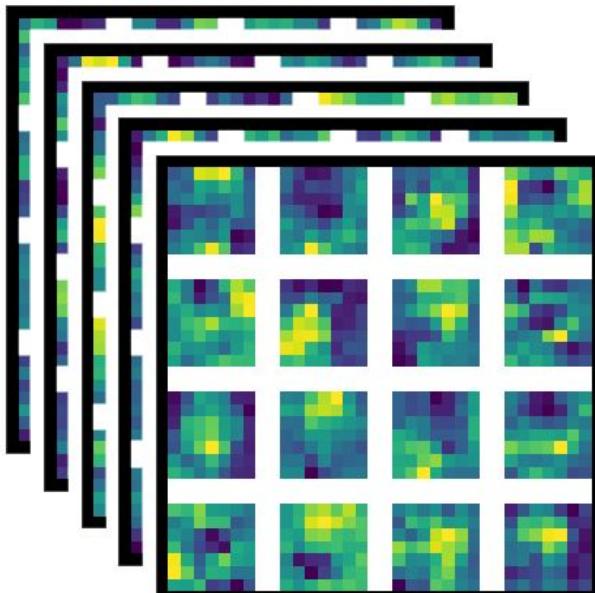
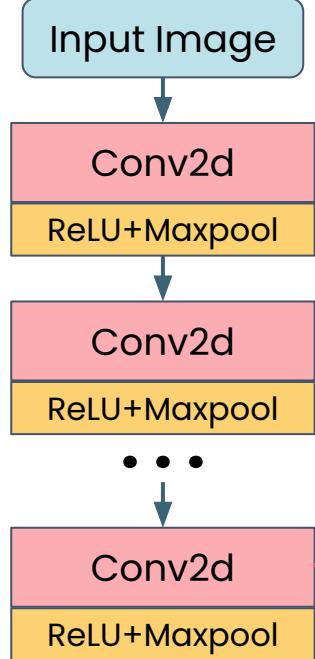
Feature Maps



Feature Maps



Feature Maps



2048 Feature Maps
7x7



224x224 Image

Class Activation Maps



GradCAM

```
# Load pretrained ResNet50 model and class labels
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

GradCAM

```
# Load pretrained ResNet50 model and class labels
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)
```

GradCAM

```
# Load pretrained ResNet50 model and class labels
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)

class GradCAM:
    def __init__(self, model, target_layer):
        . . .
```

GradCAM

```
# Load pretrained ResNet50 model and class labels
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)

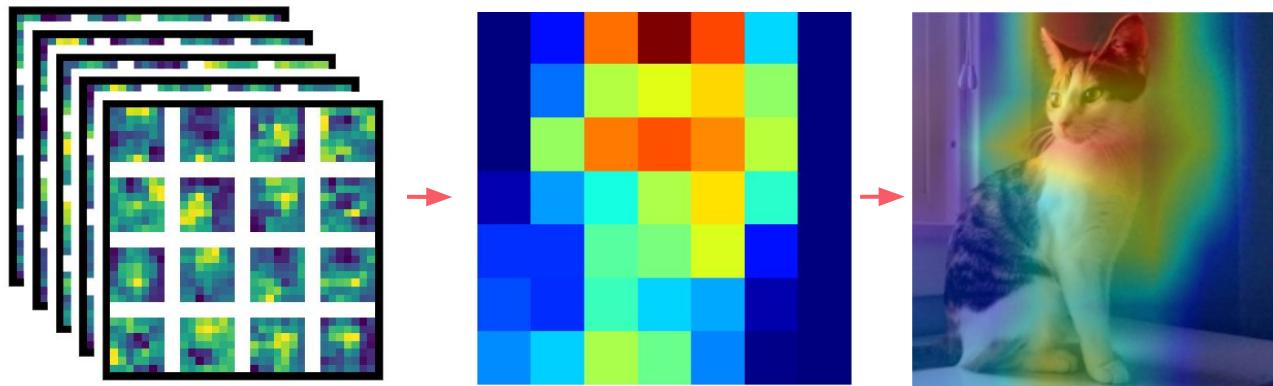
class GradCAM:
    def __init__(self, model, target_layer):
        . . .
```

GradCAM

```
# Load pretrained ResNet50 model and class labels
model = models.resnet50(weights=models.ResNet50_Weights.IMGNET1K_V1)
model.eval()
model.to(device)

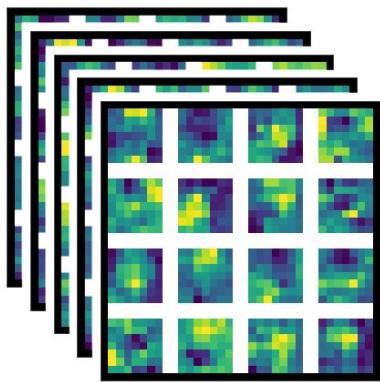
class GradCAM:
    def __init__(self, model, target_layer):
        . . .
```

GradCAM



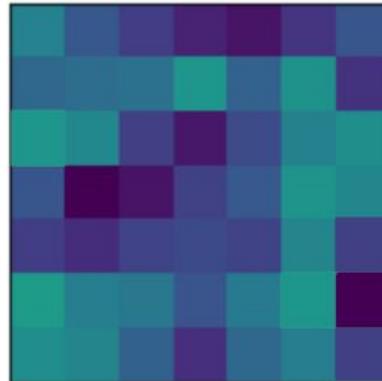
2048 Feature Maps
7x7

GradCAM

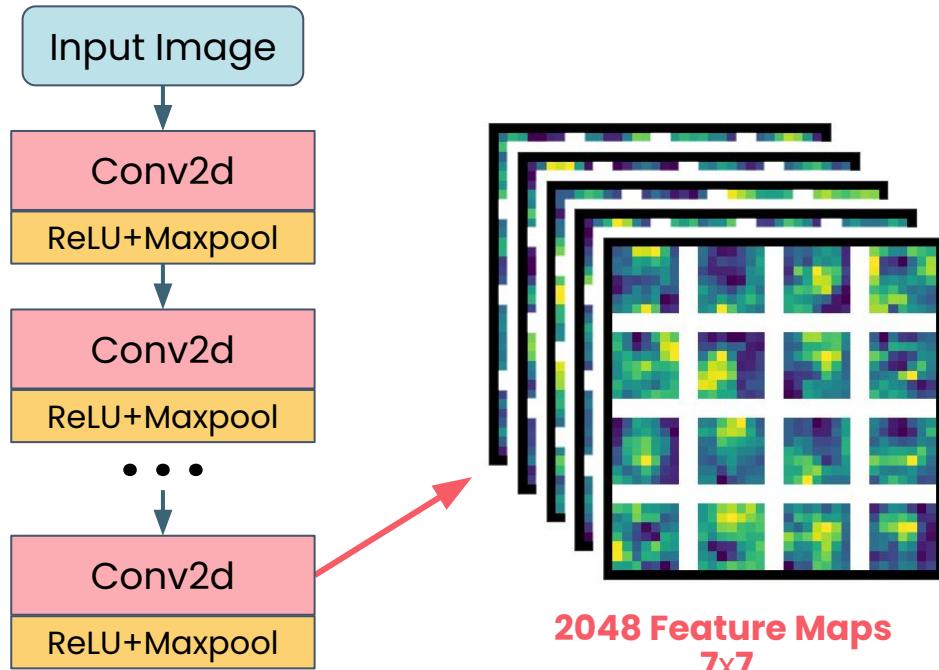


2048 Feature Maps
7x7

Feature Map #274



GradCAM



Intermediate Tensors = No saved gradients

Hooks

```
# "Hook" (register callback) into target layer for forward activations  
target_layer.register_forward_hook(self.save_activation)  
  
# Hook into the backward pass to save the gradients  
target_layer.register_full_backward_hook(self.save_gradient)
```

Hooks

```
# "Hook" (register callback) into target layer for forward activations  
target_layer.register_forward_hook(self.save_activation)  
  
# Hook into the backward pass to save the gradients  
target_layer.register_full_backward_hook(self.save_gradient)
```

Hooks

```
# "Hook" (register callback) into target layer for forward activations  
target_layer.register_forward_hook(self.save_activation)  
  
# Hook into the backward pass to save the gradients  
target_layer.register_full_backward_hook(self.save_gradient)
```

Hooks

```
# "Hook" (register callback) into target layer for forward activations  
target_layer.register_forward_hook(self.save_activation)  
  
# Hook into the backward pass to save the gradients  
target_layer.register_full_backward_hook(self.save_gradient)
```

Hooks: Forward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):

        # Forward Pass
        output = self.model(x)

# Automatically called after the forward pass
def save_activation(self, module, input, output):
    # .detach() to avoid linking this to computational graph
    self.activations = output.detach()
```

Hooks: Forward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):

        # Forward Pass
        output = self.model(x)
```

```
# Automatically called after the forward pass
def save_activation(self, module, input, output):
    # .detach() to avoid linking this to computational graph
    self.activations = output.detach()
```

Hooks: Forward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):

        # Forward Pass
        output = self.model(x)
```

```
# Automatically called after the forward pass
def save_activation(self, module, input, output):
    # .detach() to avoid linking this to computational graph
    self.activations = output.detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        ...
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        ...
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        ...
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        ...
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        . . .
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

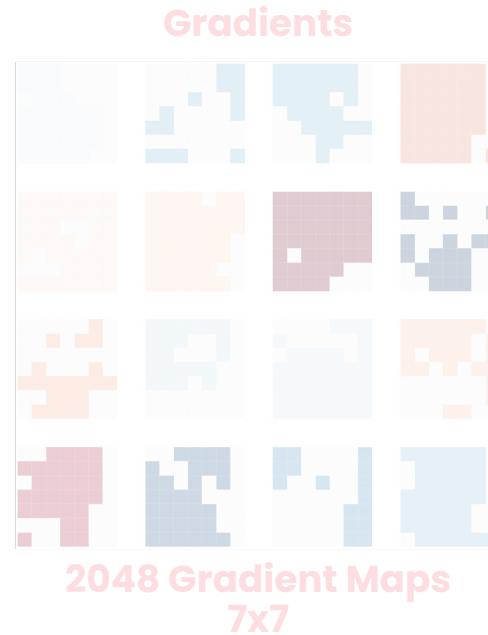
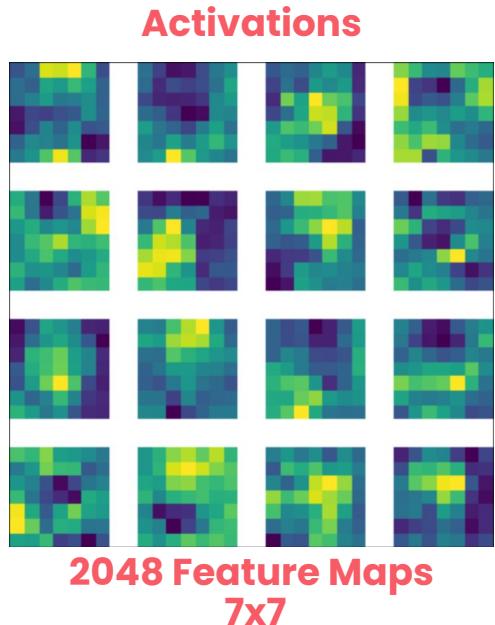
```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Hooks: Backward Pass

```
class GradCAM(nn.Module):
    def __call__(self, x, class_idx=None):
        ...
        # Backward Pass
        self.model.zero_grad()
        one_hot = torch.zeros_like(output)
        one_hot[0, class_idx] = 1
        output.backward(gradient=one_hot, retain_graph=True)
```

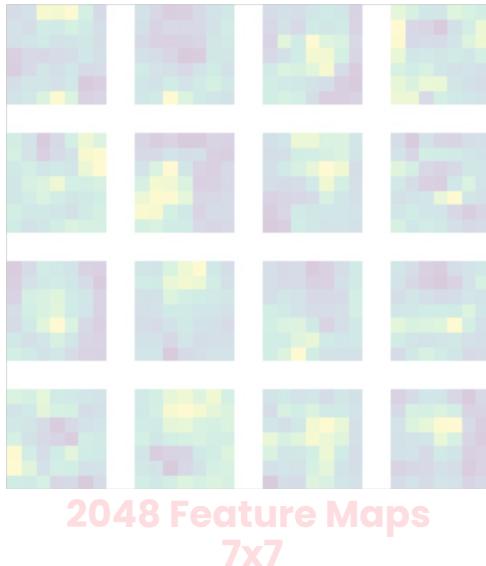
```
# Automatically called after the backward pass
def save_gradients(self, module, grad_input, grad_output):
    # grad_output is a tuple; the gradient we care about is [0]
    self.gradients = grad_output[0].detach()
```

Activations and Gradients



Activations and Gradients

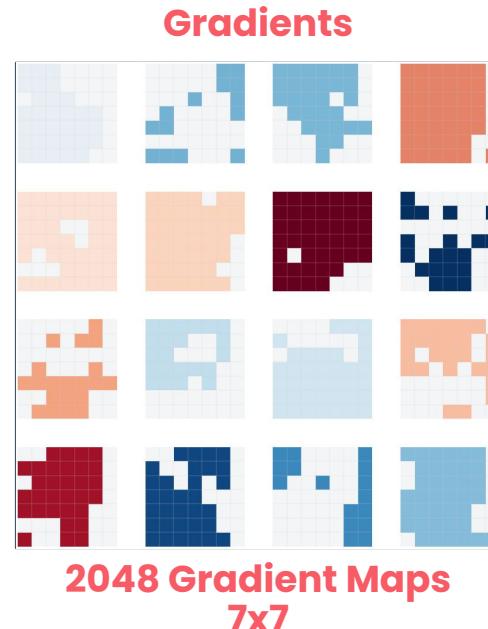
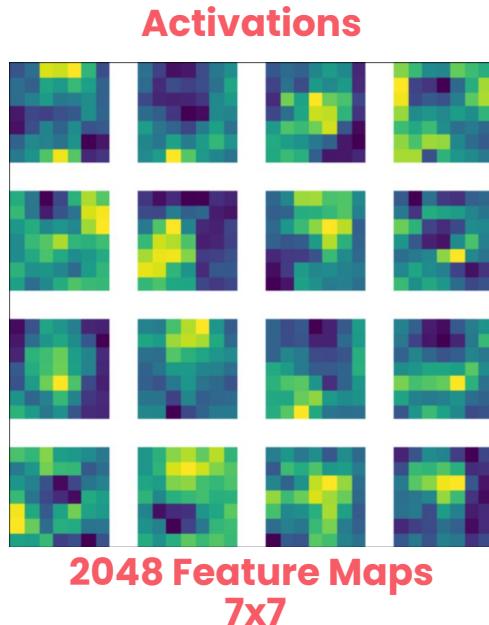
Activations



Gradients

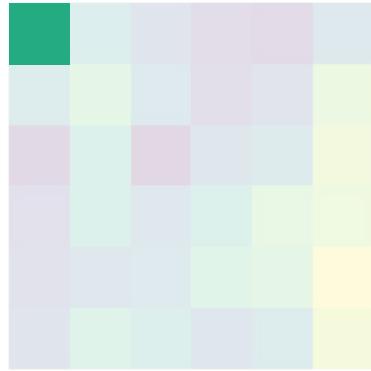


Activations and Gradients



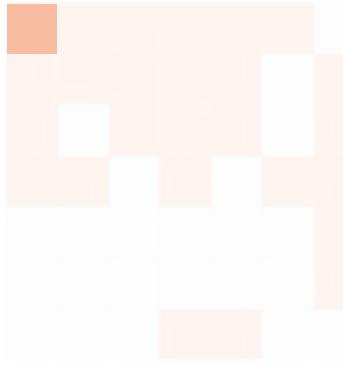
Activations and Gradients

strength



7x7 Feature Map

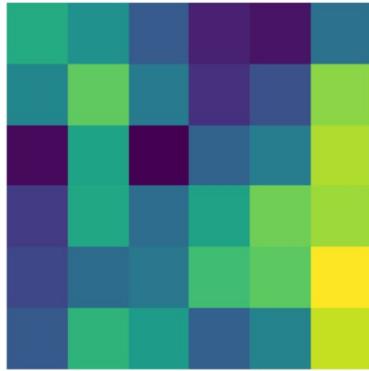
influence



7x7 Gradient Map

Activations and Gradients

How important is this map?



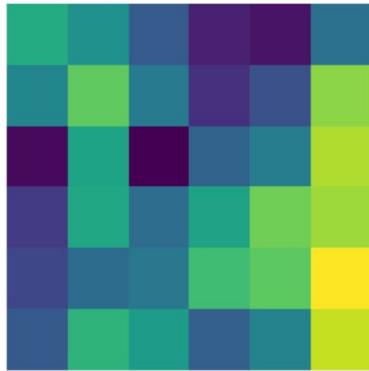
7x7 Feature Map



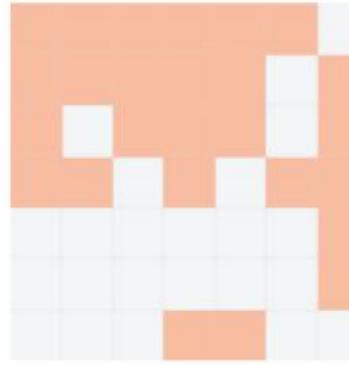
7x7 Gradient Map

Activations and Gradients

How important is this map?



7x7 Feature Map

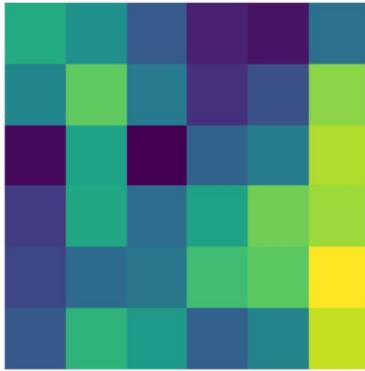


7x7 Gradient Map

```
weight = torch.mean(gradients, dim=(1, 2))
```

Activations and Gradients

How important is this map?



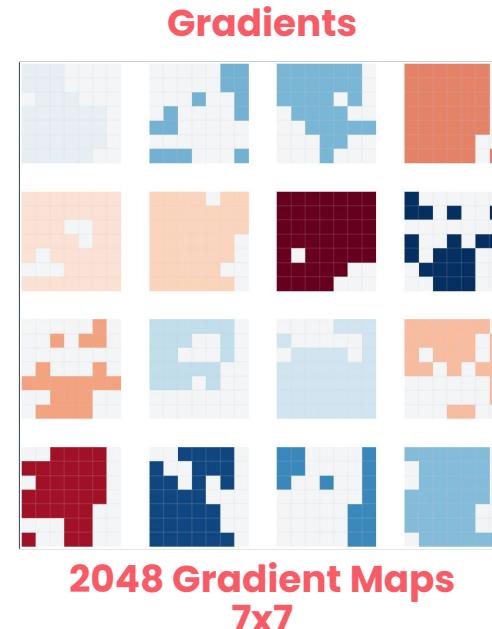
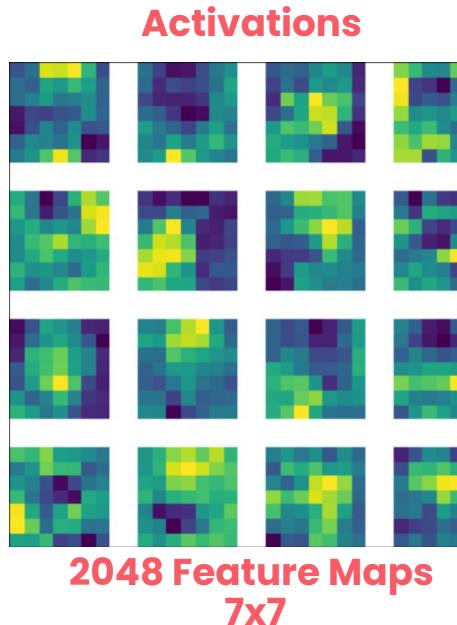
7x7 Feature Map

Average Influence

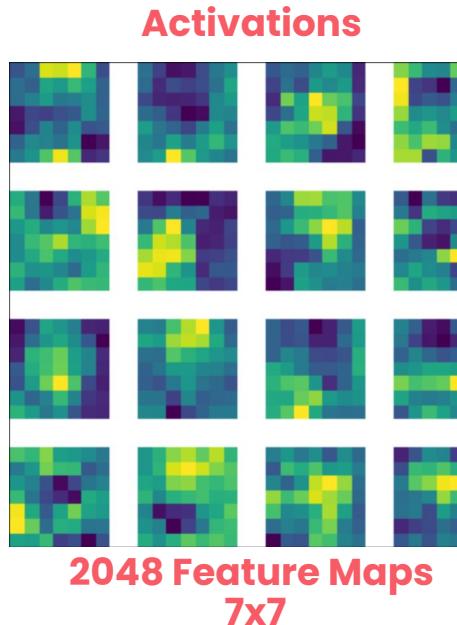
.004

```
weight = torch.mean(gradients, dim=(1, 2))
```

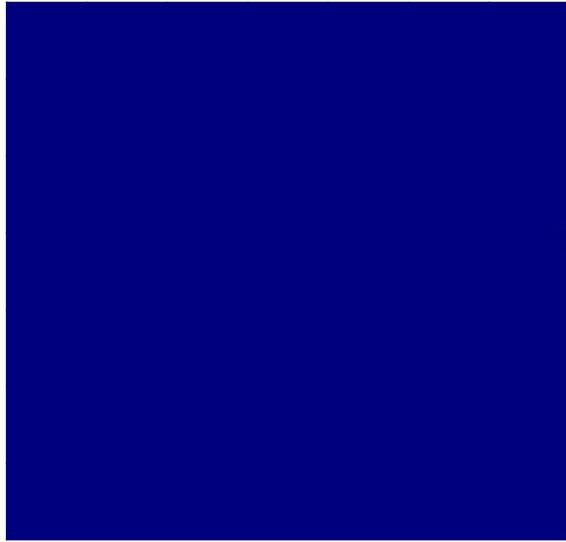
Activations and Gradients



Activations and Gradients

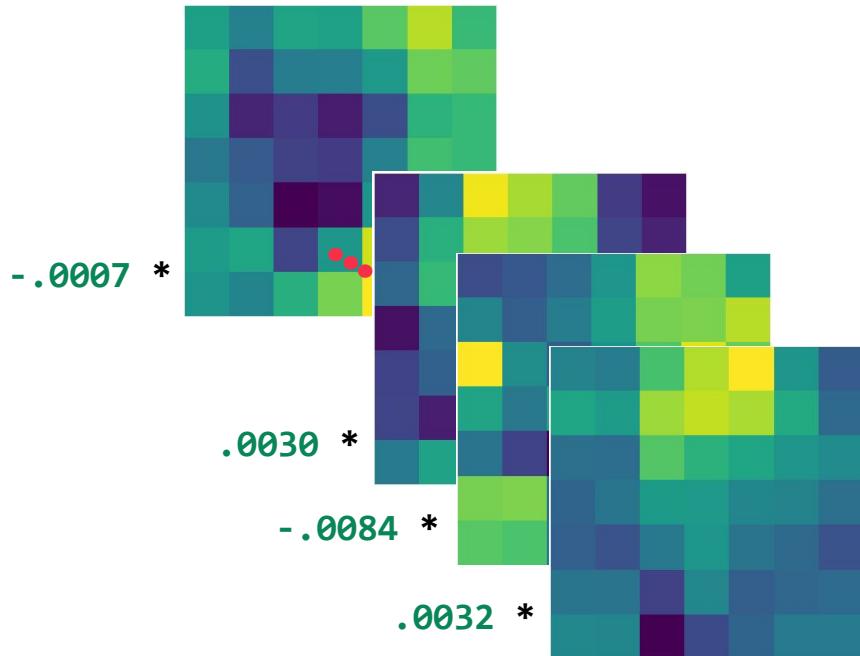


Weighted Maps

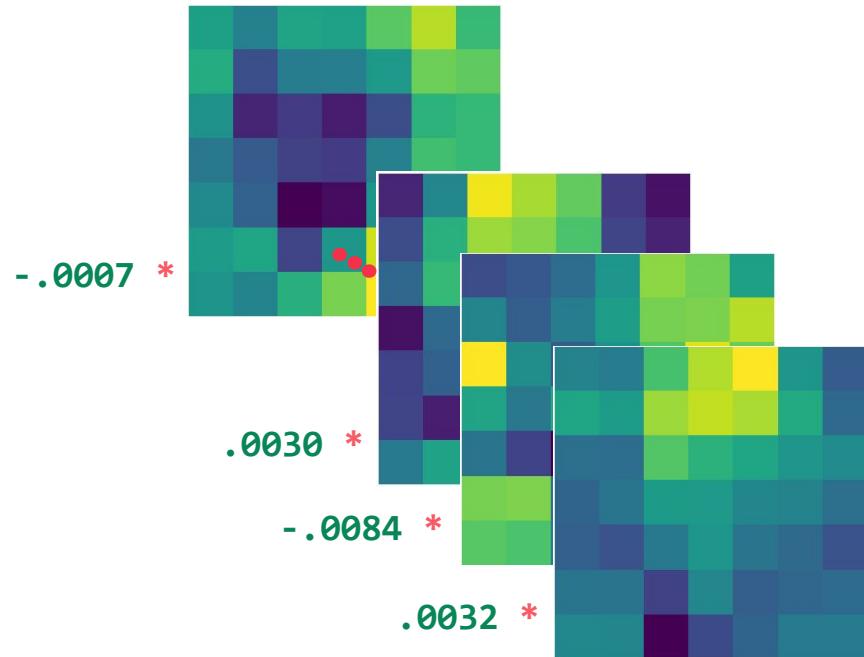


```
cam = torch.zeros(  
    activations.shape[1:],  
    dtype=activations.dtype).to(x.device)
```

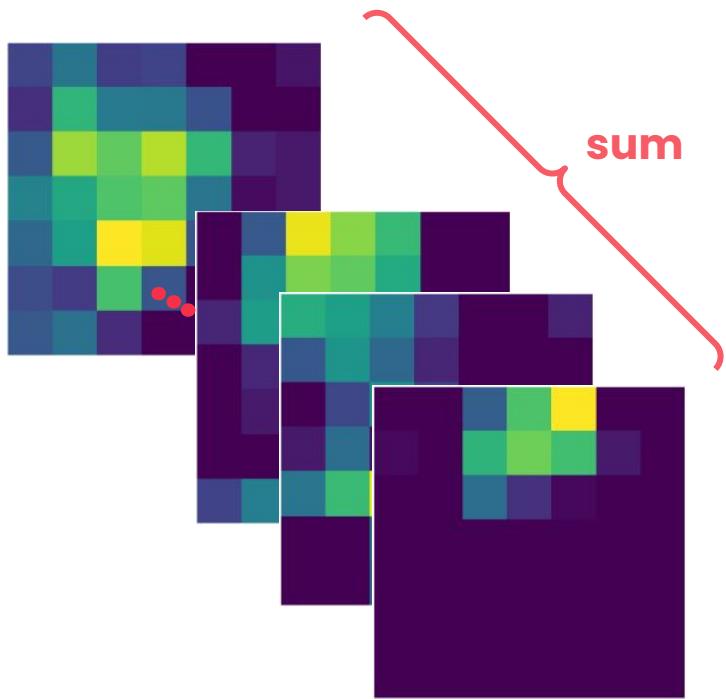
Weighted Maps



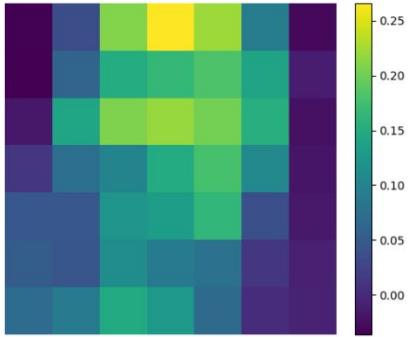
Weighted Maps



Weighted Maps

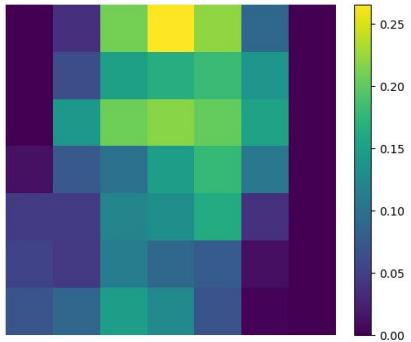


CAM Prep



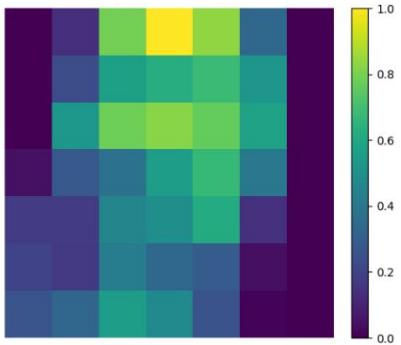
```
for i, w in enumerate(weights):
    cam += w * activations[i]
```

CAM Prep



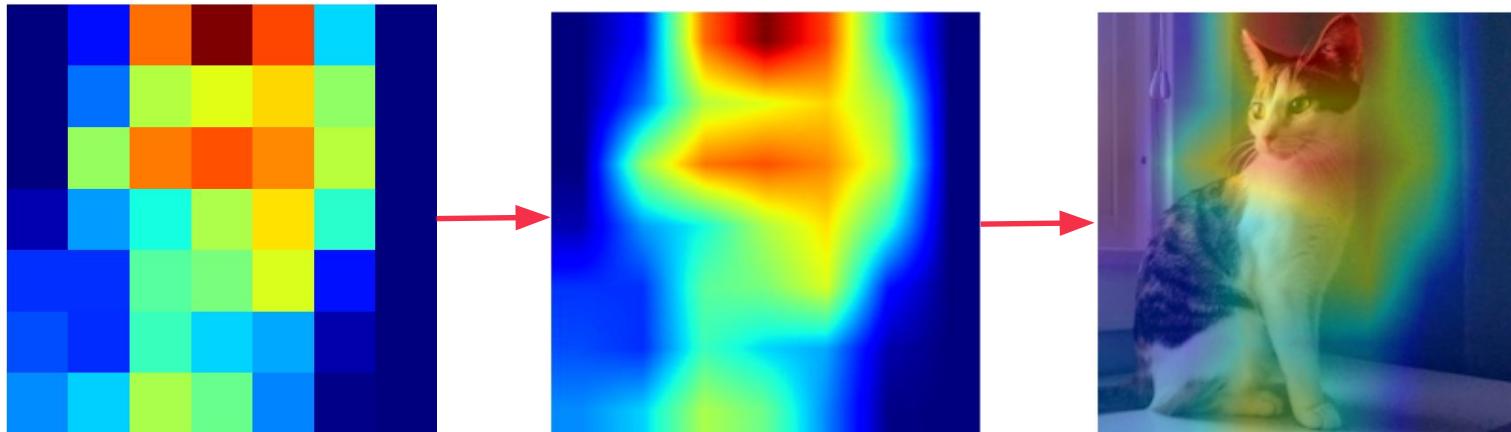
```
for i, w in enumerate(weights):
    cam += w * activations[i]
    # Keep positive values
    cam = torch.maximum(cam, torch.tensor(0, dtype=cam.dtype).to(x.device))
```

CAM Prep



```
for i, w in enumerate(weights):
    cam += w * activations[i]
    # Keep positive values
    cam = torch.maximum(cam, torch.tensor(0, dtype=cam.dtype).to(x.device))
    cam_np = cam.cpu().numpy()
    cam_np = (cam_np - cam_np.min()) / (cam_np.max() - cam_np.min() + 1e-8)
```

CAM Prep





DeepLearning.AI

Diffusion

Specialized Approaches to Vision

Classifying and Interpreting Images

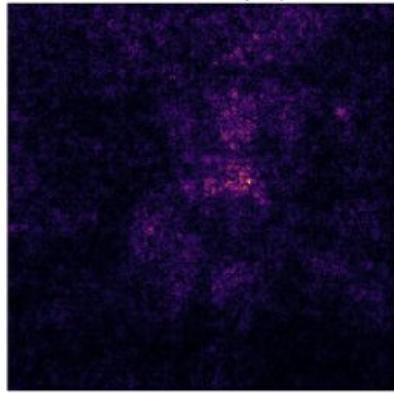
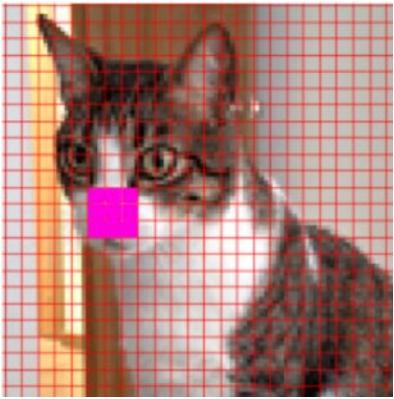


Image Generation

“A puppy riding a skateboard in Times Square”



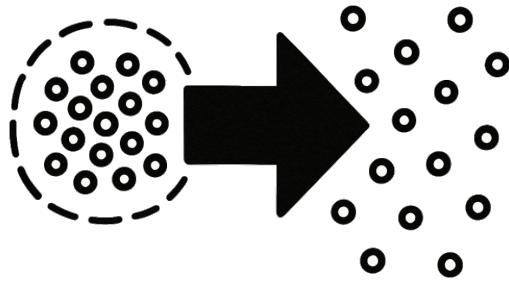
Early Image Generation

- **Produced blurry images.**
- **Unstable to train.**
- **Too slow for practical use.**



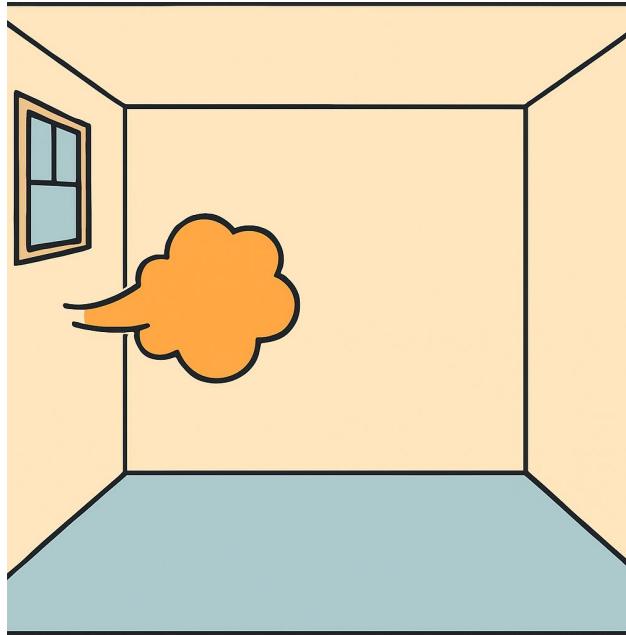
Modern Image Generation

Thermodynamics

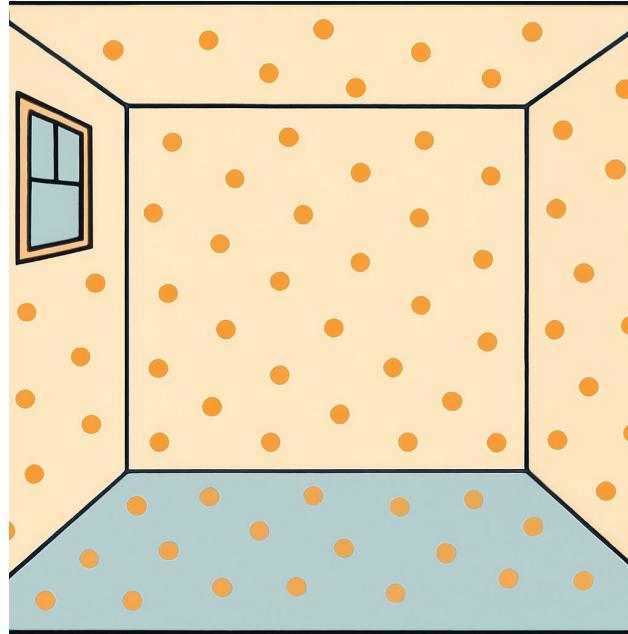


Diffusion

Diffusion

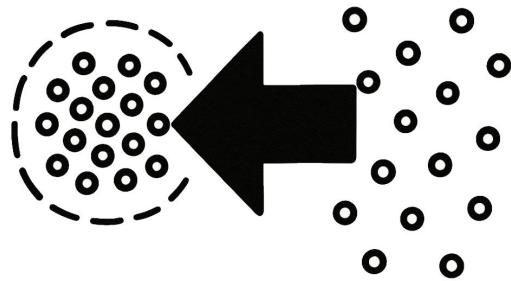


Diffusion



Follows probabilistic rules

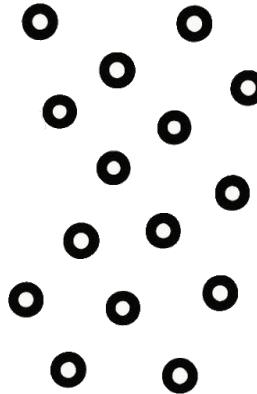
Reversing Diffusion



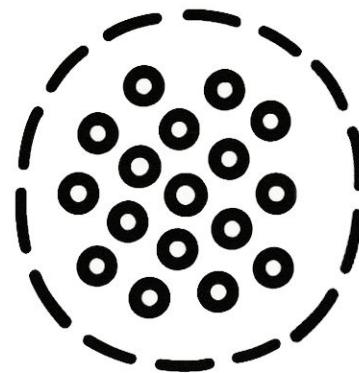
Probabilistic Rules

Reversing Diffusion

R: 102 G: 179 B: 92	R: 14 G: 106 B: 71	R: 188 G: 20 B: 102	R: 121 G: 210 B: 214	R: 74 G: 202 B: 87
R: 116 G: 99 B: 103	R: 151 G: 130 B: 149	R: 52 G: 1 B: 87	R: 235 G: 157 B: 37	R: 129 G: 191 B: 187
R: 20 G: 160 B: 203	R: 57 G: 21 B: 252	R: 235 G: 88 B: 48	R: 218 G: 58 B: 254	R: 169 G: 255 B: 219
R: 187 G: 207 B: 14	R: 189 G: 189 B: 174	R: 189 G: 50 B: 107	R: 54 G: 243 B: 63	R: 248 G: 130 B: 228
R: 50 G: 134 B: 20	R: 72 G: 166 B: 17	R: 131 G: 88 B: 59	R: 13 G: 241 B: 249	R: 8 G: 89 B: 52



Reversing Diffusion



Forward Diffusion



R:255 G:255 B:253	R:255 G:255 B:254	R:255 G:255 B:252	R:255 G:255 B:249	R:255 G:255 B:242
R:255 G:255 B:255	R:255 G:255 B:255	R:255 G:255 B:251	R:255 G:255 B:249	R:255 G:255 B:244
R:210 G:195 B:170	R:255 G:255 B:255	R:255 G:255 B:255	R:255 G:255 B:250	R:255 G:255 B:244
R: 87 G: 43 B: 11	R:164 G:146 B:121	R:255 G:255 B:252	R:255 G:255 B:255	R:255 G:255 B:247
R:136 G:104 B: 64	R: 85 G: 49 B: 10	R:143 G:119 B: 92	R:250 G:248 B:239	R:255 G:255 B:255



Latent Space

Forward Diffusion



Image Tensor

R G B
[[0.45, -0.10, 0.72],
 [-0.30, 0.80, -.05], +
 [0.00, 0.15, -0.60],
 . . .]

Random Gaussian Noise Tensor

[[0.01, -0.05, 0.22],
 [-0.03, 0.02, -.09],
 [0.18, -0.07, -0.04],
 . . .]

Forward Diffusion

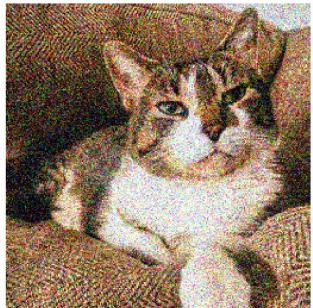


Image + Noise

R **G** **B**
[[0.46, -0.05, 0.94],
 [-0.27, 0.82, -.14],
 [0.18, 0.08, -0.64],
 . . .]

Forward Diffusion



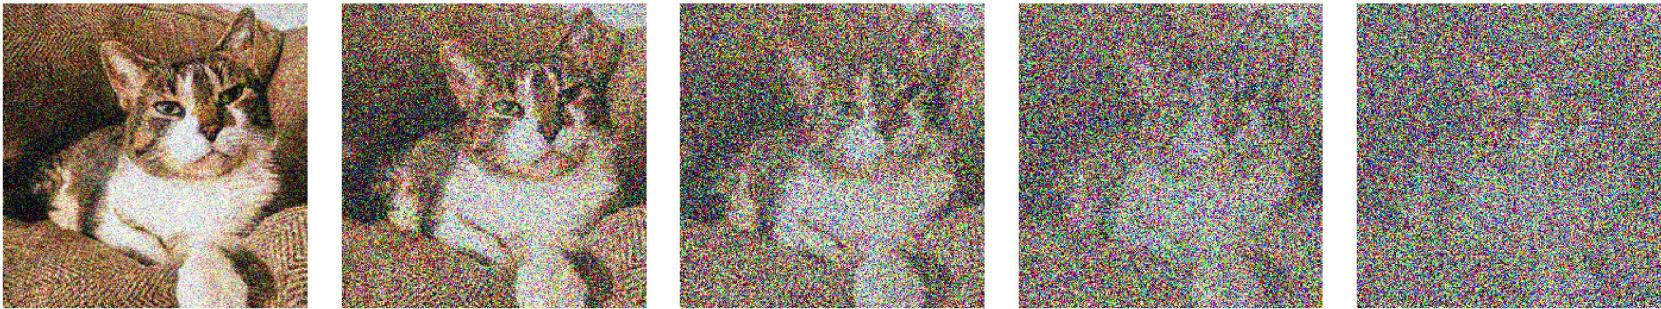
Image + Noise

R G B
[[0.46, -0.05, 0.94],
 [-0.27, 0.82, -.14],
 [0.18, 0.08, -0.64],
 . . .]

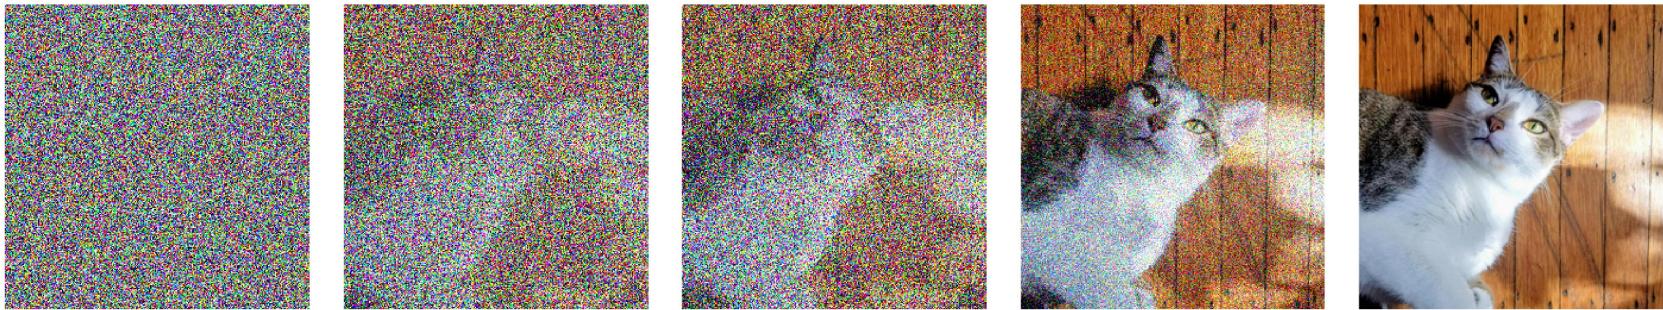
Random Gaussian Noise Tensor

[[0.01, -0.05, 0.22],
 [-0.03, 0.02, -.09],
 [0.18, -0.07, -0.04],
 . . .]

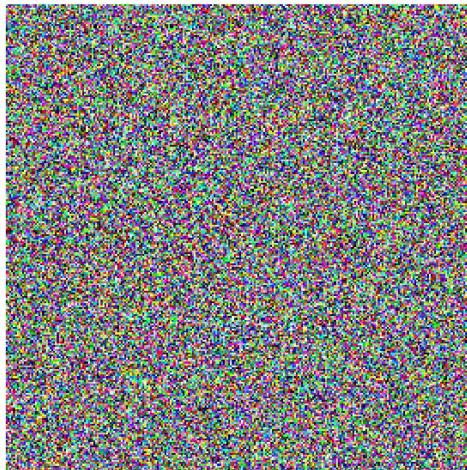
Forward Diffusion



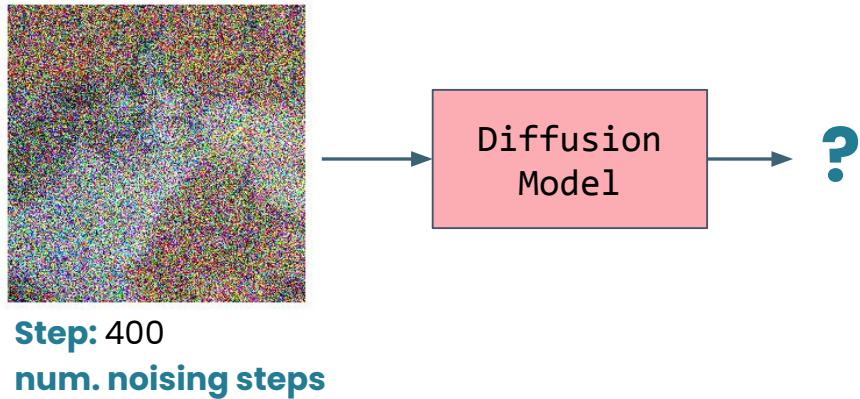
Diffusion



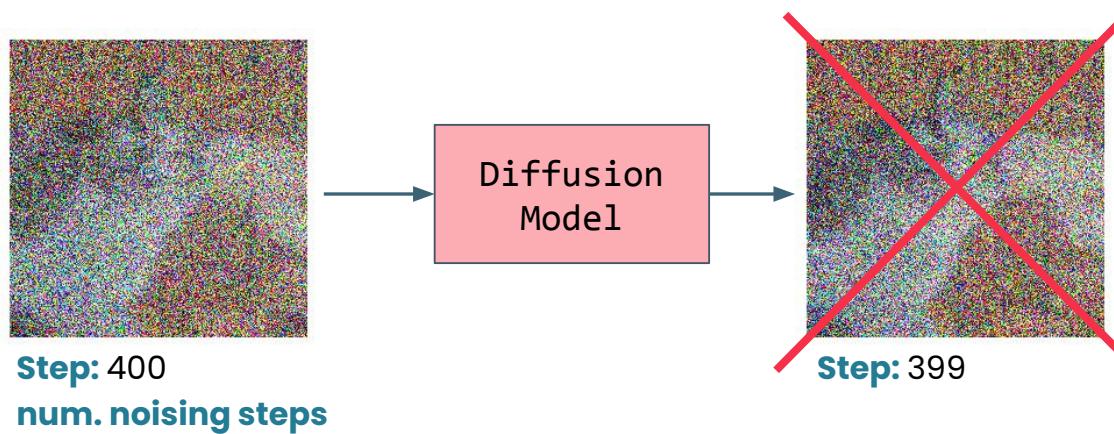
Diffusion



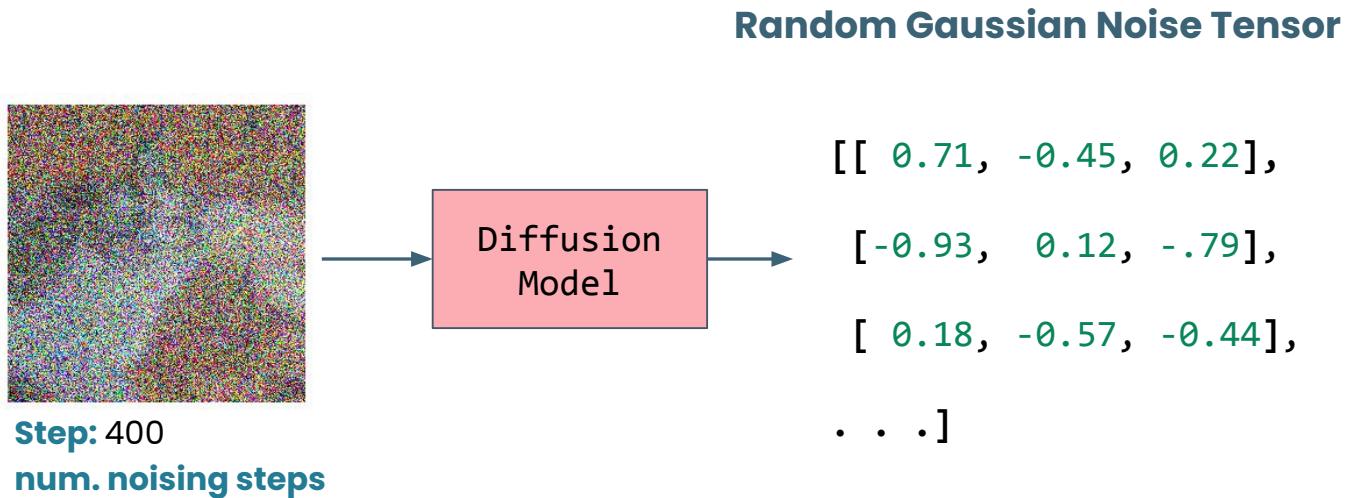
What does the model predict?



What does the model predict?

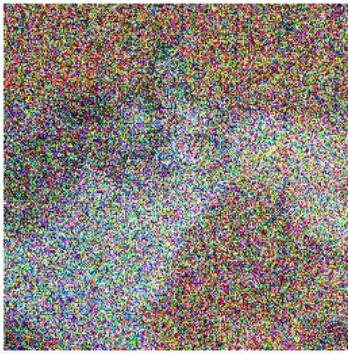


What does the model predict?



Predicting Noise

Current Image
Step: 400

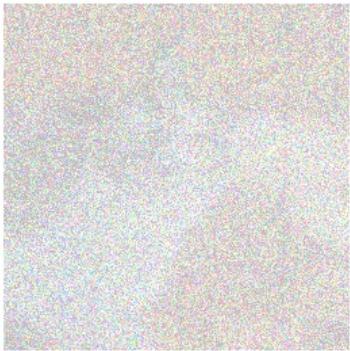


=

Earlier Image + Noise Tensor

Predicting Noise

Current Image
Step: 400



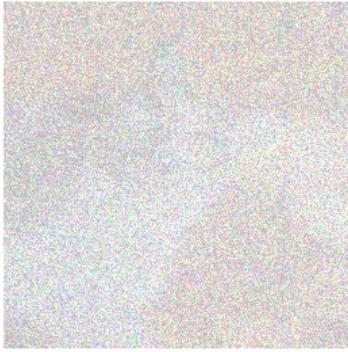
Model predicts:

=

Earlier Image + Noise Tensor

Predicting Noise

Current Image
Step: 400

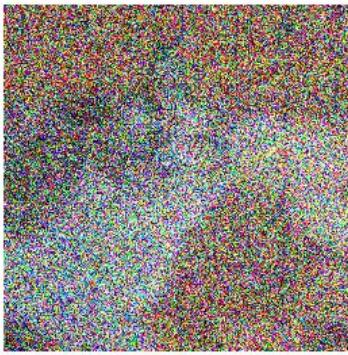


=

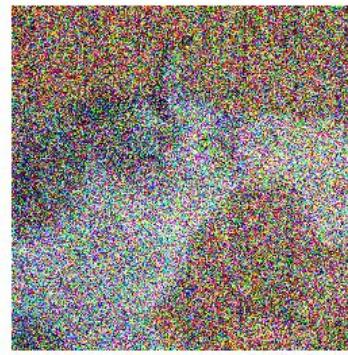
Earlier Image + Noise Tensor

Predicting Noise

Current Image
Step: 400



Earlier Image
Step: 399



=

+ **Noise Tensor**

Predicting Noise

Current Image
Step: 400



Earlier Image
Step: 399



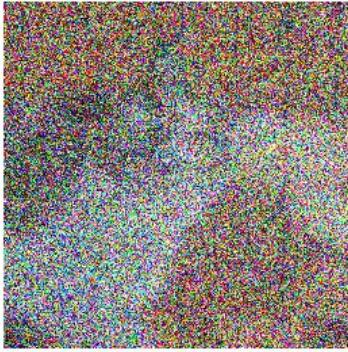
=

Small noise from previous

$\begin{bmatrix} -0.01, & -0.05, & 0.22 \\ -0.03, & 0.02, & -.09 \\ 0.18, & -0.07, & -0.04 \\ \dots \end{bmatrix}$

Predicting Noise

Current Image
Step: 400



From:
Original Image

=



Model predicts

Total Accumulated Noise

$[[1.71, -3.45, 2.22],$

$+ [-1.93, 2.12, -1.79],$

$[4.18, -1.57, -3.44],$

$\dots]$

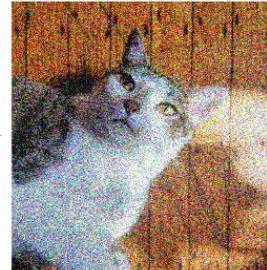
Training Data

Original Image



Time Steps: 125

Noisy image



$t=125$

Training Data

Original Image



Total Accumulated Noise

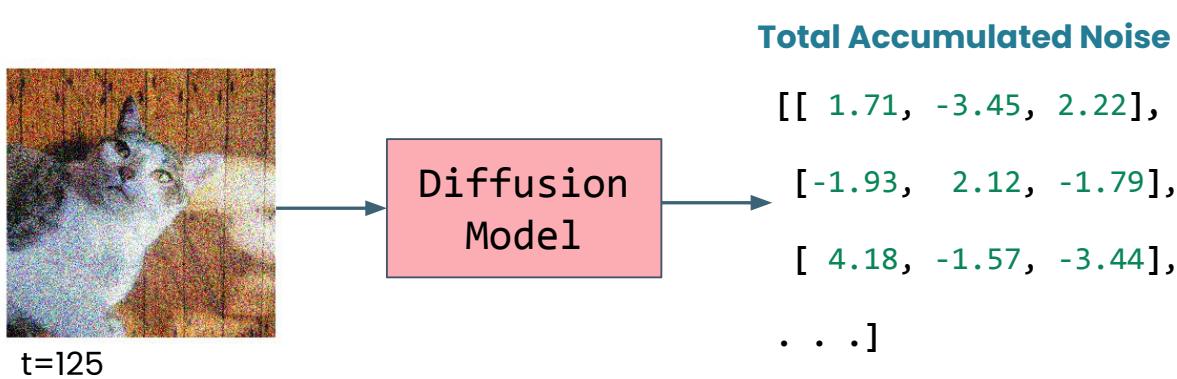
$$\begin{bmatrix} [1.71, -3.45, 2.22], \\ [-1.93, 2.12, -1.79], \\ [4.18, -1.57, -3.44], \\ \dots \end{bmatrix}$$


Noisy image

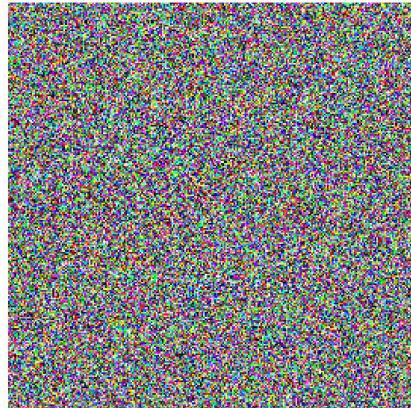


t=125

Training Data



Total Accumulated Noise



Total Accumulated Noise

$[[1.71, -3.45, 2.22],$

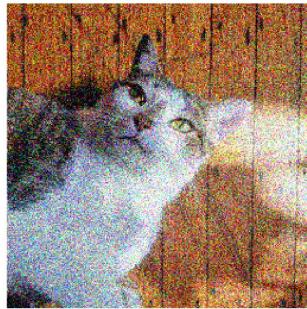
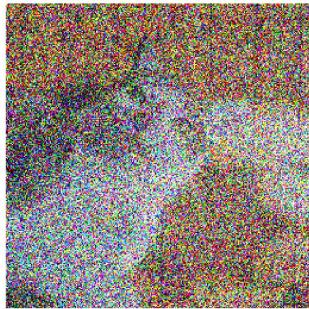
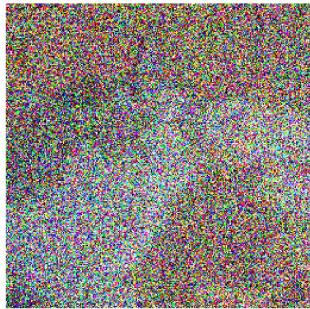
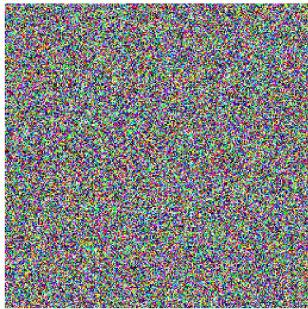
$- [-1.93, 2.12, -1.79],$

$[4.18, -1.57, -3.44],$

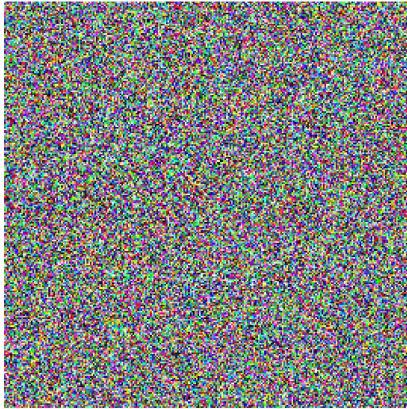
$\dots]$

$=$ **Clean image?**

Image Generation



Why Predict Total Accumulated Noise?



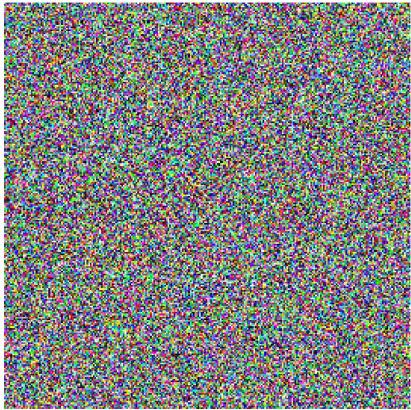
-

**Total
Accumulated
Noise**

=

?

Removing Total Noise

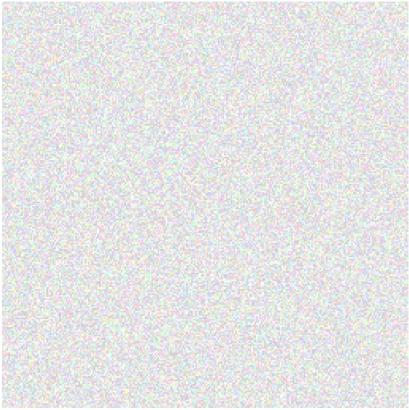


750 noise steps

$$\text{---} \quad \begin{matrix} \text{Total} \\ \text{Accumulated} \\ \text{Noise} \end{matrix} \quad =$$

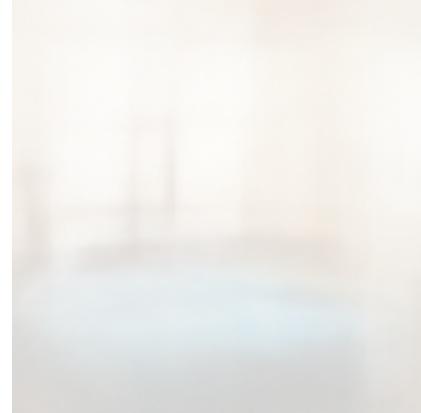


Removing Total Noise

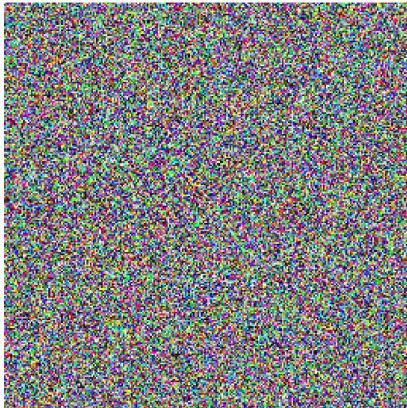


750 noise steps

**Total
Accumulated
Noise**



Gradual Generation



750 noise steps

**Total
Accumulated
Noise**

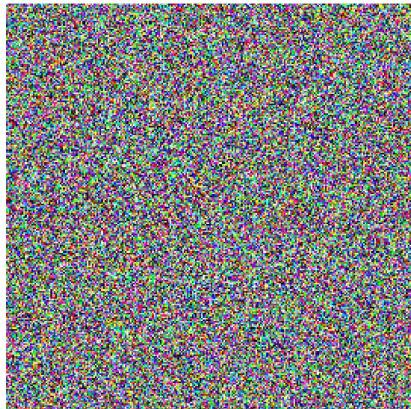
-

=

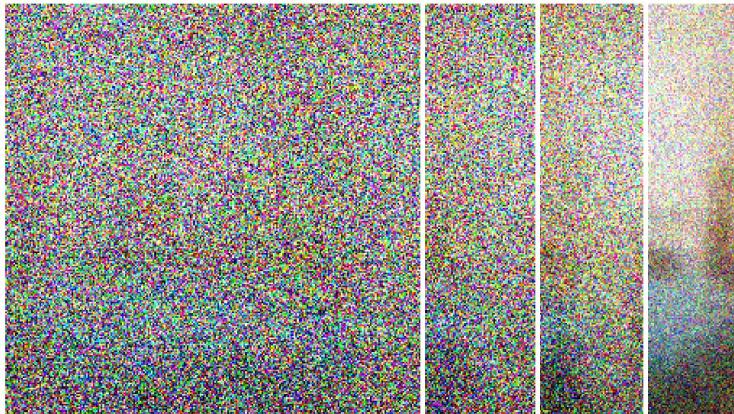
Predicted Original



Gradual Generation



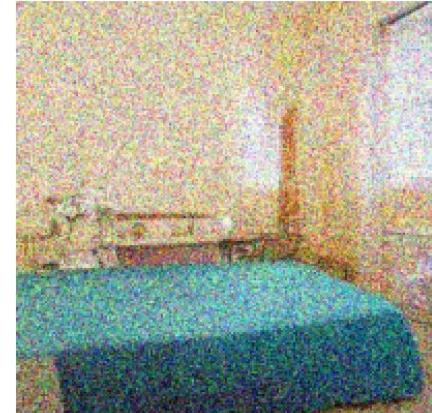
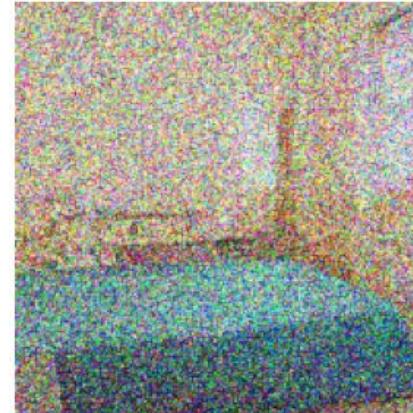
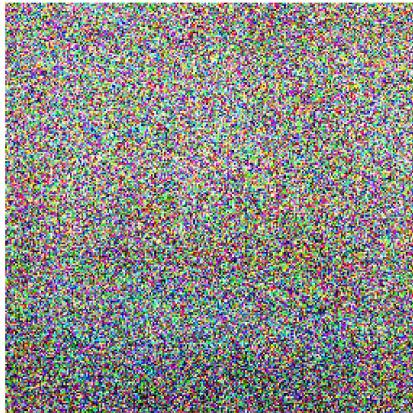
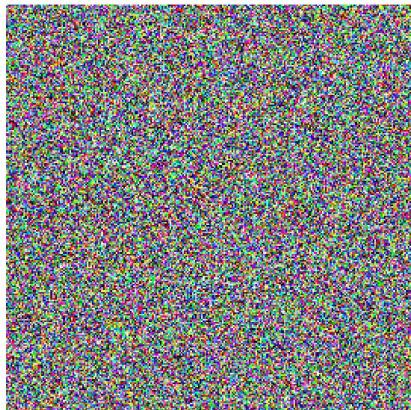
750 noise steps



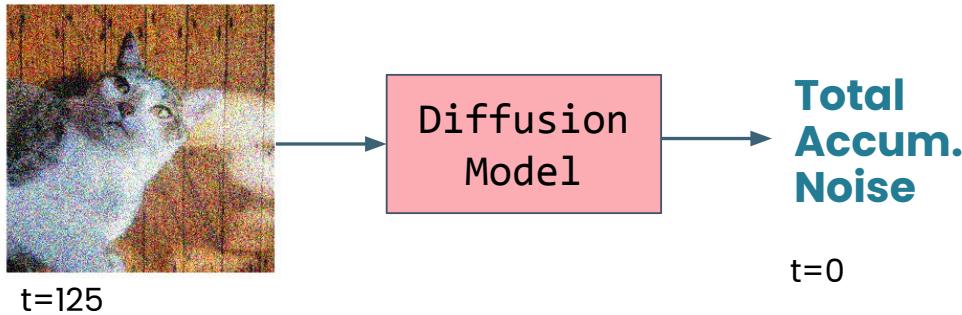
Predicted Original



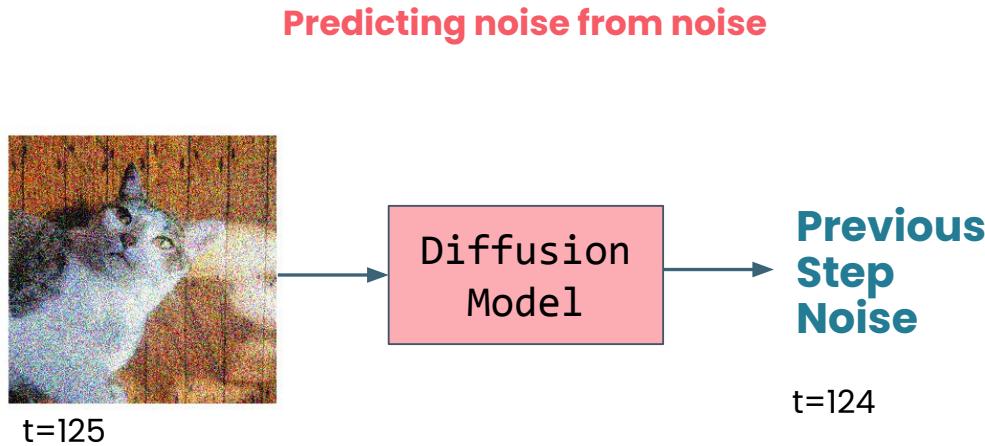
Gradual Generation



Why Predict Total Accumulated Noise?



Why Predict Total Accumulated Noise?



Why Predict Total Accumulated Noise?

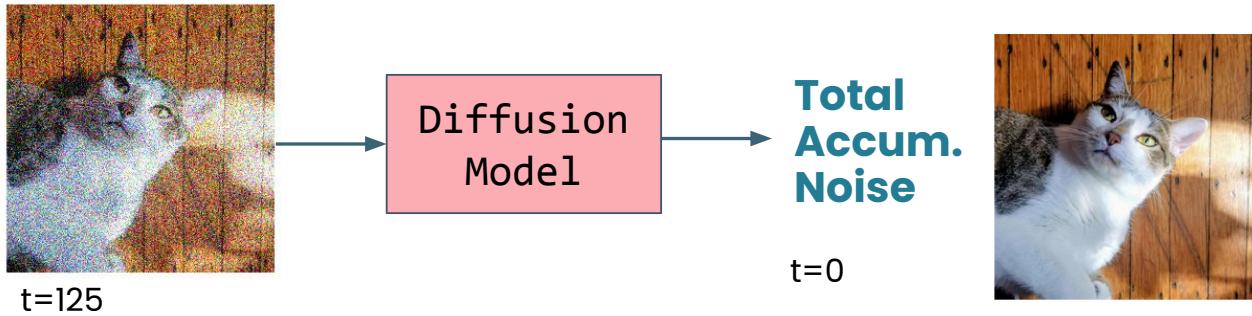
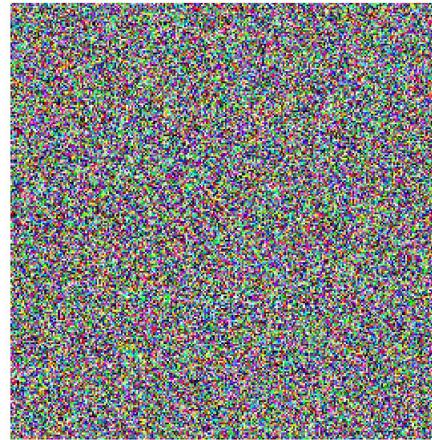


Image Generation

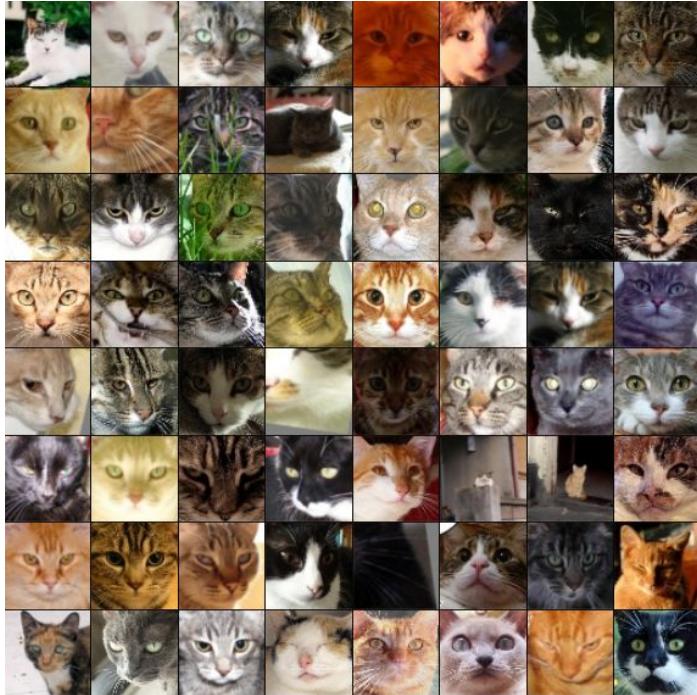
“A puppy riding a skateboard in Times Square”



Unconditional Models



Unconditional Models



Cat Dataset © 2022 AnnikaV9, licensed under the MIT License



Conditioning

“A puppy riding a skateboard in Times Square”



Step: 400



DeepLearning.AI

Image Generation Walkthrough

Specialized Approaches to Vision

Diffusion

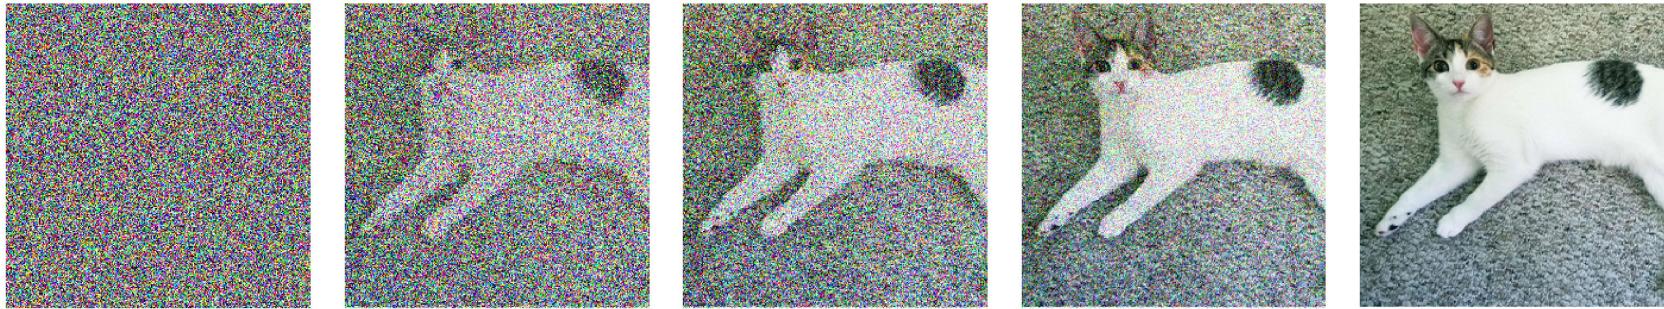


Image Generation

“A puppy riding a skateboard in Times Square”



Stable Diffusion

```
from diffusers import StableDiffusionPipeline
```

Stable Diffusion

```
from diffusers import StableDiffusionPipeline
```

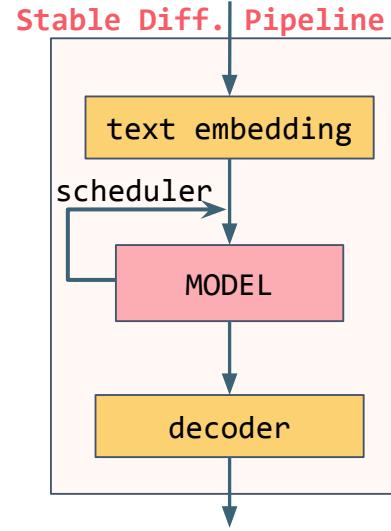
Stable Diffusion

```
from diffusers import StableDiffusionPipeline
```

Pipeline

```
from diffusers import StableDiffusionPipeline
```

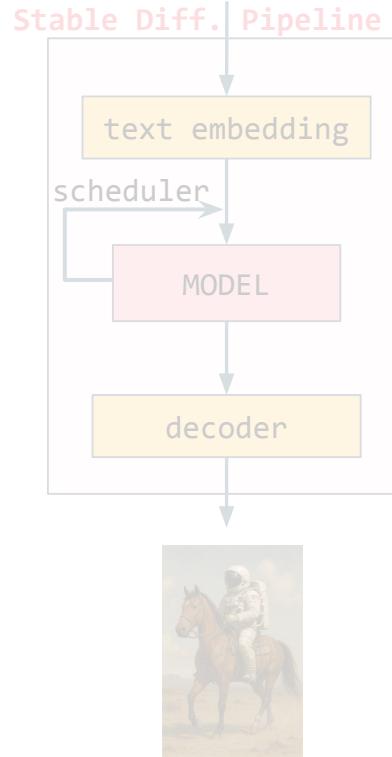
"An astronaut riding a horse"



Pipeline

```
from diffusers import StableDiffusionPipeline
```

“An astronaut riding a horse”



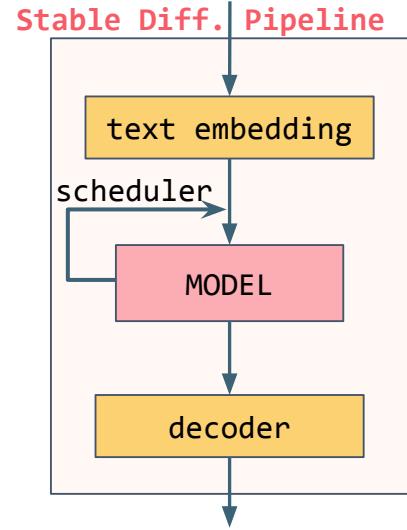
Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
"stable-diffusion-3.5-medium"  
"stable-diffusion-3.5-large"  
"stable-diffusion-3.5-large-turbo"
```

Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16  
                                              variant="fp16",  
                                              cache_dir=".models")
```

"An astronaut riding a horse"

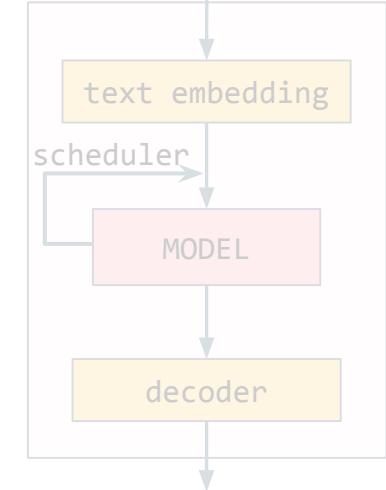


Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16  
                                              variant="fp16",  
                                              cache_dir=".models")
```

"An astronaut riding a horse"

Stable Diff. Pipeline

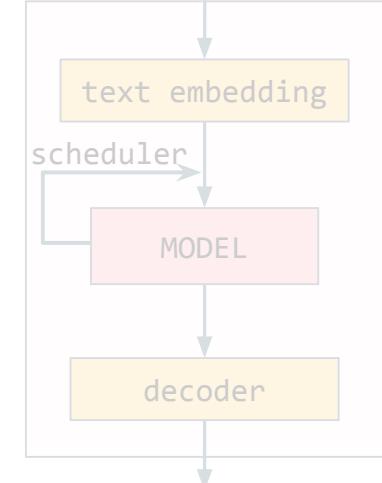


Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16,  
                                              variant="fp16",  
                                              cache_dir=".models")
```

"An astronaut riding a horse"

Stable Diff. Pipeline

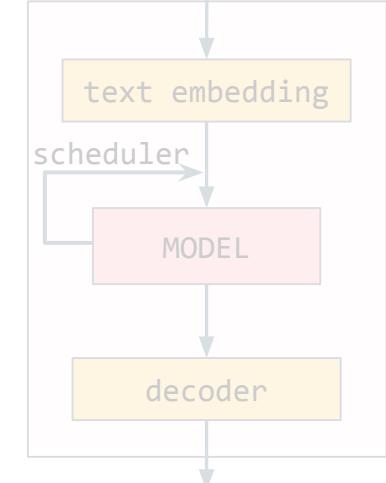


Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16  
                                              variant="fp16",  
                                              cache_dir=".models")
```

"An astronaut riding a horse"

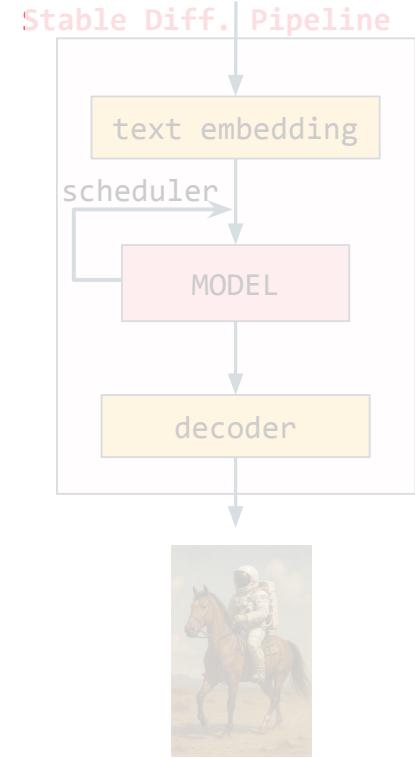
Stable Diff. Pipeline



Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16  
                                              variant="fp16",  
                                              cache_dir="../models")
```

"An astronaut riding a horse"

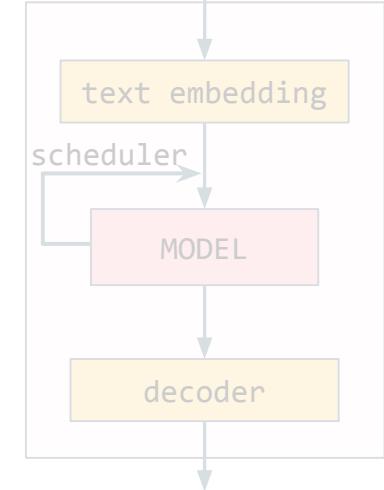


Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id,  
                                              torch_dtype=torch.float16,  
                                              variant="fp16",  
                                              cache_dir=".models")
```

"An astronaut riding a horse"

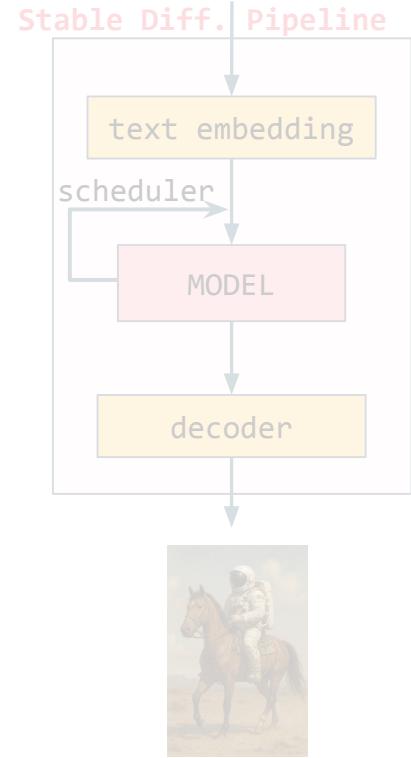
Stable Diff. Pipeline



Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)
```

"An astronaut riding a horse"



Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)
```



Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)  
pipe = pipe.to(device)  
prompt = "A puppy riding a skateboard in Times Square."  
  
images = pipe(  
    prompt,  
    num_inference_steps=40,  
    generator=generator  
).images
```

Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)  
pipe = pipe.to(device)  
prompt = "A puppy riding a skateboard in Times Square."  
  
images = pipe(  
    prompt,  
    num_inference_steps=40,  
    generator=generator  
).images
```

Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)  
pipe = pipe.to(device)  
prompt = "A puppy riding a skateboard in Times Square."  
  
images = pipe(  
    prompt,  
    num_inference_steps=40,  
    generator=generator  
).images
```

Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)  
pipe = pipe.to(device)  
prompt = "A puppy riding a skateboard in Times Square."  
  
images = pipe(  
    prompt,  
    num_inference_steps=40,      more steps = more refinement  
    generator=generator  
).images
```

Diffusers Library

```
from diffusers import StableDiffusionPipeline  
  
repo_id = "stabilityai/stable-diffusion-2-base"  
  
pipe = StableDiffusionPipeline.from_pretrained(repo_id, . . .)  
  
generator = torch.Generator(device=device).manual_seed(42)  
pipe = pipe.to(device)  
prompt = "A puppy riding a skateboard in Times Square."  
  
images = pipe(  
    prompt,  
    num_inference_steps=40,  
    generator=generator  
).images
```

Stable Diffusion

```
generator = torch.Generator(device=device).manual_seed(42)
```



Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im,  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im,  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=3,  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```



Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

```
prompt = "A cute dog with a red bandana, sitting in a lush park"
```

Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

prompt = "A surreal landscape with floating clocks, melting trees, and a purple sky, in the style of Salvador Dali"

guidance_scale = 5



Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

```
neg_prompt = "deformed hands"
```



Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs,  
    negative_prompt=neg_prompt  
).images
```

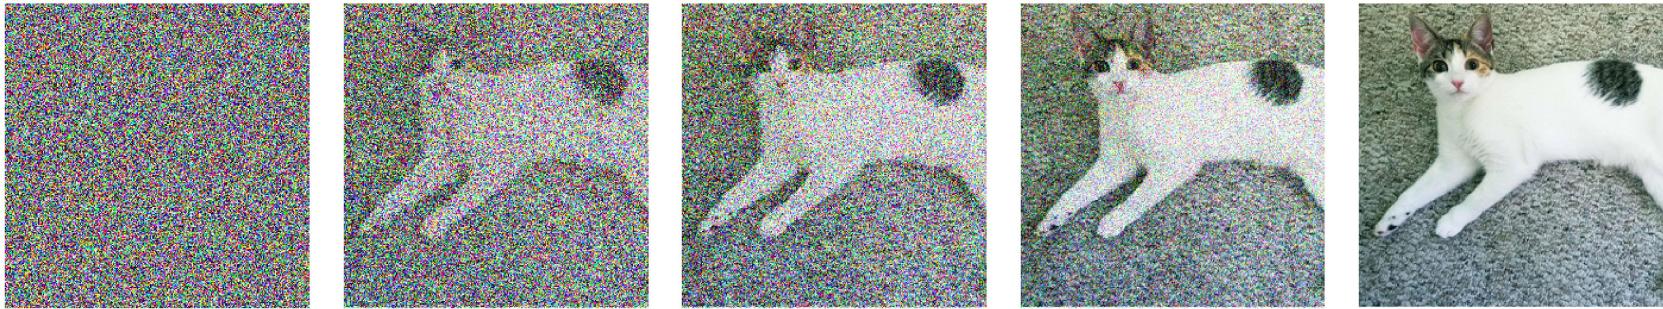
```
prompt = "Teddy Bears eating pizza at a birthday party"  
neg_prompt = "no pepperoni pizza"
```



Parameters

```
image_list = pipe(  
    prompt,  
    num_images_per_prompt=num_im,  
    num_inference_steps=num_inf,  
    generator=generator,  
    guidance_scale=gs  
    negative_prompt=neg_prompt  
).images
```

Visualizing Denoising



Latent Space



Latent Space



- **Smaller**
- **Faster and Uses Less Memory**
- **Scalable and Practical**
- **Check out the resources!**