DeepLearning.AI

Transformers

Specialized Approaches to
Natural Language Processing in Pytorch

Module 2

Specialized Approaches to **Vision**

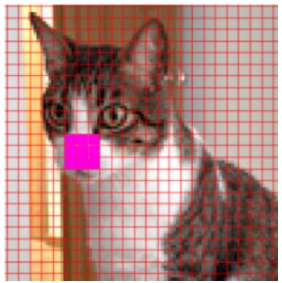
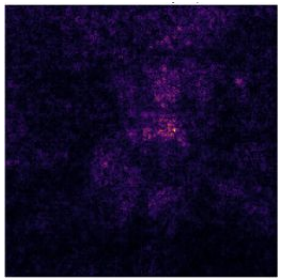

Module 3

Specialized Approaches to **Language**

Transformers = Backbone of modern NLP

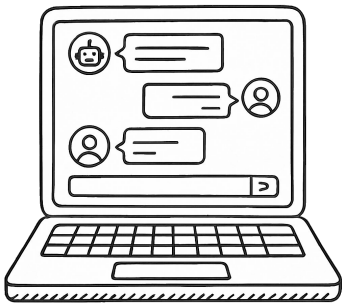

Tokenization and Embeddings

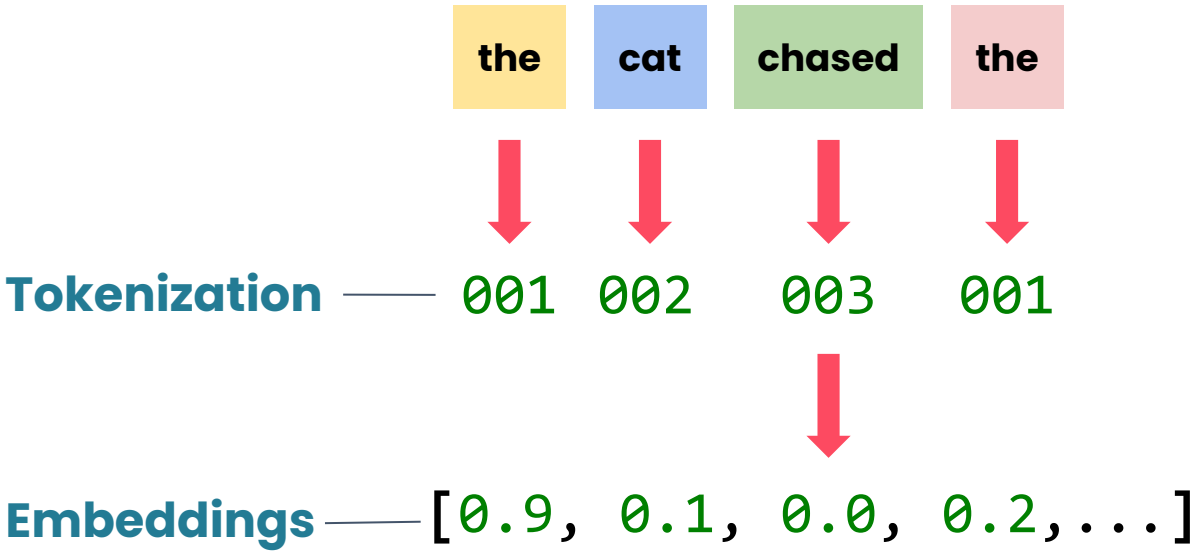

Pre-Trained Models

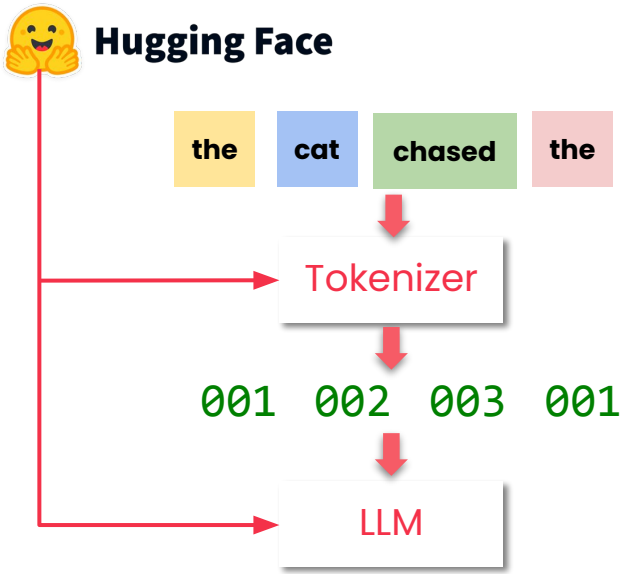


```python
from transformers import AutoTokenizer,
AutoModelForCausalLM

# Load a pretrained GPT-2 model and its tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Encode a short prompt
inputs = tokenizer("The cat chased the",
                   return_tensors="pt")

# Generate up to 10 tokens
outputs = model.generate(**inputs, max_length=10)
```

Pre-Trained Models

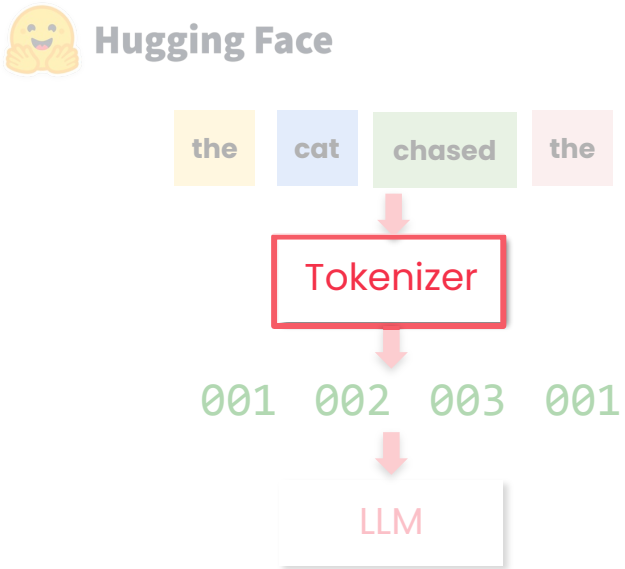


```
from transformers import AutoTokenizer,
AutoModelForCausalLM

# Load a pretrained GPT-2 model and its tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Encode a short prompt
inputs = tokenizer("The cat chased the",
                   return_tensors="pt")

# Generate up to 10 tokens
outputs = model.generate(**inputs, max_length=10)
```

Pre-Trained Models

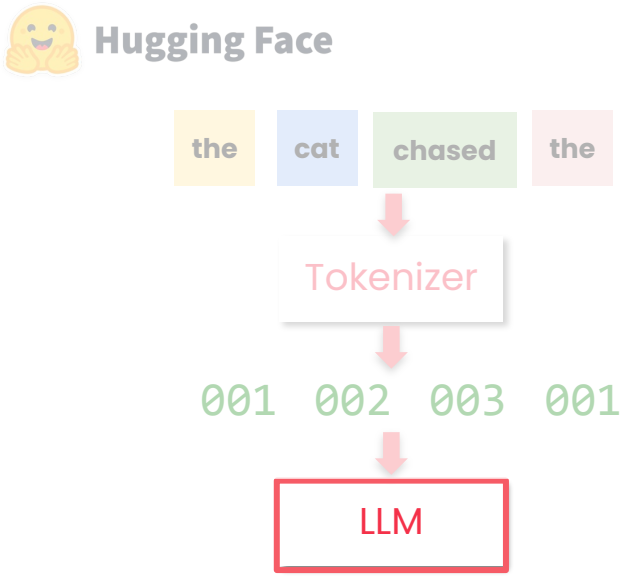


```
from transformers import AutoTokenizer,
AutoModelForCausalLM

# Load a pretrained GPT-2 model and its tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Encode a short prompt
inputs = tokenizer("The cat chased the",
                   return_tensors="pt")

# Generate up to 10 tokens
outputs = model.generate(**inputs, max_length=10)
```

Pre-Trained Models

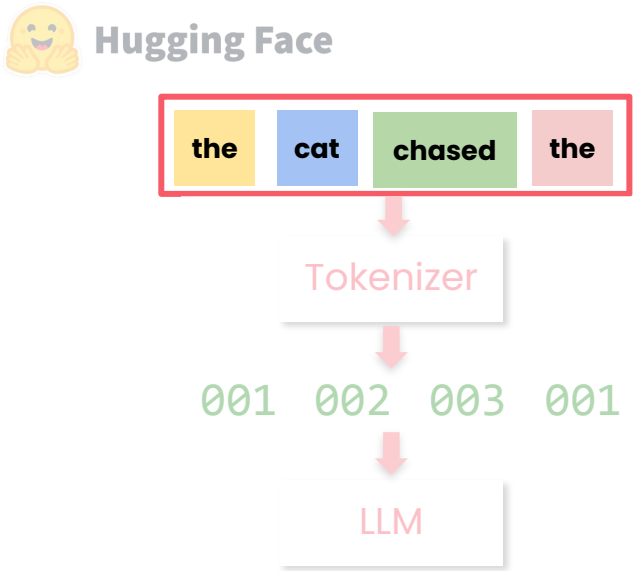


```
from transformers import AutoTokenizer,
AutoModelForCausalLM

# Load a pretrained GPT-2 model and its tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Encode a short prompt
inputs = tokenizer("The cat chased the",
                   return_tensors="pt")

# Generate up to 10 tokens
outputs = model.generate(**inputs, max_length=10)
```

Pre-Trained Models

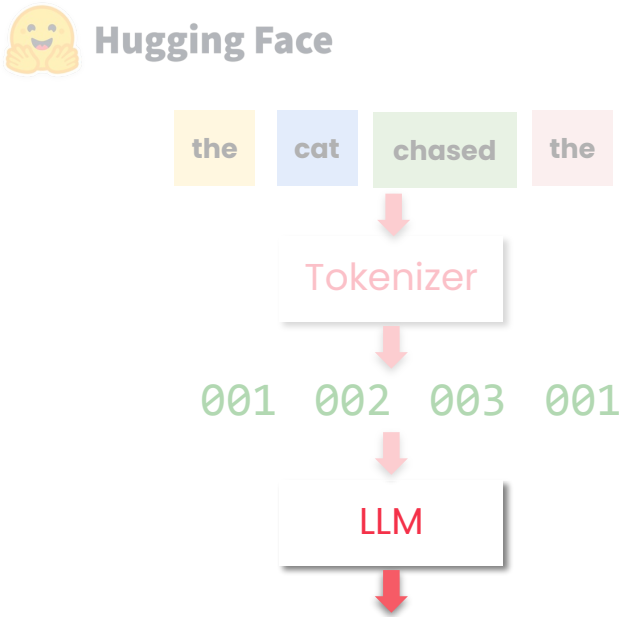

“The cat chased the mouse . . .”

```
from transformers import AutoTokenizer,
AutoModelForCausalLM

# Load a pretrained GPT-2 model and its tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Encode a short prompt
inputs = tokenizer("The cat chased the",
                   return_tensors="pt")

# Generate up to 10 tokens
outputs = model.generate(**inputs, max_length=10)
```

Transformers

You can't...

- Experiment
- Adapt
- Debug
- Explain

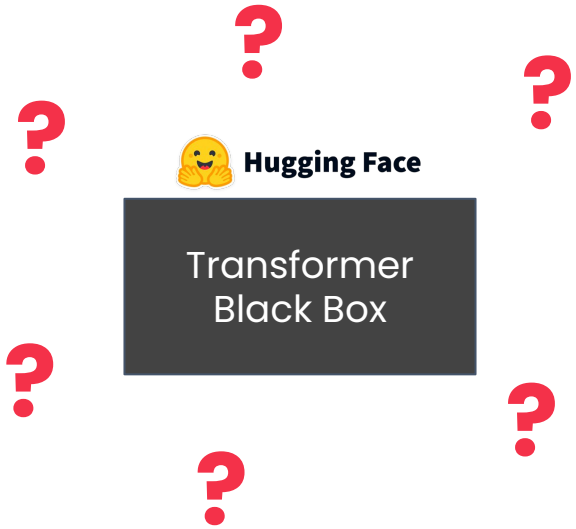

Transformers

Early Language Models

RNNs & LSTMs

The cat chased the mouse quickly across the yard

Transformers

The cat chased the mouse

Attention

“the cat chased the mouse”

Attention

"the **cat** chased the mouse"

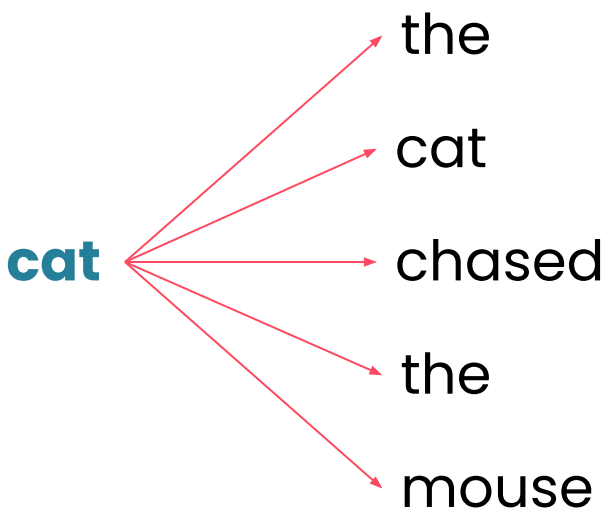

Attention

"the **cat** chased the mouse"

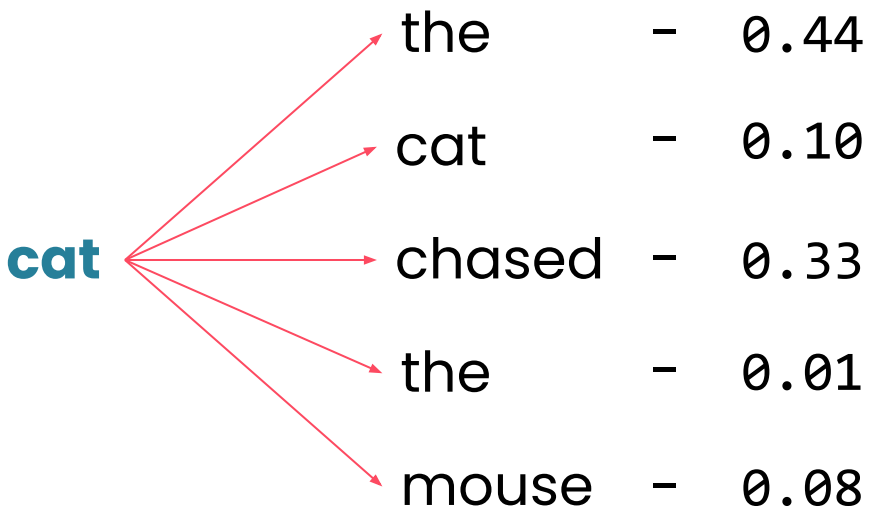

Attention

"the **cat** chased the mouse"

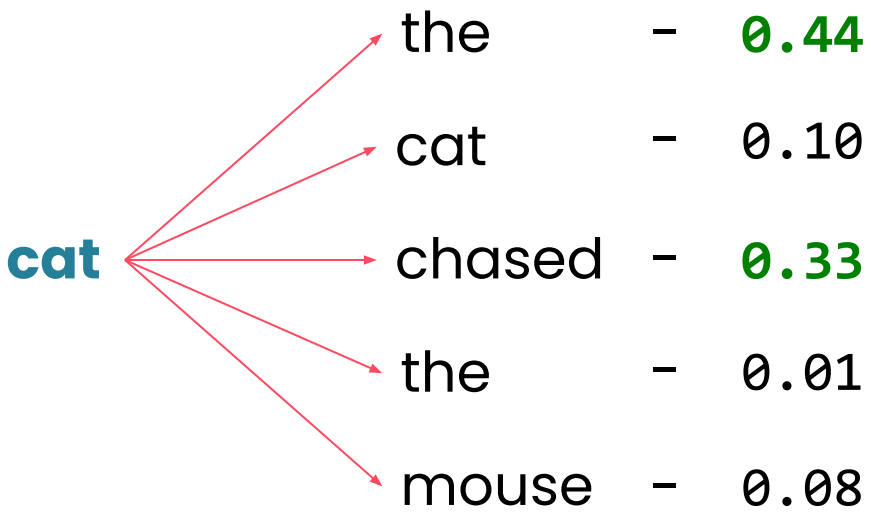

Attention

"the **cat** chased the mouse"

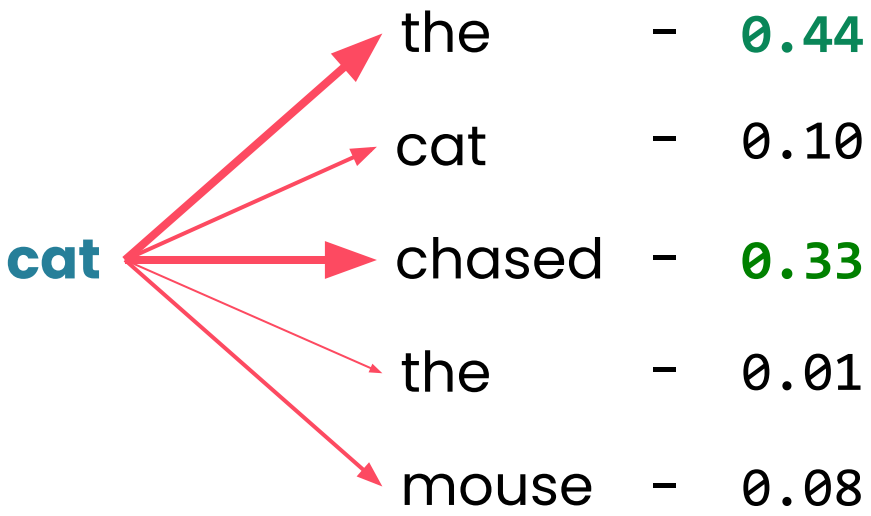

Attention

The animal didn’t cross the street because it was too wide

Attention

The animal didn't cross the street because it was too wide

Attention

The animal didn't cross the street because it was too wide

Attention

Attention

The animal didn’t cross the street because it was too wide

Attention


adjectives
The animal didn't cross the street because it was too wide

Attention


Subject / verb
The animal didn't cross the street because it was too wide

Attention Heads

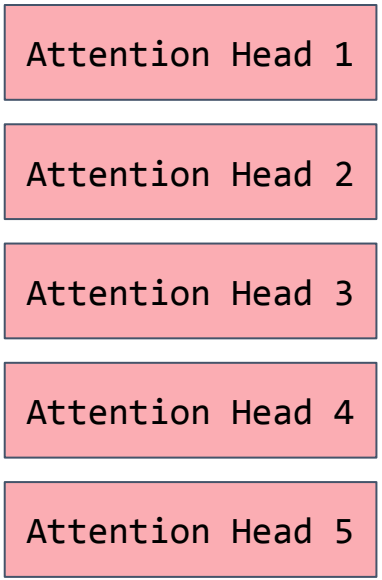

Each head learns different patterns

CNN Filters

Horizontal Edge

Vertical Edge

Ridge

Attention Heads

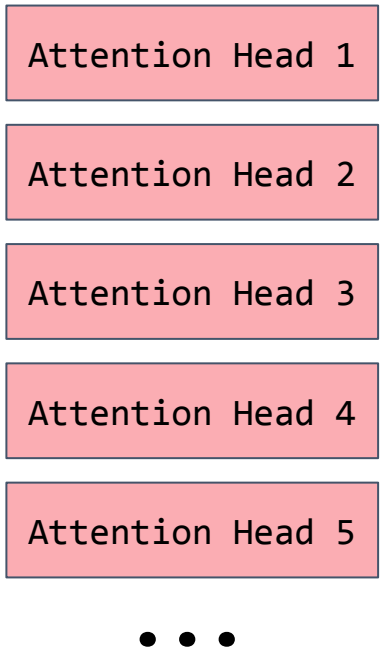

Attention Heads

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)
```

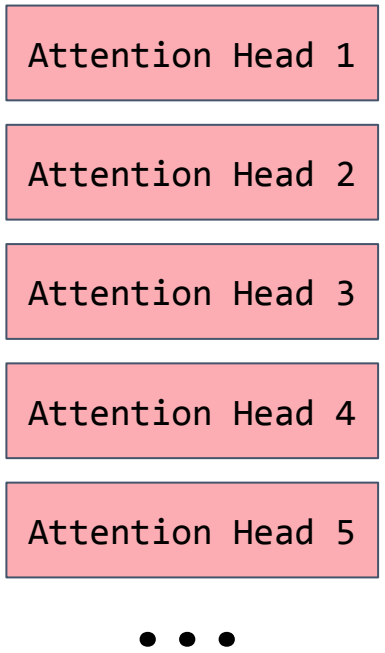

Position

Position: RNNs

Position: Transformers

The cat chased the dog

Position: Transformers

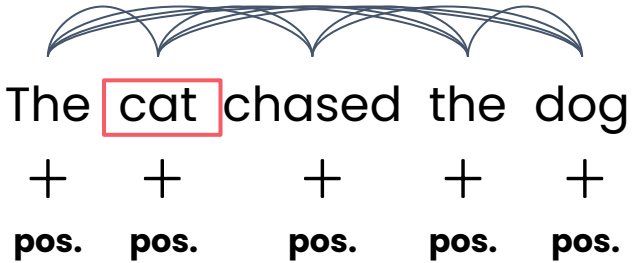

Position: Transformers

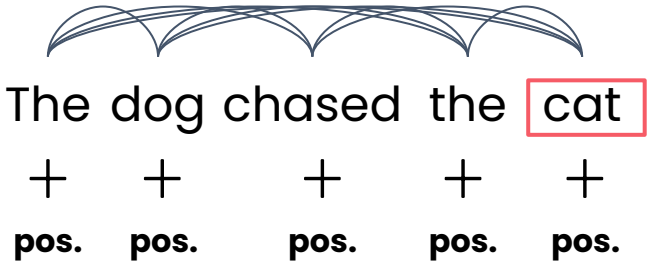

Position: Transformers

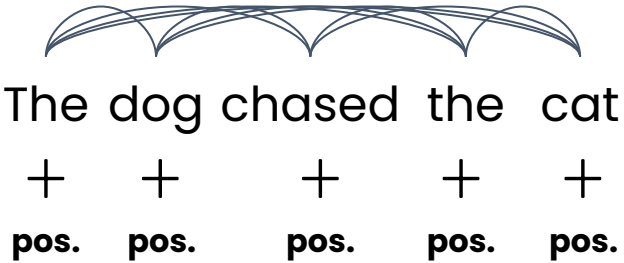

Positional info gives a sense of order

Transformer Architectures

- **Encoders**
- **Decoders**
- **Encoder-Decoders**

Encoders

- **Understanding**
- **Classification**
- **Named Entity**
- **BERT**

Decoders

Decoders

- **Generation**
- **Chatbots**
- **GPT**

Encoder - Decoders

- **Translation**
- **Summarization**

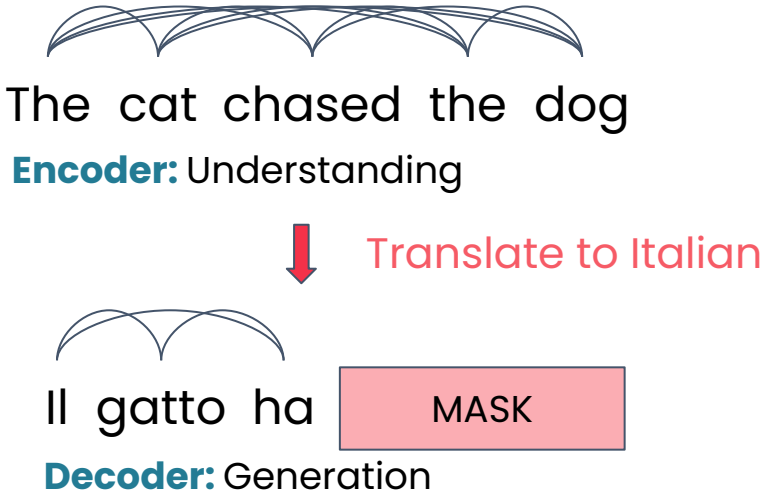

Transformer Architectures

- **Encoders:** Understanding
- **Decoders:** Generation
- **Encoder-Decoders:** Transformation

Transformer Architectures

- **Encoders:** Understanding
- **Decoders:** Generation
- **Encoder-Decoders:** Transformation

Laurence Moroney

Transformer Architectures

- **Encoders:** Understanding
- **Decoders:** Generation
- **Encoder-Decoders:** Transformation

Transformer Architectures

- **Encoders:** Understanding
- **Decoders:** Generation
- **Encoder-Decoders:** Transformation

Laurence Moroney

DeepLearning.AI

Attention

Specialized Approaches to
Natural Language Processing in Pytorch

Attention

Attention Head 1

Attention Head 2

Attention Head 3

Attention Head 4

Attention Head 5

• • •

Attention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)
```

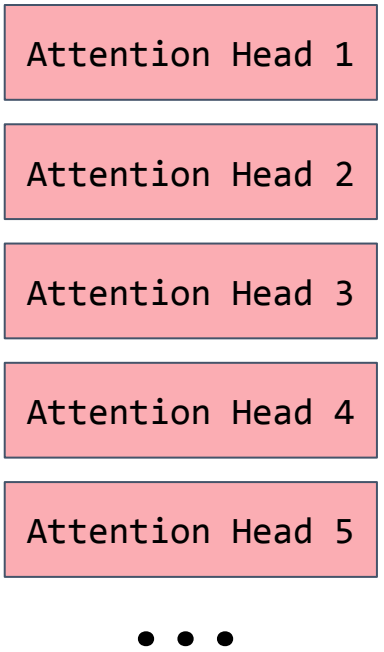

Attention

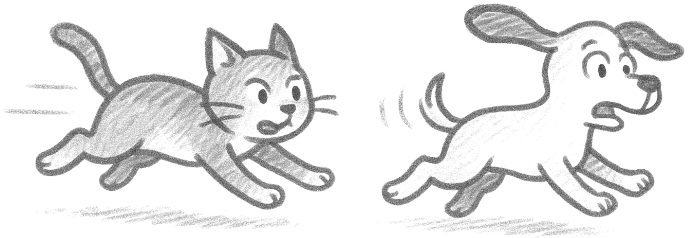

“The cat chased the dog”

Tokens

Tokens

Tokens

Tokens

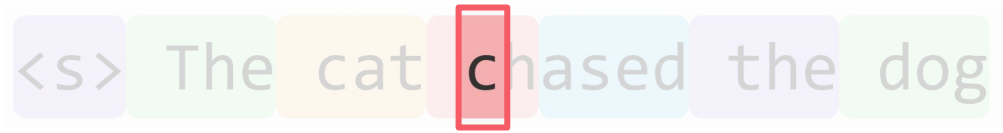

Tokens

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']]
```

Tokens

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']]

token_sequences = torch.tensor([[0, 1, 2, 0, 3]])
```

Tokens

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']]

token_sequences = torch.tensor([[0, 1, 2, 0, 3]])

tok_embed = nn.Embedding(num_embeddings=len(vocab), embedding_dim=4)
```

Tokens

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']]

token_sequences = torch.tensor([[0, 1, 2, 0, 3]])

tok_embed = nn.Embedding(num_embeddings=len(vocab), embedding_dim=4)
```

Tokens

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']]

token_sequences = torch.tensor([[0, 1, 2, 0, 3]])

tok_embed = nn.Embedding(num_embeddings=len(vocab), embedding_dim=4)
```

```
'cat'

[0.2,  0.8, -0.5, 0.1]
```

```
'dog'

[0.3,  0.7, -0.6, 0.2]
```

```
'carburetor'

[-0.9,  0.1, 0.8, -0.4]
```

Token Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Token Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Token Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Token Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Which word came first?

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

5

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the     0
               [-0.44,  0.91, -0.12, -0.77],   # cat     1
               [ 0.48,  0.02,  0.05,  0.39],   # chased  2
               [ 0.12, -0.55,  0.33,  0.10],   # the     3
               [-0.30,  0.14, -0.70,  0.81]]]) # dog     4
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the     0
               [-0.44,  0.91, -0.12, -0.77],   # cat     1
               [ 0.48,  0.02,  0.05,  0.39],   # chased  2
               [ 0.12, -0.55,  0.33,  0.10],   # the     3
               [-0.30,  0.14, -0.70,  0.81]]]) # dog     4


positions = torch.arange(seq_len, device=word_vecs.device)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [
               [
               [

positions = torch
```

```
# In practice setup
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")

# Tokenize with max_length=512
inputs = tokenizer(
    "Your text here",
    max_length=512,
    truncation=True,
    padding="max_length",
    return_tensors="pt"
)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [
               [
               [_

positions = torch
```

```
# In practice setup
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
model = AutoModel.from_pretrained("bert-base-uncased")

# Tokenize with max_length=512
inputs = tokenizer(
    "Your text here",
    max_length=512,
    truncation=True,
    padding="max_length",
    return_tensors="pt"
)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
```

```
#positions
tensor([0, 1, 2, 3, 4])
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
positions = positions.unsqueeze(0).expand(batch_size, seq_len)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
positions = positions.unsqueeze(0).expand(batch_size, seq_len)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
positions = positions.unsqueeze(0).expand(batch_size, seq_len)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
positions = positions.unsqueeze(0).expand(batch_size, seq_len)
```

Positional Embeddings

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog


positions = torch.arange(seq_len, device=word_vecs.device)
positions = positions.unsqueeze(0).expand(batch_size, seq_len)

pos_embed = nn.Embedding(num_embeddings=seq_len, embedding_dim=emb_dim)
```

Token and Positional Embeddings

```
pos_vecs = pos_embed(positions)
word_vecs = tok_embed(token_ids)
```

Token and Positional Embeddings

```
pos_vecs = pos_embed(positions)
word_vecs = tok_embed(token_ids)


# add them together
input_vecs = word_vecs + pos_vecs
```

Attention

The cat chased the dog

Pay attention to what?

... and how much attention?

Attention Heads

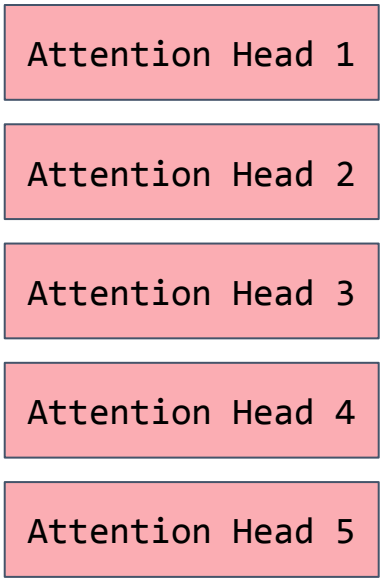

Each head learns different patterns

Attention Heads

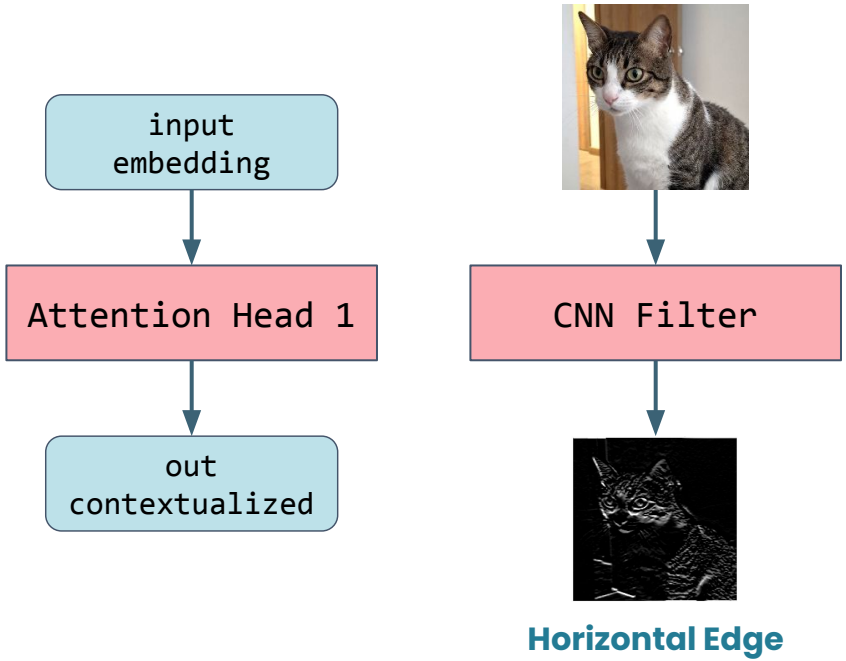

Attention Heads

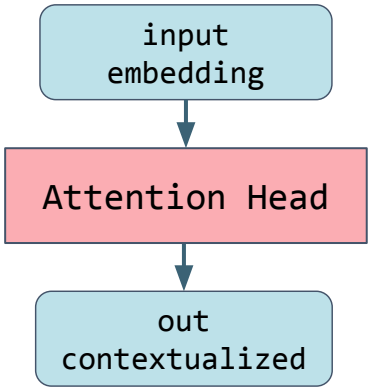

Attention Heads

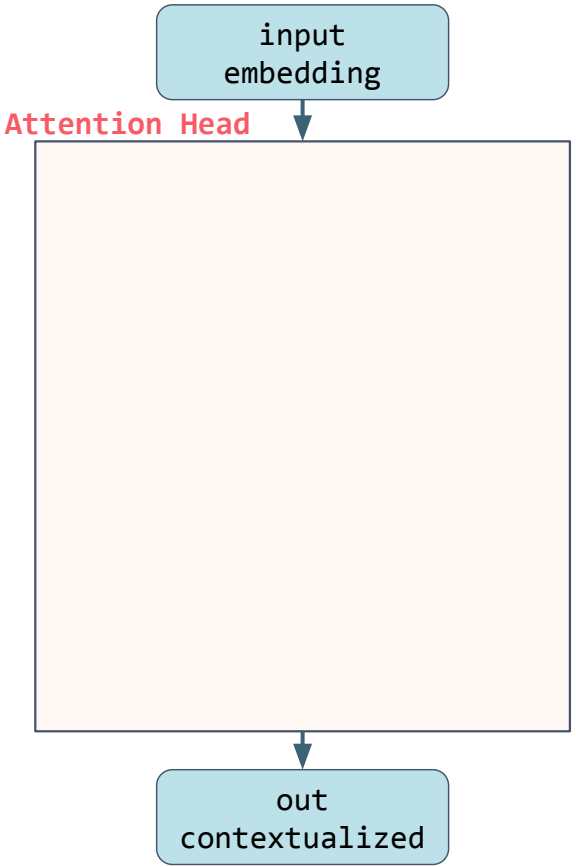

Attention Heads

```
to_k = nn.Linear(4, 4)
```

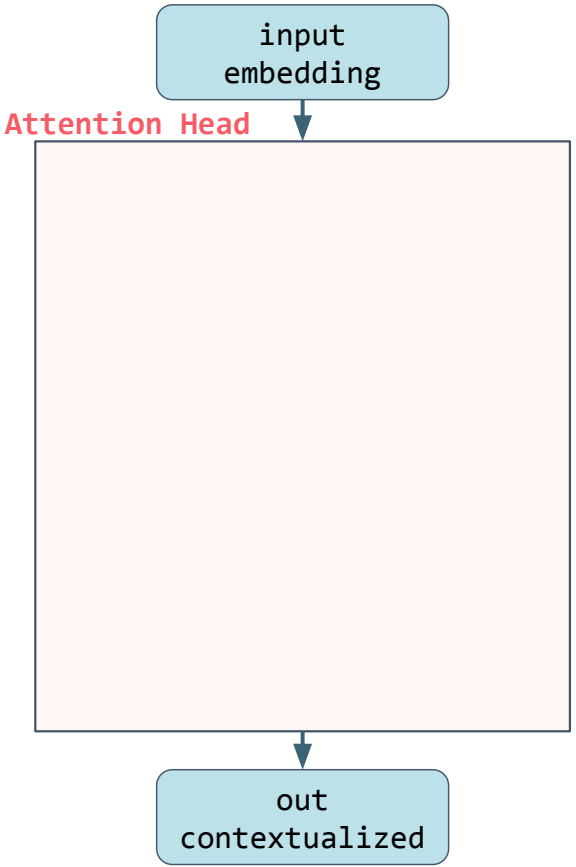

Attention Heads

```
to_k = nn.Linear(4, 4)
```

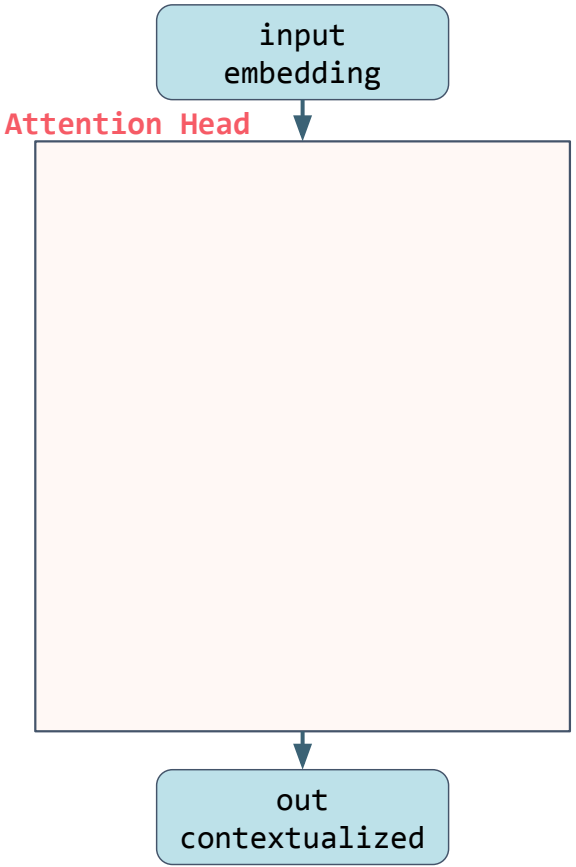

Attention Heads

```
to_k = nn.Linear(4, 4)

K = to_k(input_embedding)
```

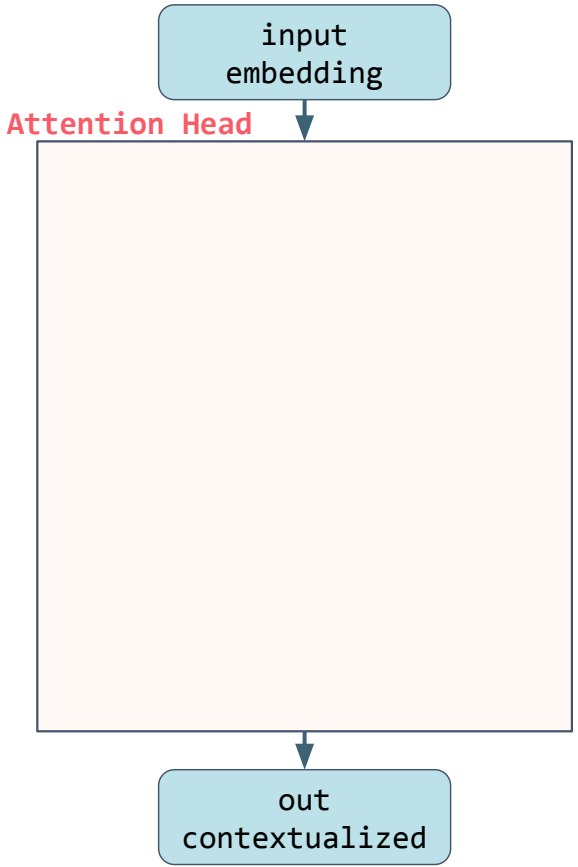

Attention Heads

```
to_k = nn.Linear(4, 4)

K = to_k(input_embedding)
```

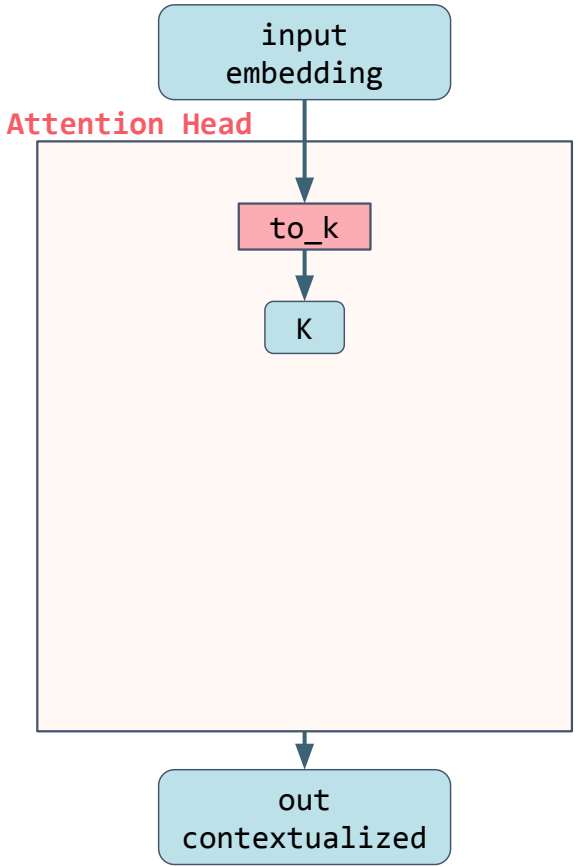

Key
What I look like in this space

But who should I pay attention to?

What am I looking for?

Attention Heads

```
to_k = nn.Linear(4, 4)

K = to_k(input_embedding)
```

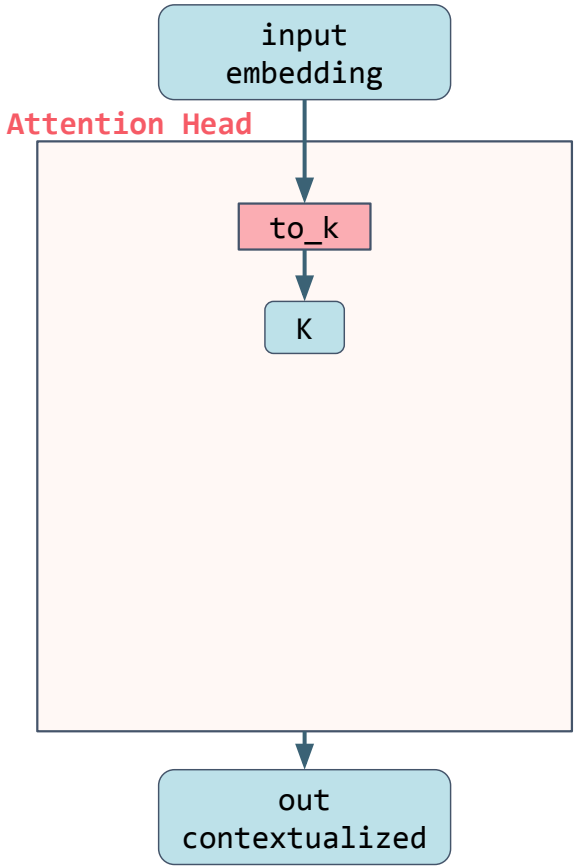

Query
What I'm looking for

Key
What I look like in this space

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)
```

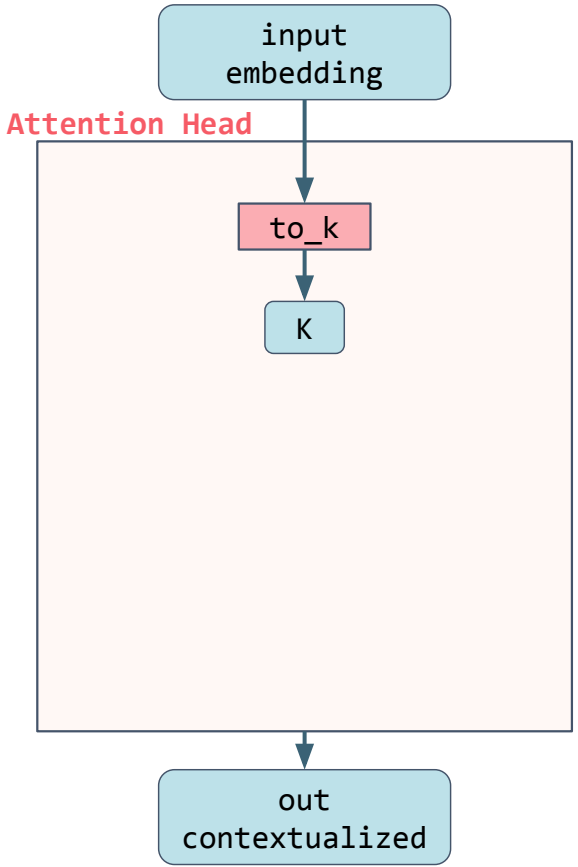

Query
What I'm looking for

Key
What I look like in this space

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)
```

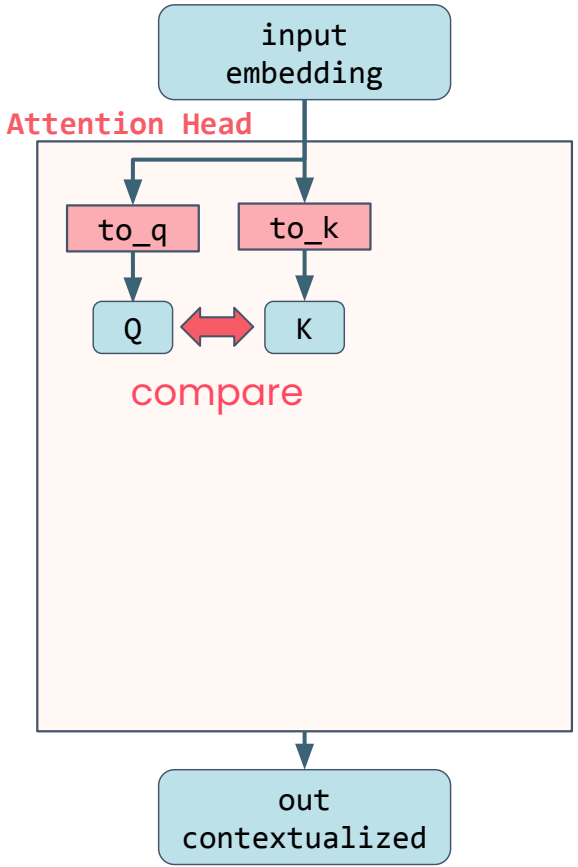

Query
What I'm looking for

Key
What I look like in this space

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)
```

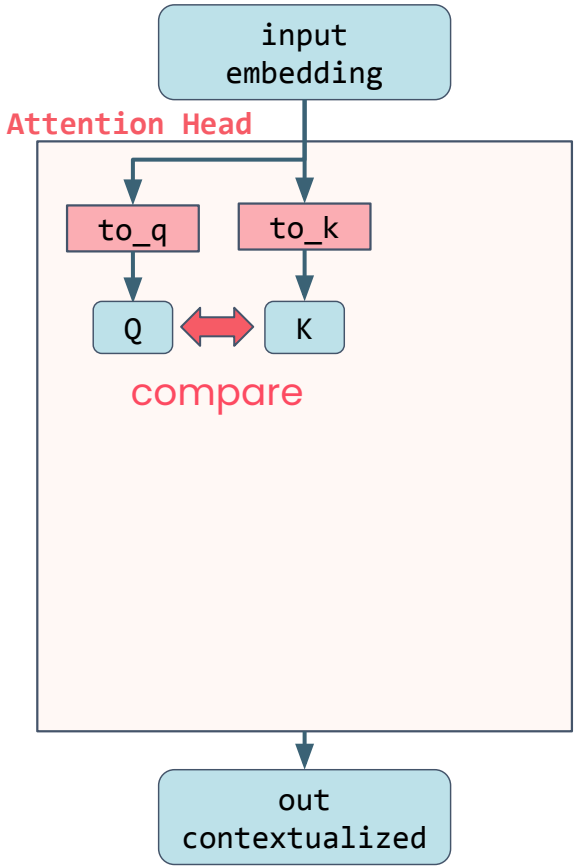

Query
What I'm looking for

Key
What I look like in this space

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)
```

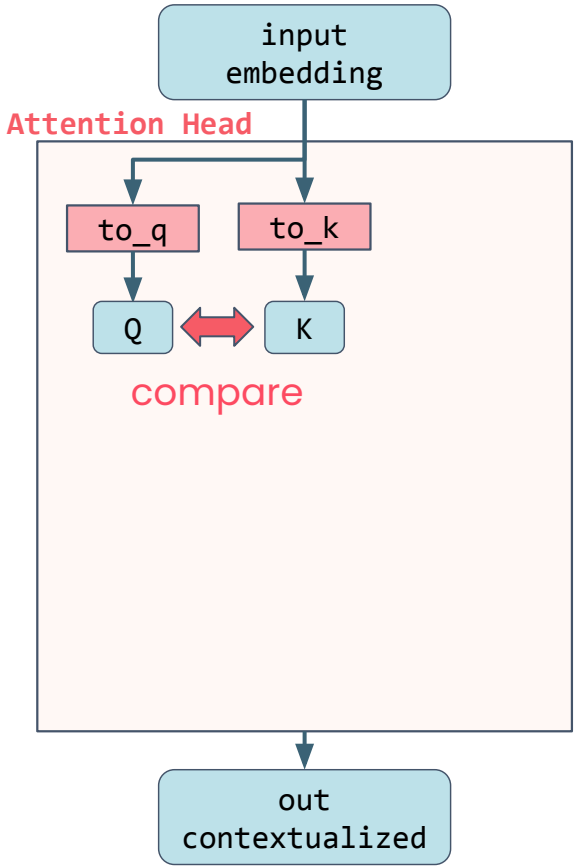

Query
What I'm looking for

Key
What I look like in this space

Calculating Attention Weights

compare

scaled dot product

"the **cat** chased the dog"

"cat" Query = [0.2, -0.1, 0.5, 0.3] what 'cat' is looking for

Calculating Attention Weights

compare

scaled dot product

"**the cat** chased the dog"

"cat" Query = [0.2, -0.1, 0.5, 0.3] what 'cat' is looking for

* * * *

"the" Key = [0.1, 0.5, 0.4, 0.2] what 'the' looks like

0.02 -0.05 0.2 0.06

Calculating Attention Weights

compare

scaled dot product

"**the cat** chased the dog"

"cat" Query = [0.2, -0.1, 0.5, 0.3] what 'cat' is looking for

* * * *

"the" Key = [0.1, 0.5, 0.4, 0.2] what 'the' looks like

```
sum([0.02, -0.05, 0.2, 0.06])
```

Calculating Attention Weights

compare

scaled dot product

"**the cat** chased the dog"

"cat" Query = [0.2, -0.1, 0.5, 0.3] what 'cat' is looking for

* * * *

"the" Key = [0.1, 0.5, 0.4, 0.2] what 'the' looks like

0.23 attention score

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale
```

0.23 attention score

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale
```

$$\frac{0.23}{\text{scale}}$$ attention score

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale
```

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale
```

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

attn = scores.softmax(dim=-1)
```

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Calculating Attention Weights

compare

scaled dot product

```
K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))

# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

attn = scores.softmax(dim=-1)
```

Attention Weights

	the	cat	chased	the	dog	
the	.557	.129	.077	.061	.175	Sum=1
cat	.083	.189	.020	.036	.673	Sum=1
chased	.540	.083	.046	.047	.285	Sum=1
the	.524	.036	.225	.019	.196	Sum=1
dog	.329	.205	.227	.073	.167	Sum=1

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)
```

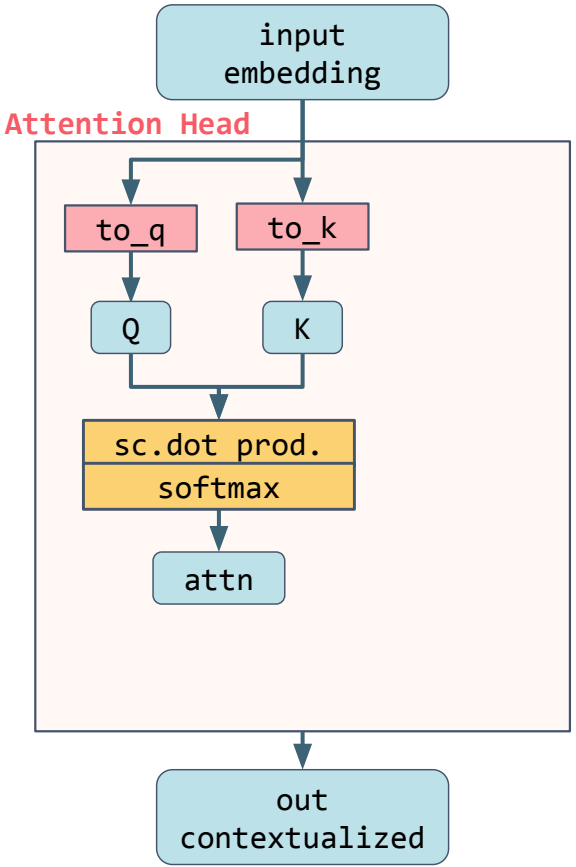

Query
What I'm looking for

Key
What I look like in this space

But what do tokens contribute?

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)
```

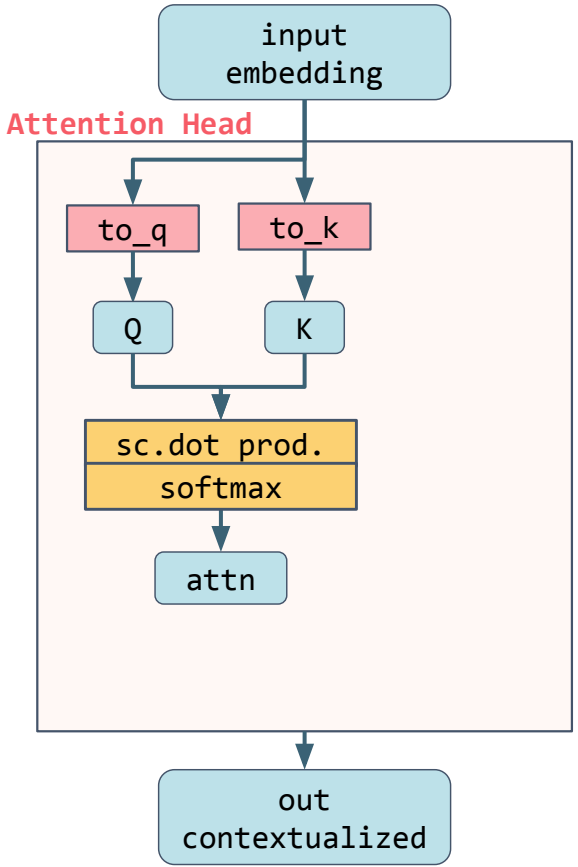

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me,
here's what I'll give you

Attention Heads

```
to_q = nn.Linear(4, 4)
to_k = nn.Linear(4, 4)
to_v = nn.Linear(4, 4)

Q = to_q(input_embedding)
K = to_k(input_embedding)
V = to_v(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)
```

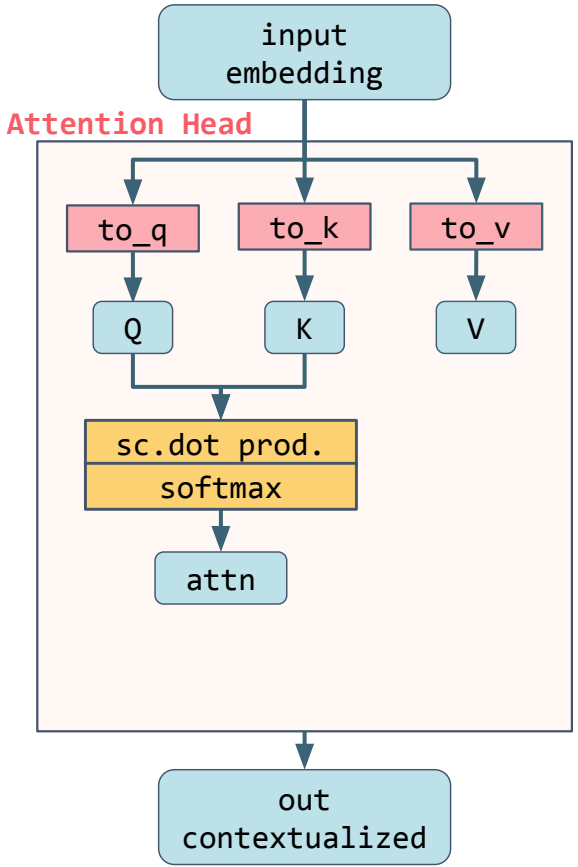

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me,
here's what I'll give you

Attention Output

The **cat** chased the dog

"cat" **attn.** to others

the	.083
cat	.189
chased	.020
the	.036
dog	.673

Value (v) vector of other tokens

"the" **value**

"cat" **value**

"chased" **value**

"the" **value**

"dog" **value**

Attention Output

The **cat** chased the dog

"Cat" **attn.** * **Value** vector of other tokens

.083 * "the" **value**

.189 * "cat" **value**

.020 * "chased" **value**

.036 * "the" **value**

.673 * "dog" **value**

SUM → [0.27, -0.35, 0.80, -0.58]

"cat" contextualized output

```
# contextualized output: dot product
out = torch.matmul(attn, V)
```

Attention Heads

```
to_k = nn.Linear(4, 4)
to_q = nn.Linear(4, 4)
to_v = nn.Linear(4, 4)

K = to_k(input_embedding)
Q = to_q(input_embedding)
V = to_v(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)

# contextualized output: dot product
out = torch.matmul(attn, V)
return out, attn
```

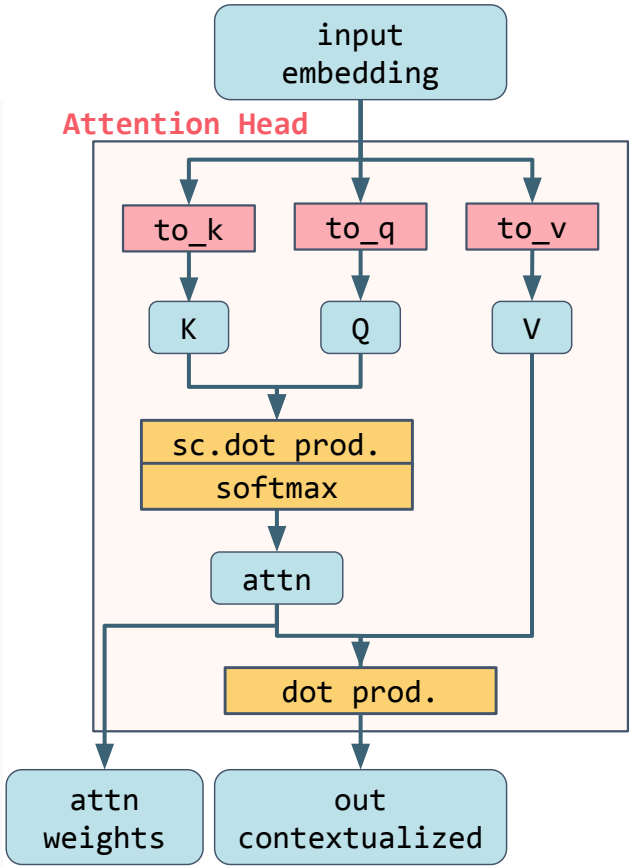

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me, here's what I'll give you

Attention Heads

```
to_k = nn.Linear(4, 4)
to_q = nn.Linear(4, 4)
to_v = nn.Linear(4, 4)

K = to_k(input_embedding)
Q = to_q(input_embedding)
V = to_v(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)

# contextualized output: dot product
out = torch.matmul(attn, V)
return out, attn
```

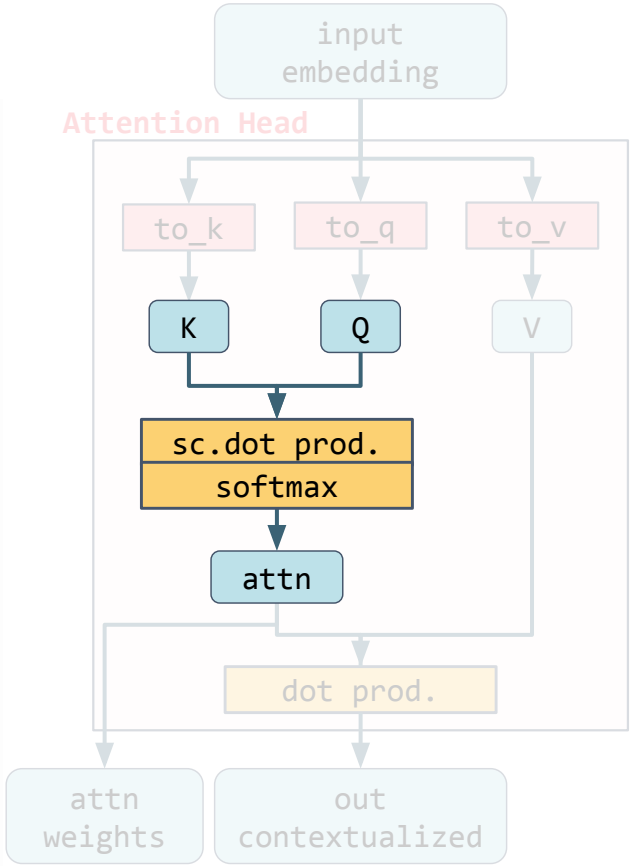

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me, here's what I'll give you

Attention Heads

```
to_k = nn.Linear(4, 4)
to_q = nn.Linear(4, 4)
to_v = nn.Linear(4, 4)

K = to_k(input_embedding)
Q = to_q(input_embedding)
V = to_v(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)

# contextualized output: dot product
out = torch.matmul(attn, V)
return out, attn
```

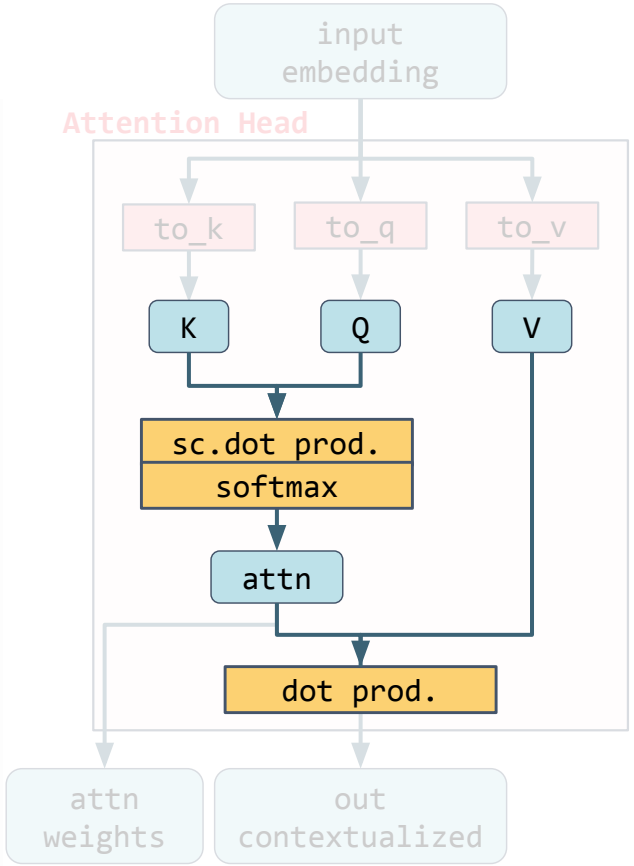

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me, here's what I'll give you

Attention Heads

```
to_k = nn.Linear(4, 4)
to_q = nn.Linear(4, 4)
to_v = nn.Linear(4, 4)

K = to_k(input_embedding)
Q = to_q(input_embedding)
V = to_v(input_embedding)

K_t = K.transpose(-2, 1)
scale = math.sqrt(Q.size(-1))
# KQ compare: scaled dot product
scores = torch.matmul(Q, K_t)/scale

# attn weights: softmax on scaled dp
attn = F.softmax(scores, dim=-1)

# contextualized output: dot product
out = torch.matmul(attn, V)
return out, attn
```

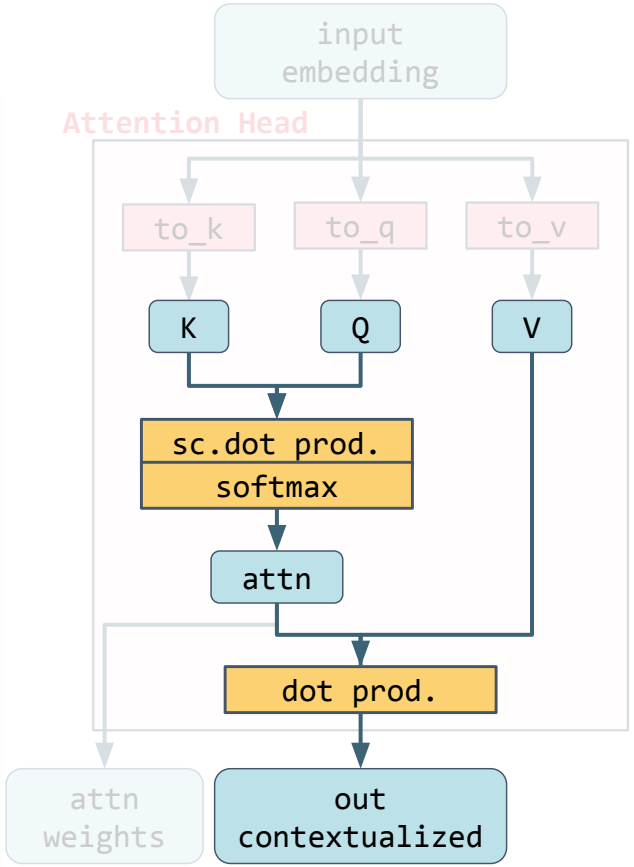

Query
What I'm looking for

Key
What I look like in this space

Value
If you attend to me, here's what I'll give you

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)


torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)
```

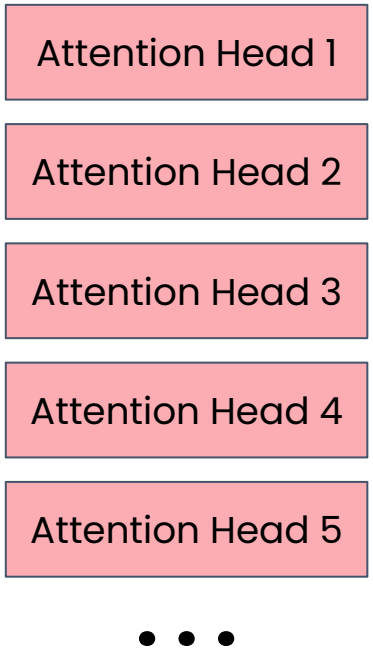

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)
```

(batch_size, seq_len, embed_dim)

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)


                                Q  K  V
attn_output, attn_weights = mha(x, x, x) self attention
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)


                                Q  K  V
attn_output, attn_weights = mha(x, x, x)  cross attention
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)


                                Q  K  V
attn_output, attn_weights = mha(x, x, x) cross attention
```

nn.MultiheadAttention

```
import torch
import torch.nn as nn

mha = nn.MultiheadAttention(embed_dim=4,
                            num_heads=2,
                            dropout=0.1,
                            batch_first=True)


                                Q  K  V
attn_output, attn_weights = mha(x, y, y)
```

cross attention

DeepLearning.AI

Encoders

Specialized Approaches to
Natural Language Processing in Pytorch

Attention

Attention

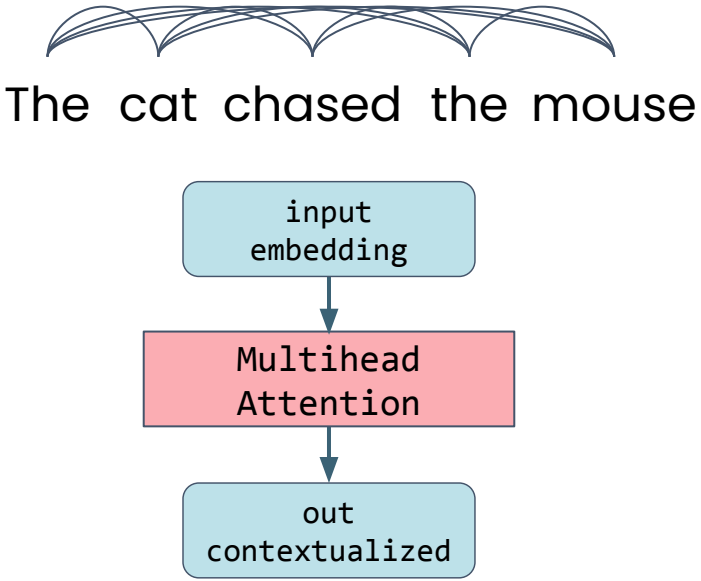
The cat chased the mouse
input embedding
Multihead Attention
out contextualized

Encoder Block

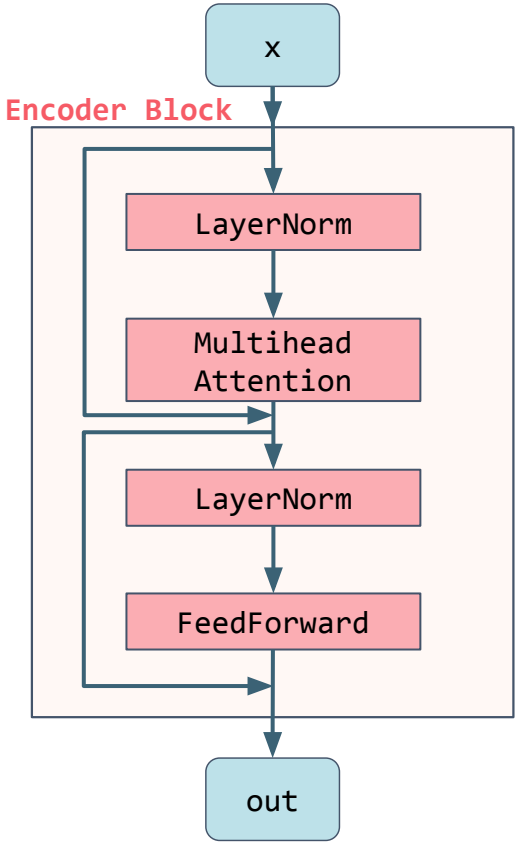

Encoder Block

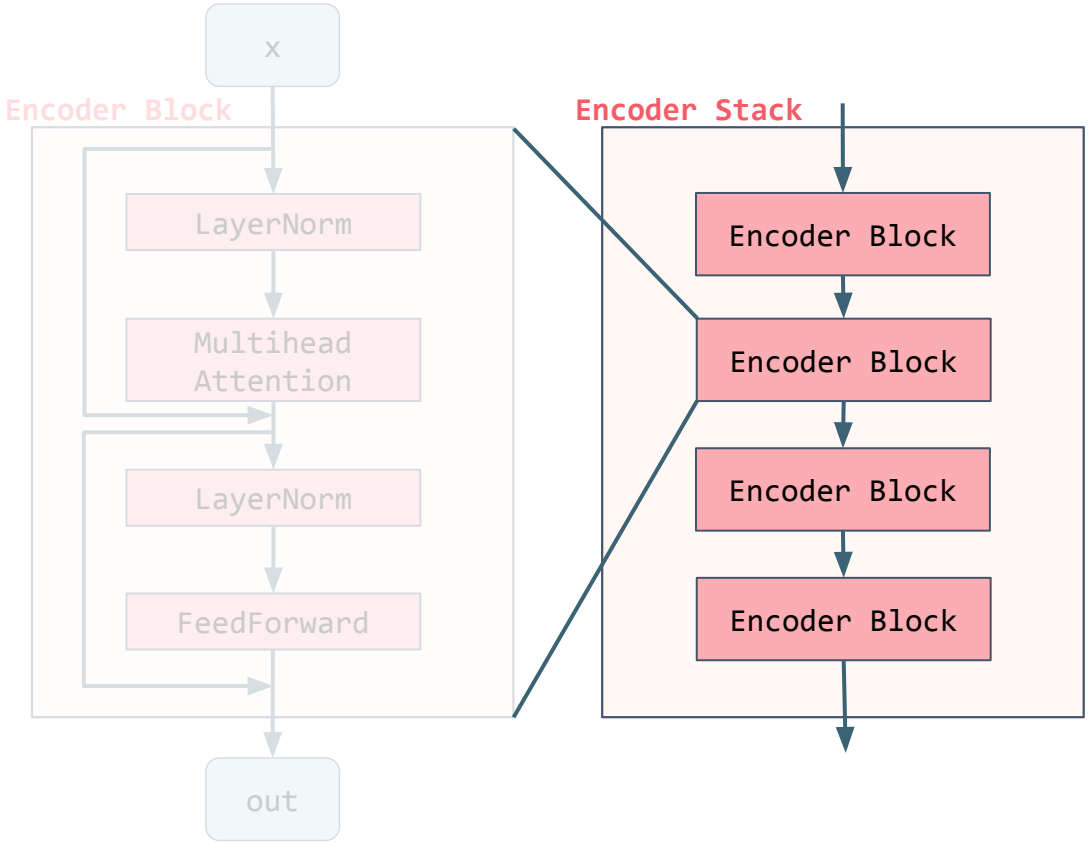

Encoder Block

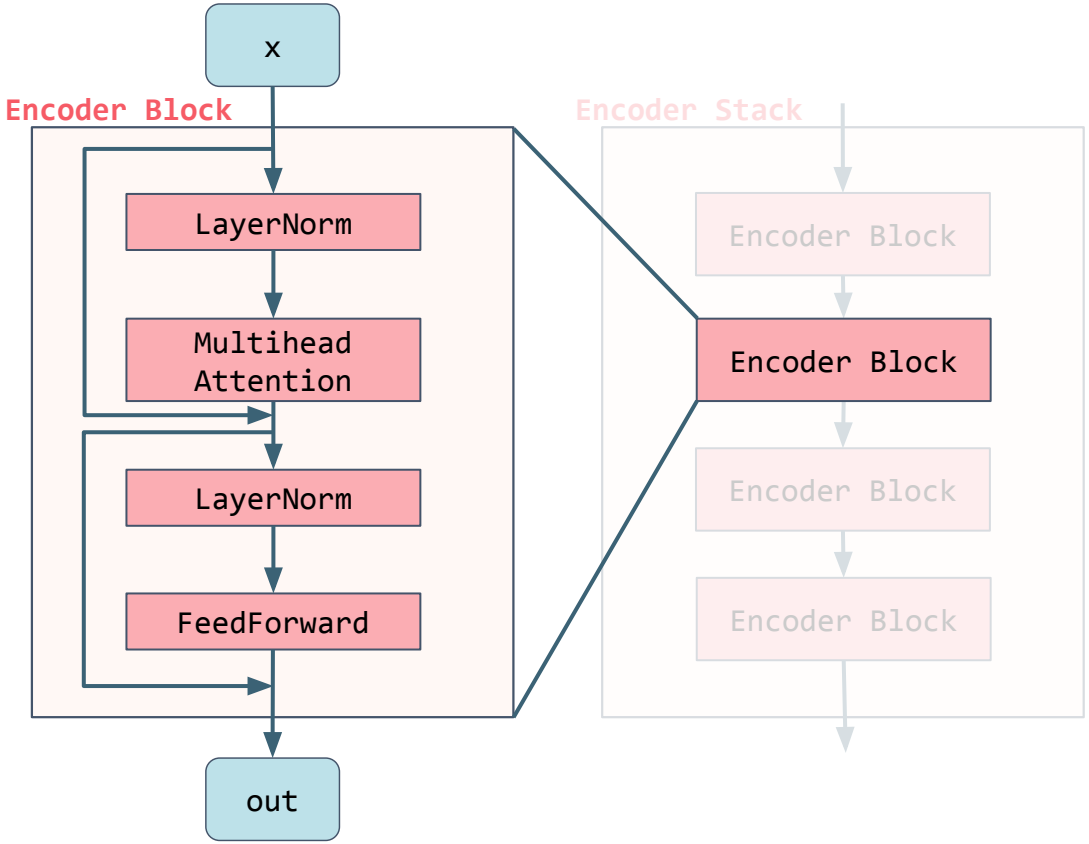

Encoder Block

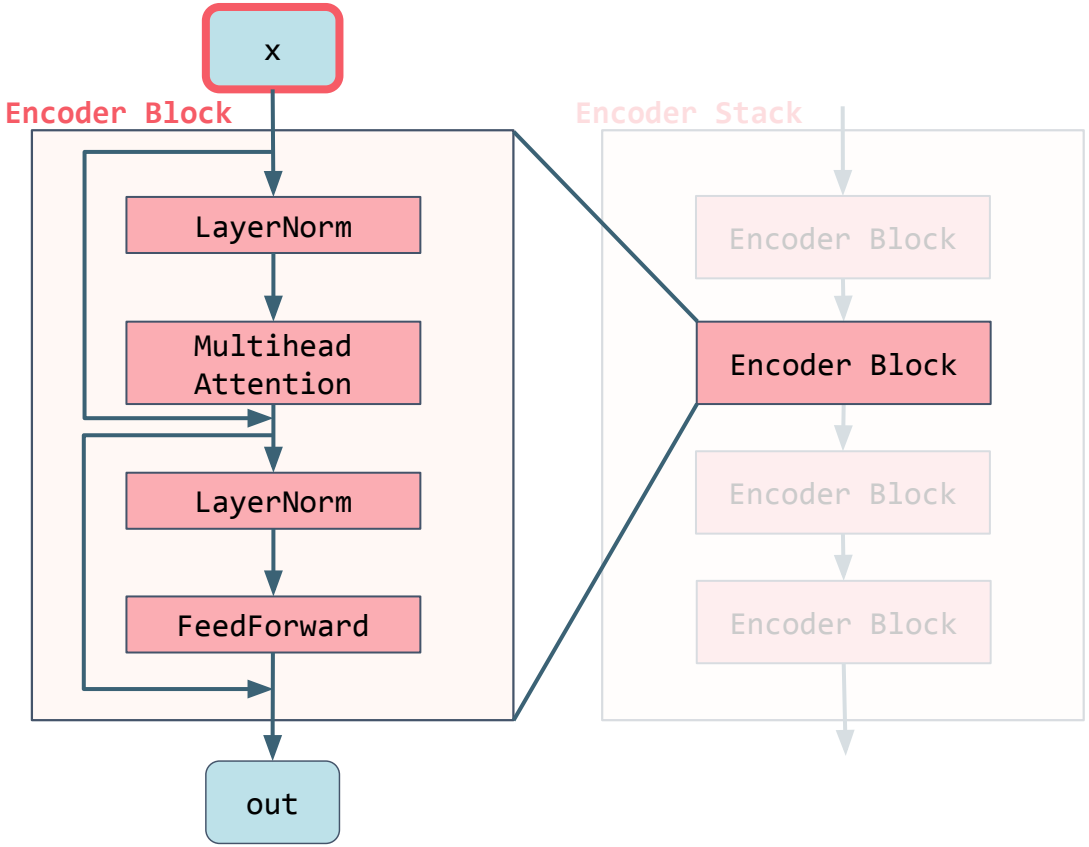

Encoder Block

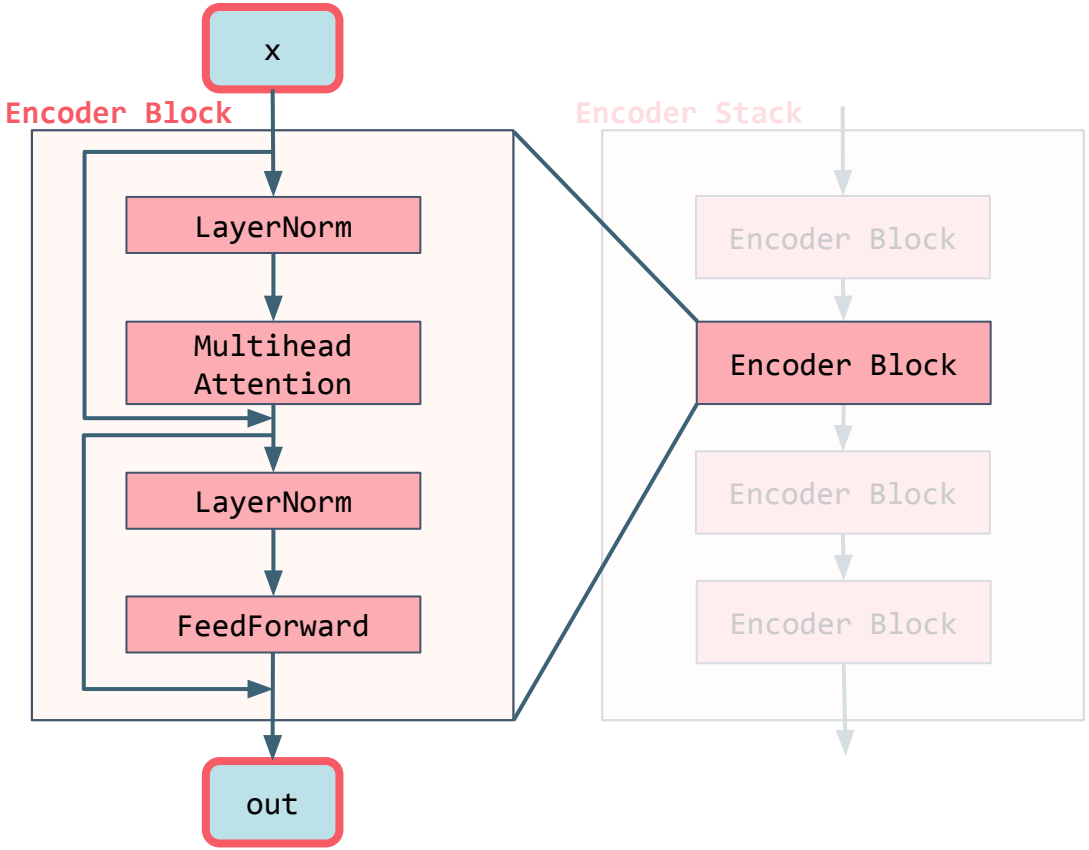

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)
```

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)


layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)


layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

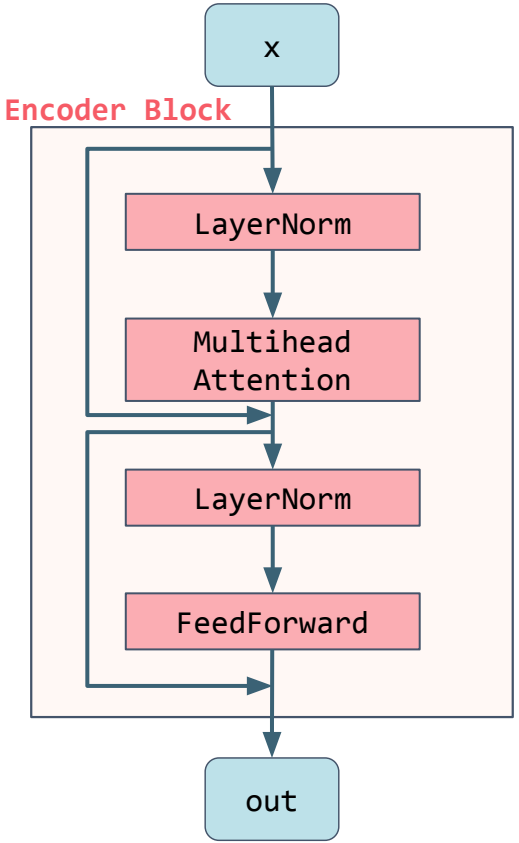

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)


layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

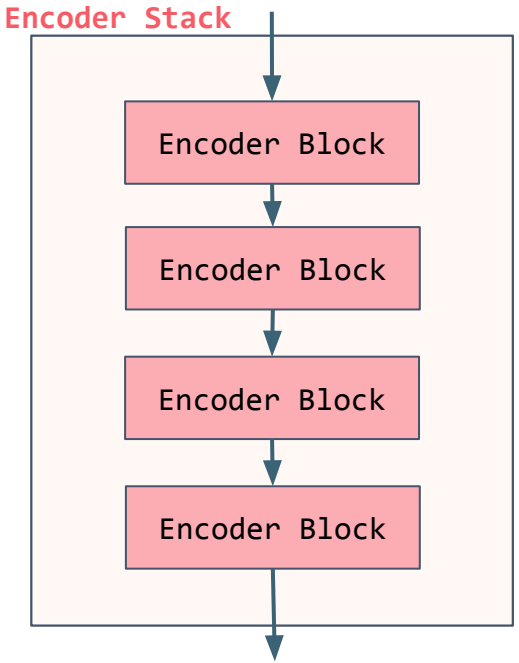

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)


layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

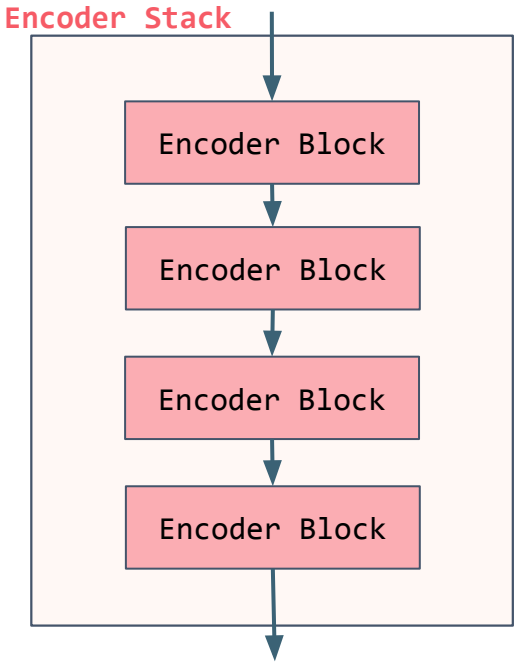

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .

layer = nn.TransformerEncoderLaye

stack = nn.TransformerEncoder(lay
```

Attention Is All You Need

Ashish Vaswani* Google Brain avaswani@google.com

Noam Shazeer* Google Brain noam@google.com

Niki Parmar* Google Research nikip@google.com

Jakob Uszkoreit* Google Research usz@google.com

Llion Jones* Google Research llion@google.com

Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu

Łukasz Kaiser* Google Brain lukaszkaiser@google.com

Illia Polosukhin* ‡ illia.polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

nn.TransformerEncoder

```
import torch.nn as nn

mha = nn.MultiheadAttention(. . .)


layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()




    def forward(self, x):
```

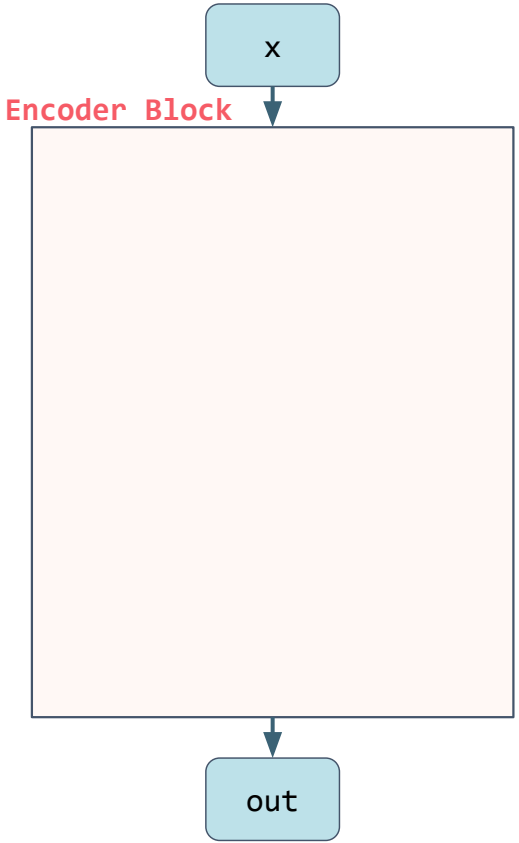

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()




    def forward(self, x):
```

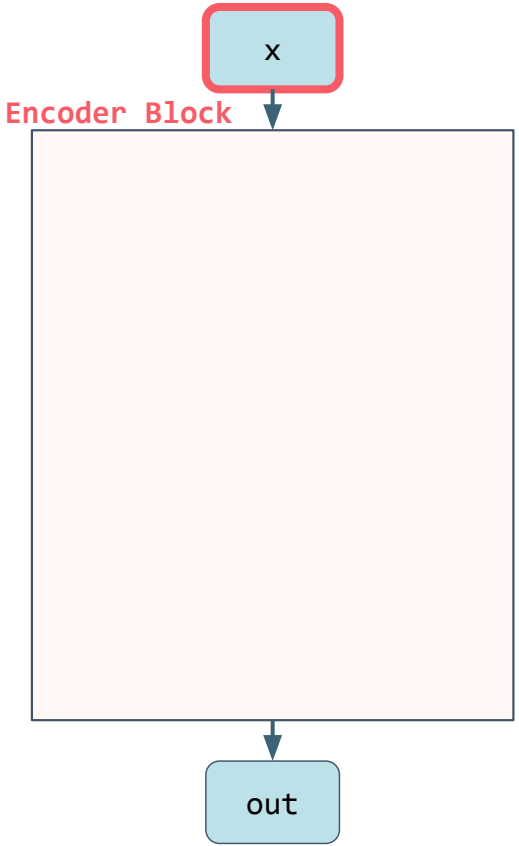

Input

x

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Input

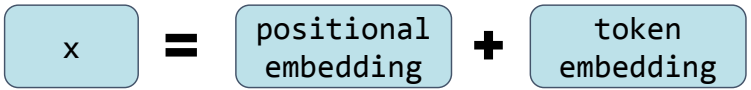


```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()



    def forward(self, x):
```

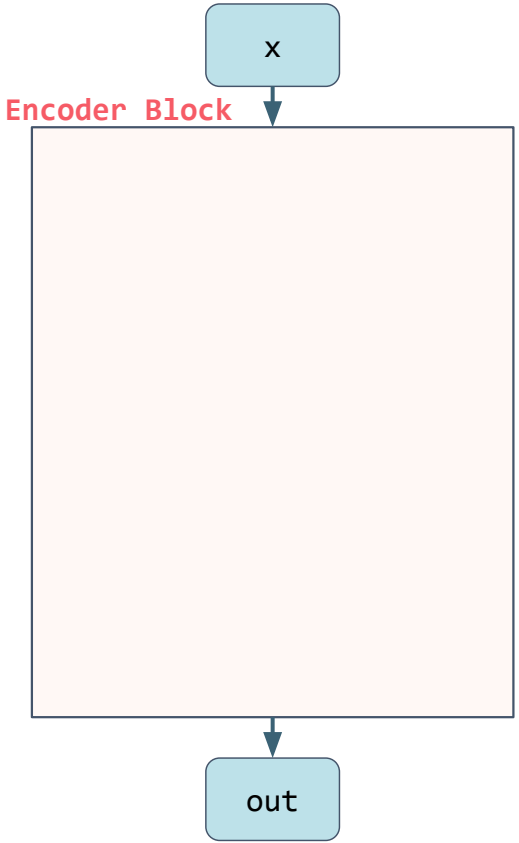

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)



    def forward(self, x):
        x_norm = self.ln1(x)
```

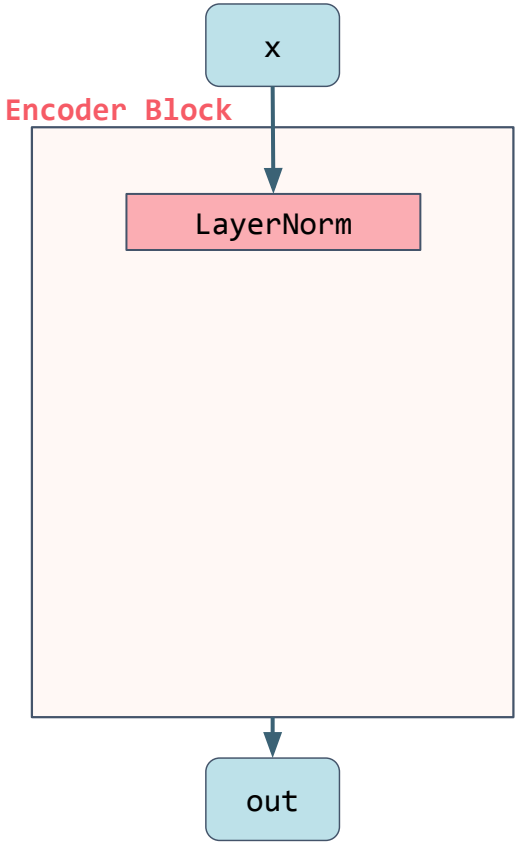

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.53,  1.61, -0.02, -1.05],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.30,  0.14, -0.70,  0.81]]]) # dog
```

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.53,  1.61, -0.02, -1.05],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.53,  1.61, -0.02, -1.05],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

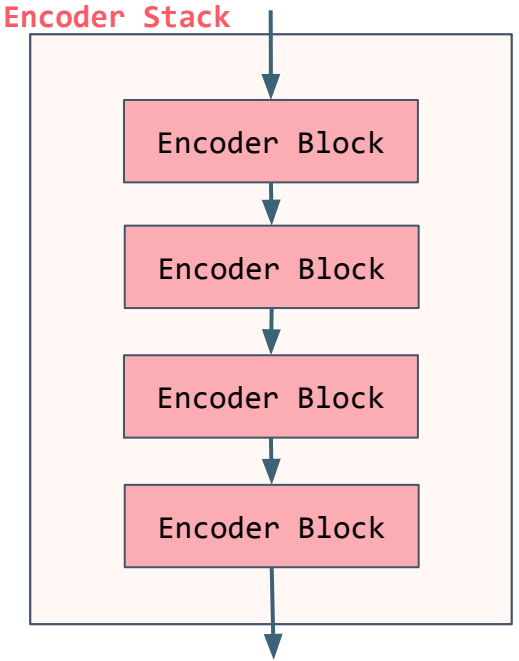

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.53,  1.61, -0.02, -1.05],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

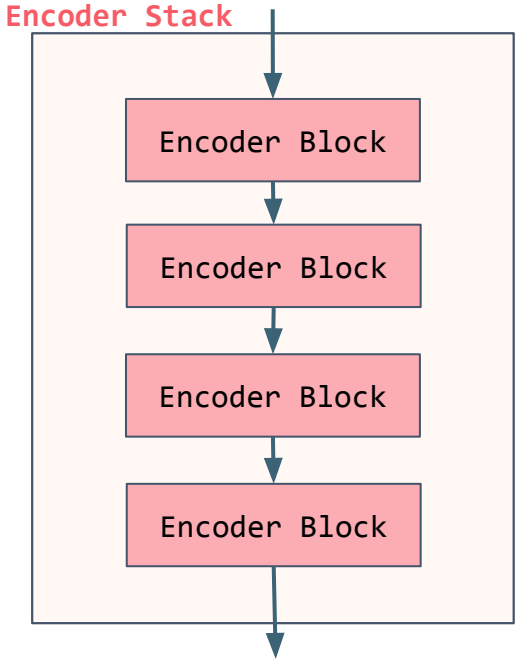

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [ 12.4,  -57.8,  30.2,  8.9],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

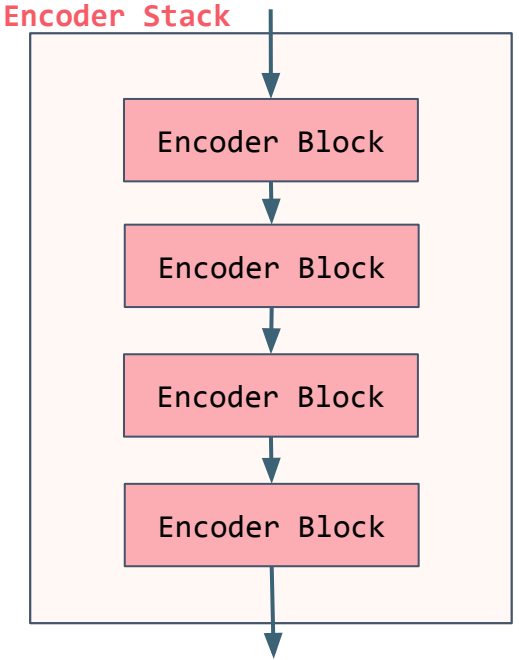

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [ .001, -.004,  .002,  0.009],  # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

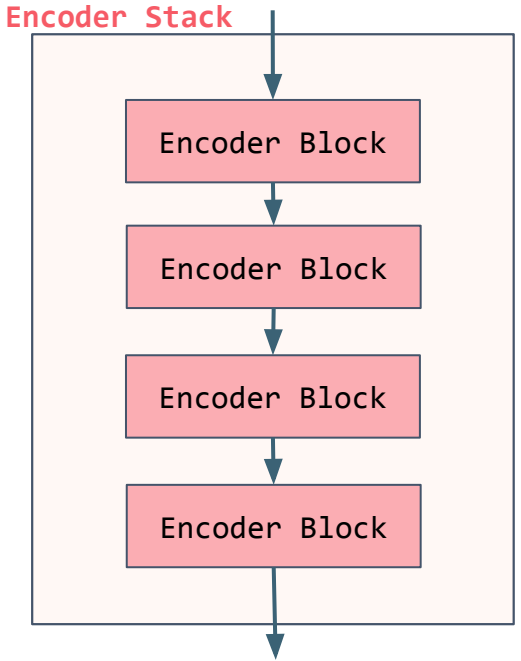

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

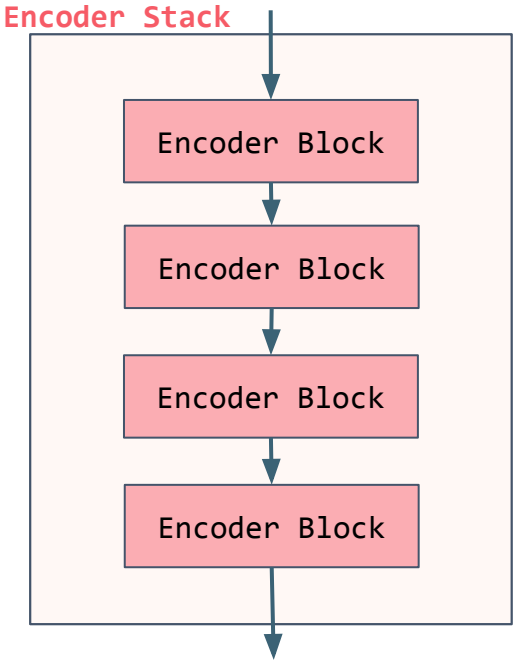

LayerNorm

LayerNorm

Tokens normalized independently

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

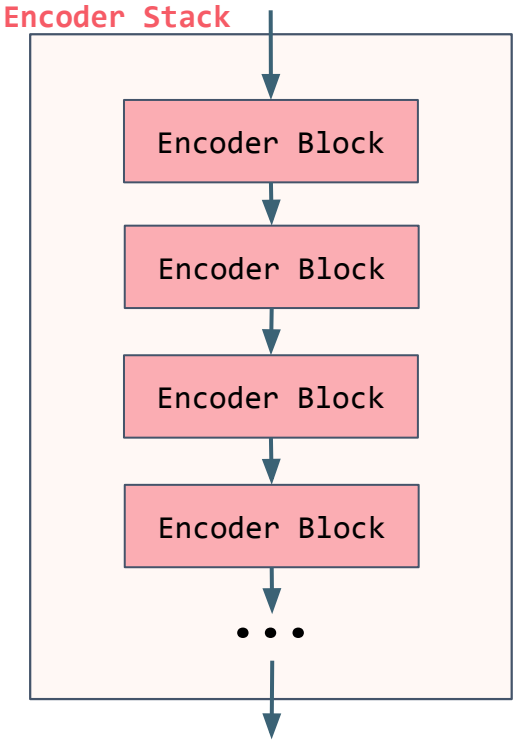

LayerNorm

LayerNorm

mean: 0 **variance:** 1

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

LayerNorm

LayerNorm **mean:** 0 **variance:** 1

```
torch.tensor([[[ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.44,  0.91, -0.12, -0.77],   # cat
               [ 0.48,  0.02,  0.05,  0.39],   # chased
               [ 0.12, -0.55,  0.33,  0.10],   # the
               [-0.51,  0.27, -1.23,  1.47]]]) # dog
```

scale γ_1 γ_2 γ_3 γ_4

shift β_1 β_2 β_3 β_4

$\gamma_1 \, x_{t,1} + \beta_1$

Encoder Block

```python
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)



    def forward(self, x):
        x_norm = self.ln1(x)
```

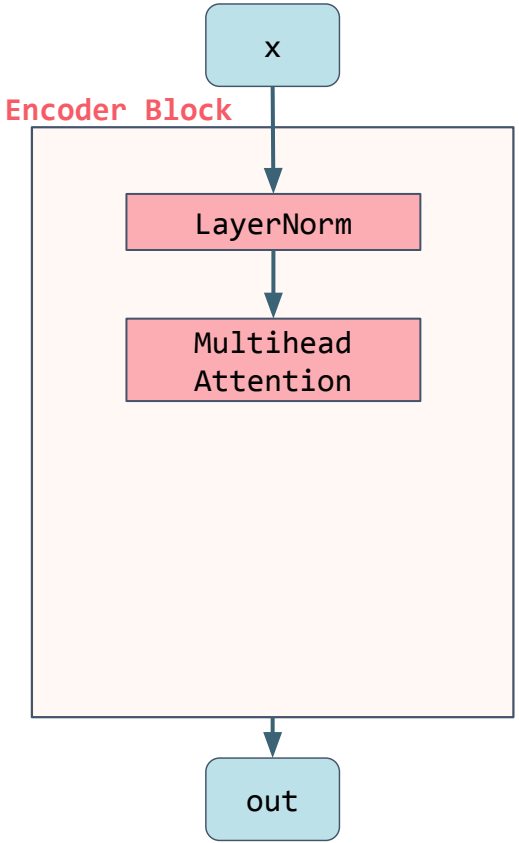

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
```

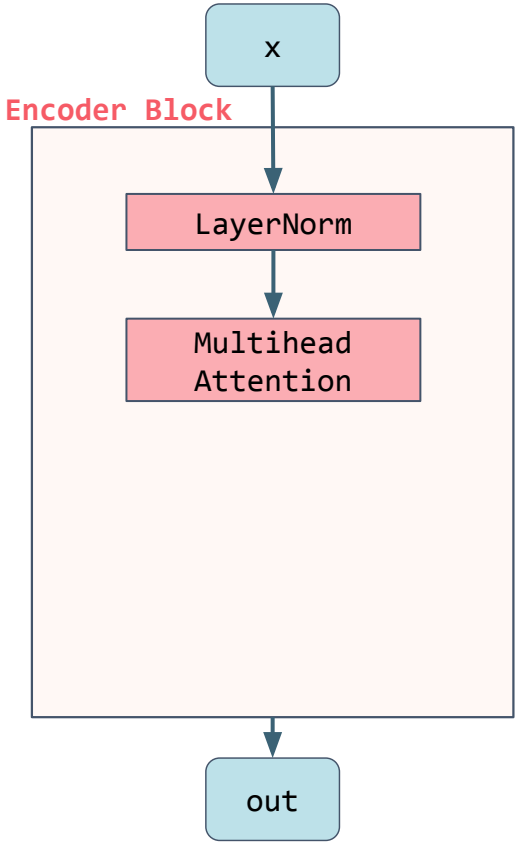

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
```

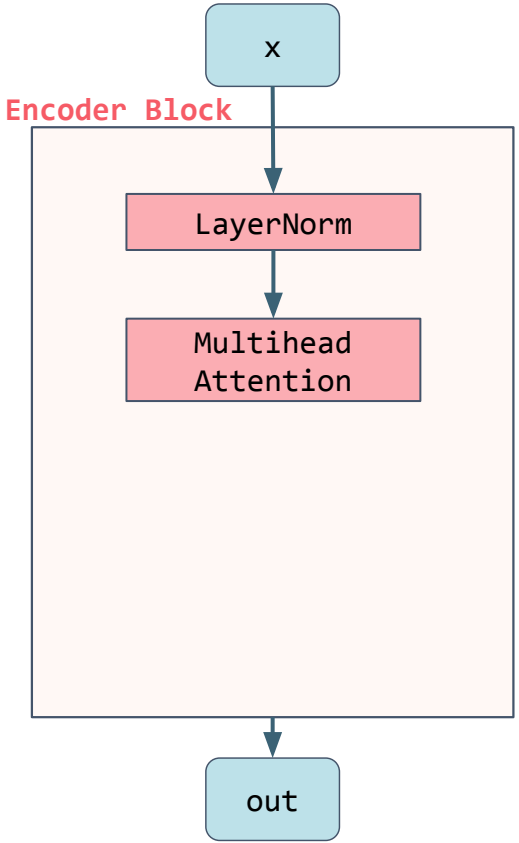

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
```

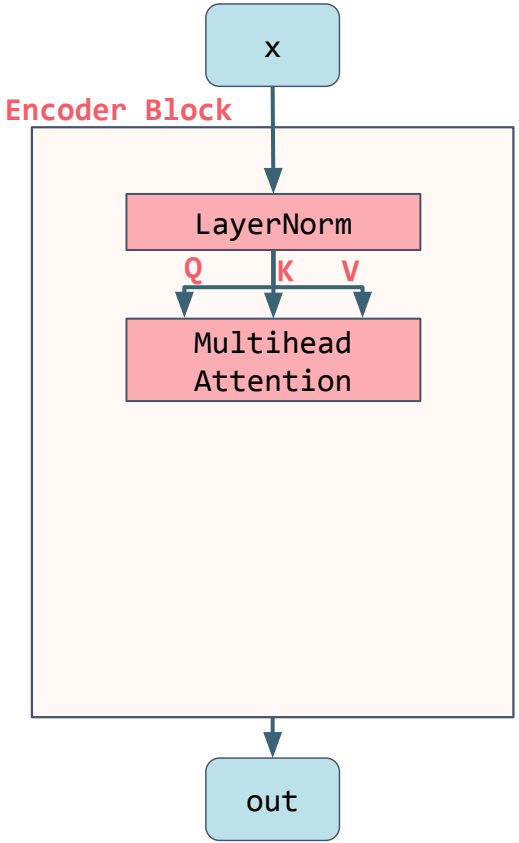

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
```

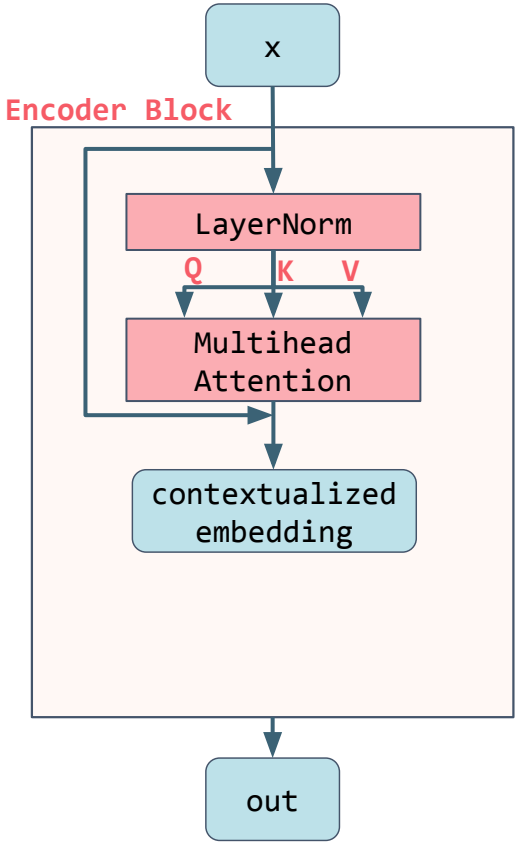

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
```

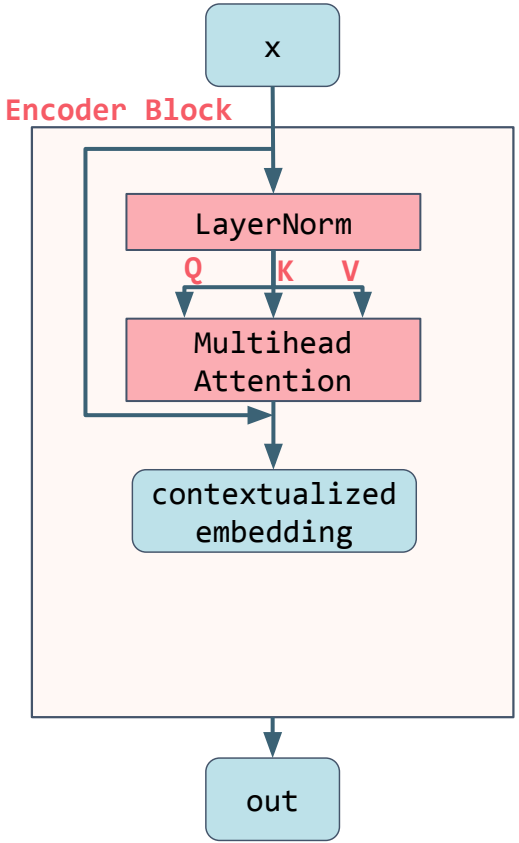

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
```

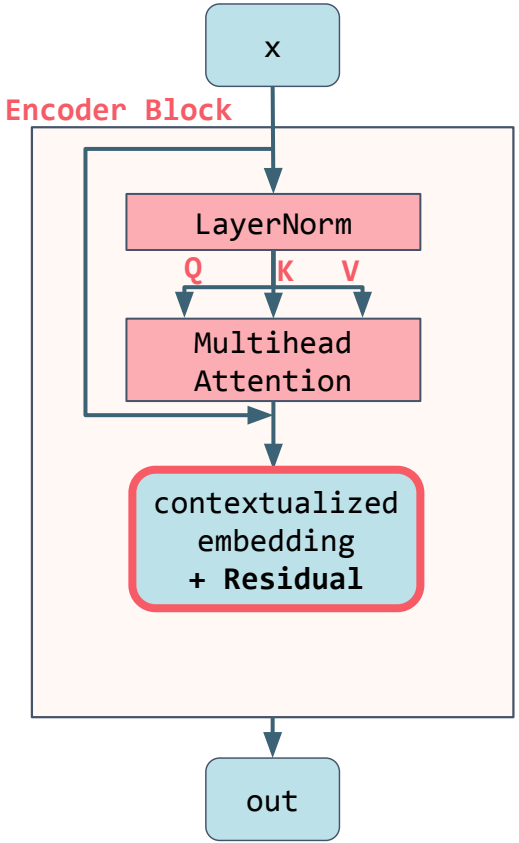

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
```

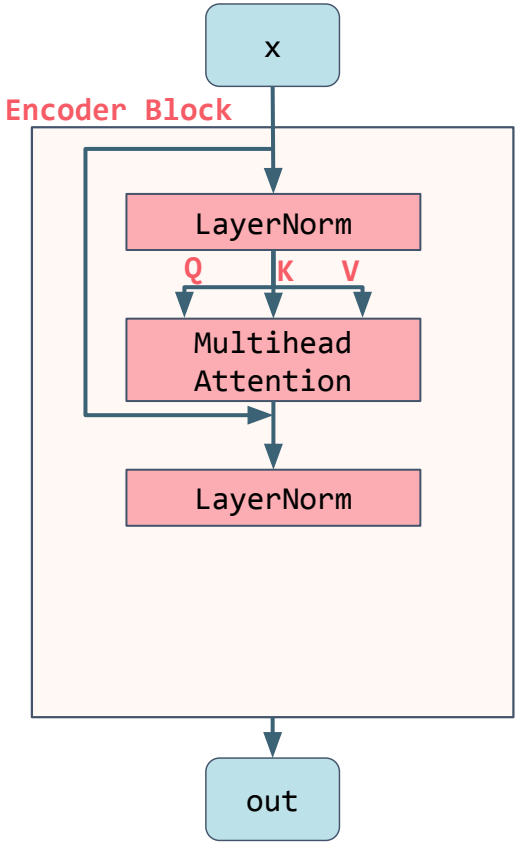

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)


    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
```

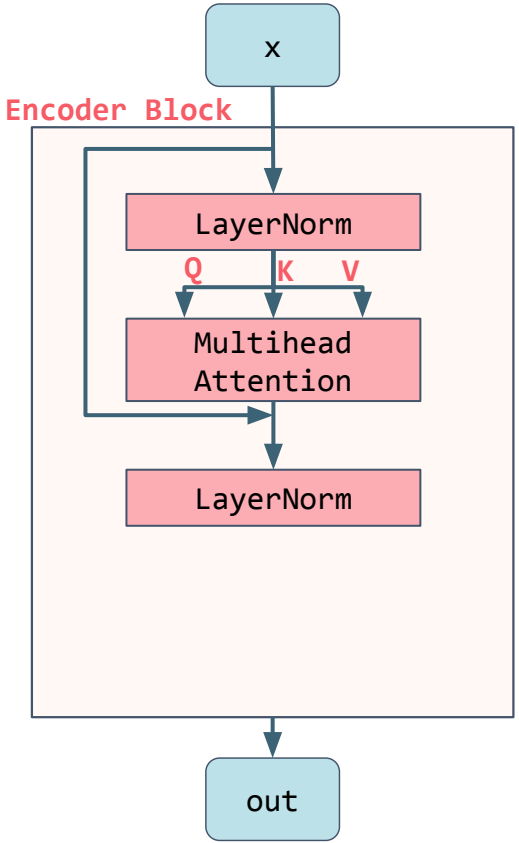

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
```

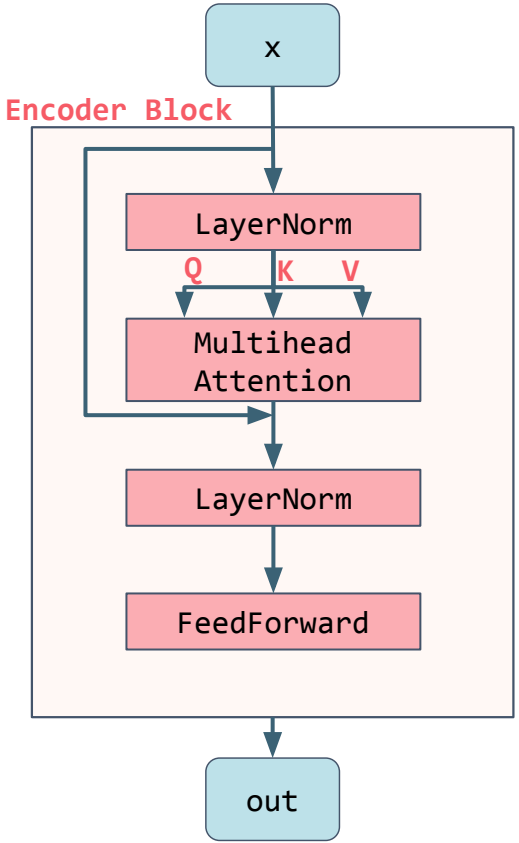

Encoder Block

```
nn.Sequential(
            nn.Linear(d_model, hidden),
            nn.ReLU(),
            nn.Linear(hidden, d_model))
```

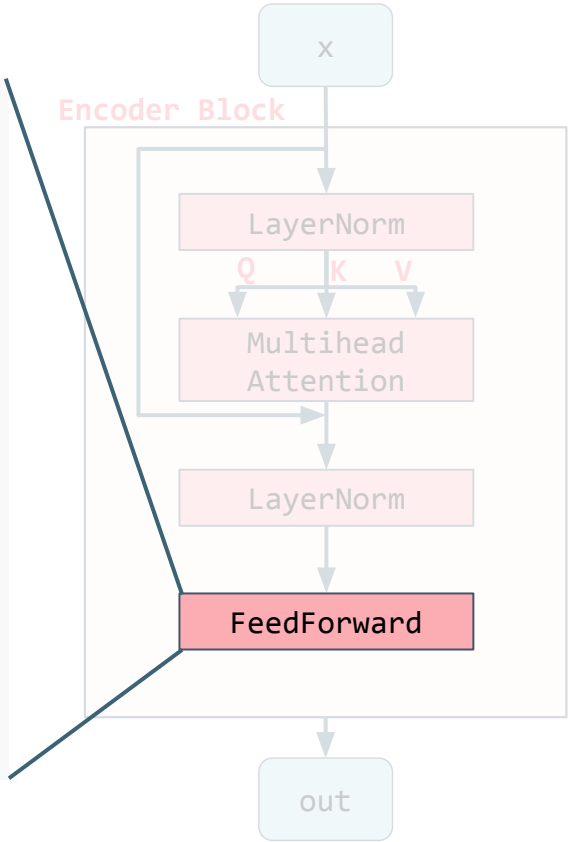

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
```

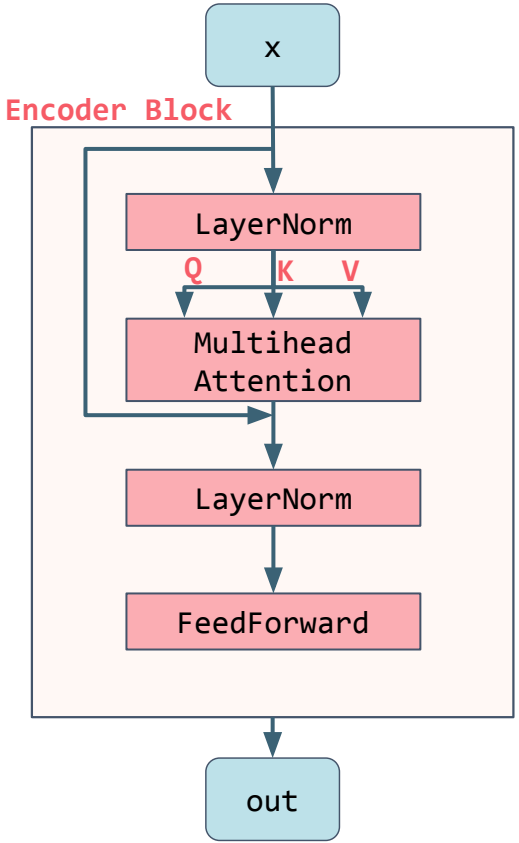

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

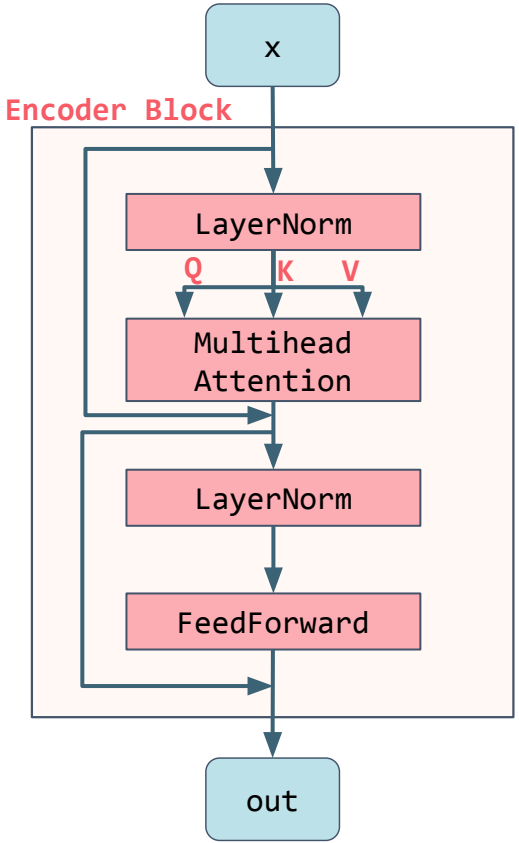

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

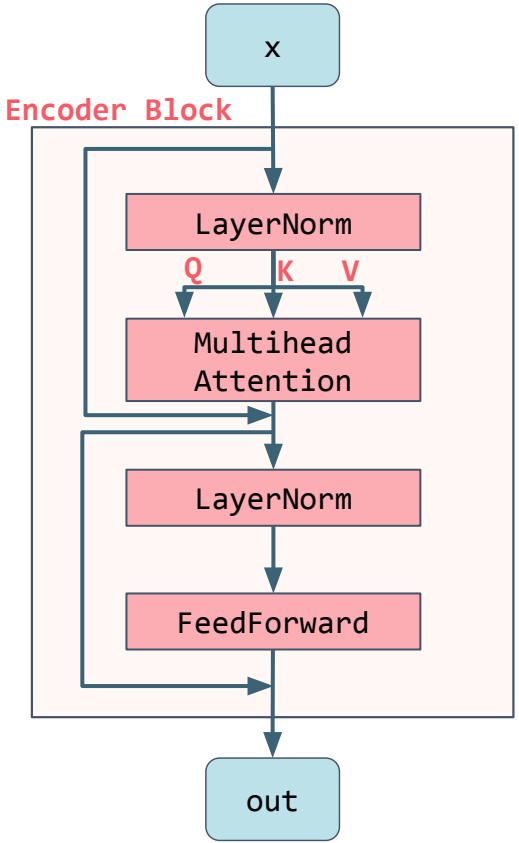

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

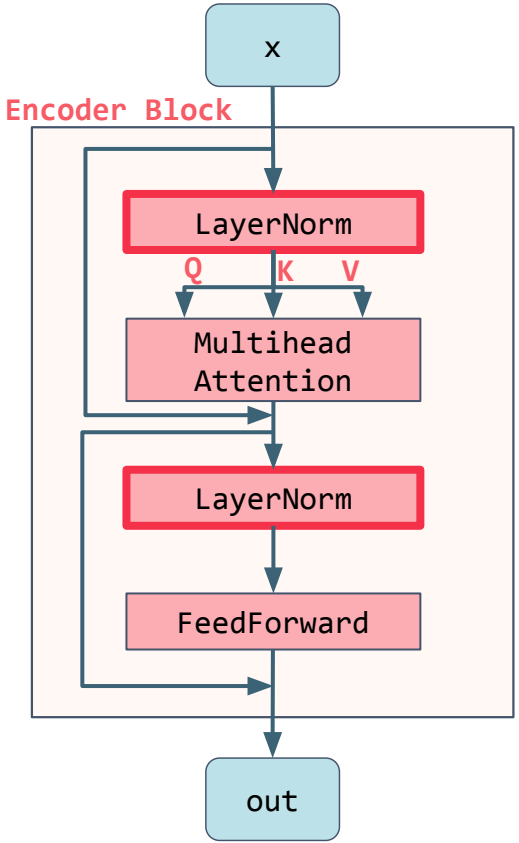

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

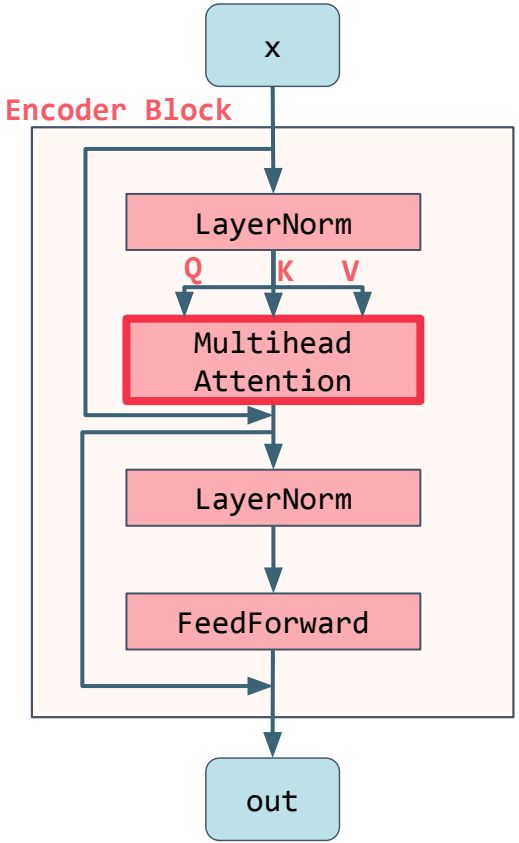

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

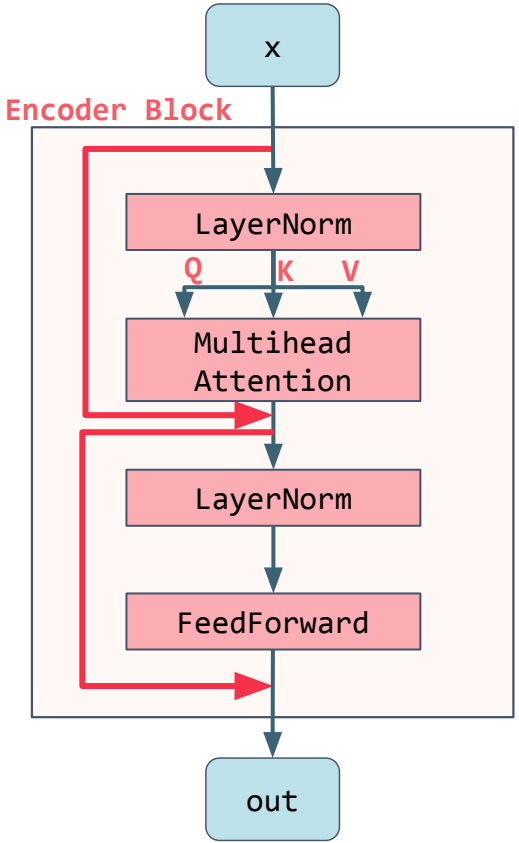

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

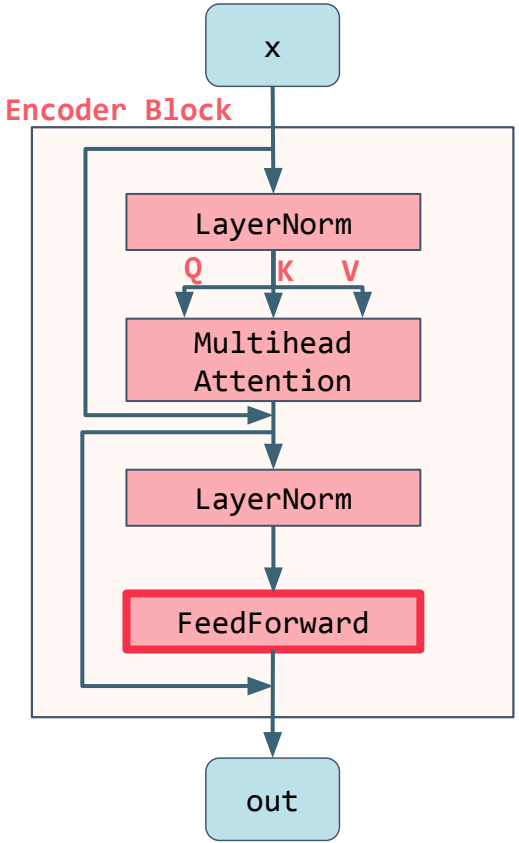

Encoder Block

```
class EncoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

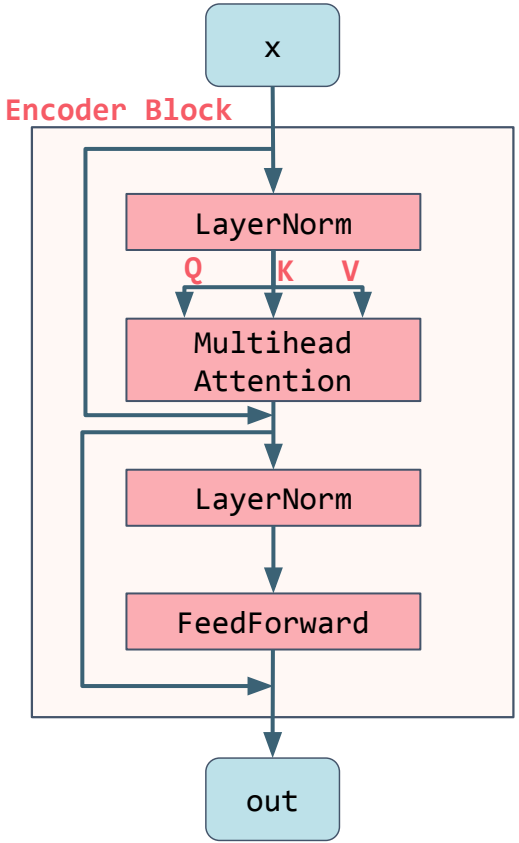

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

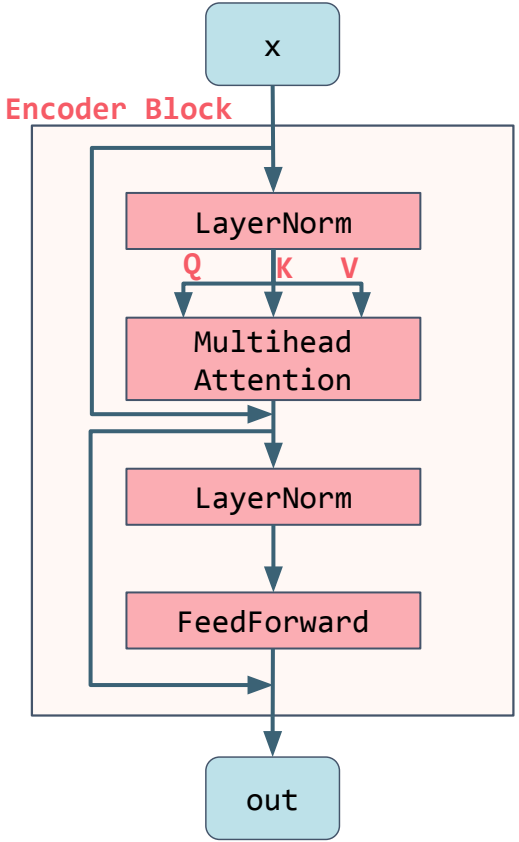

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

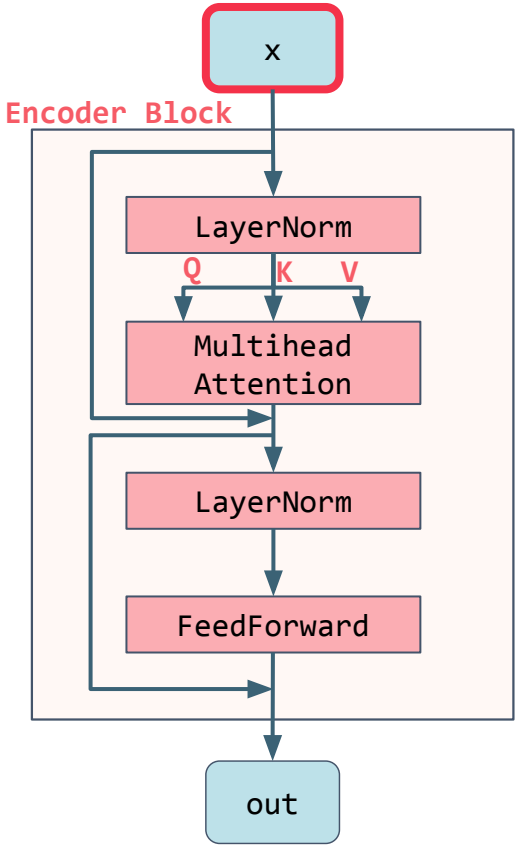

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

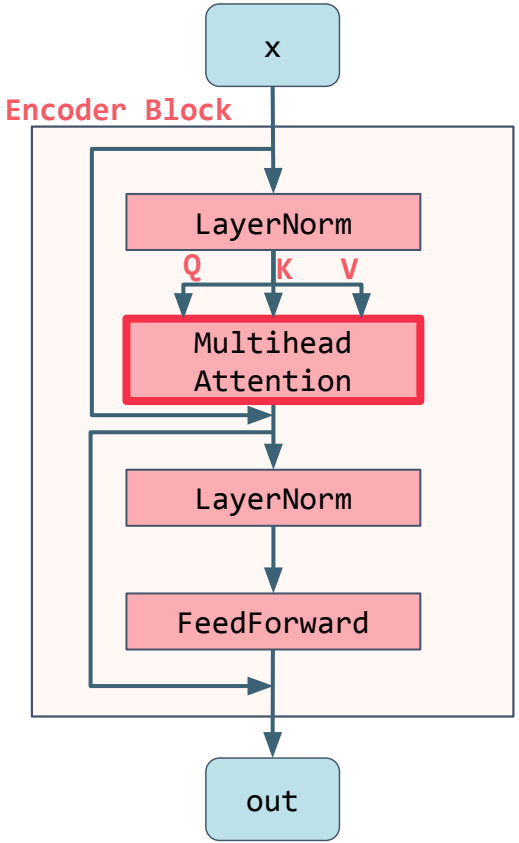

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

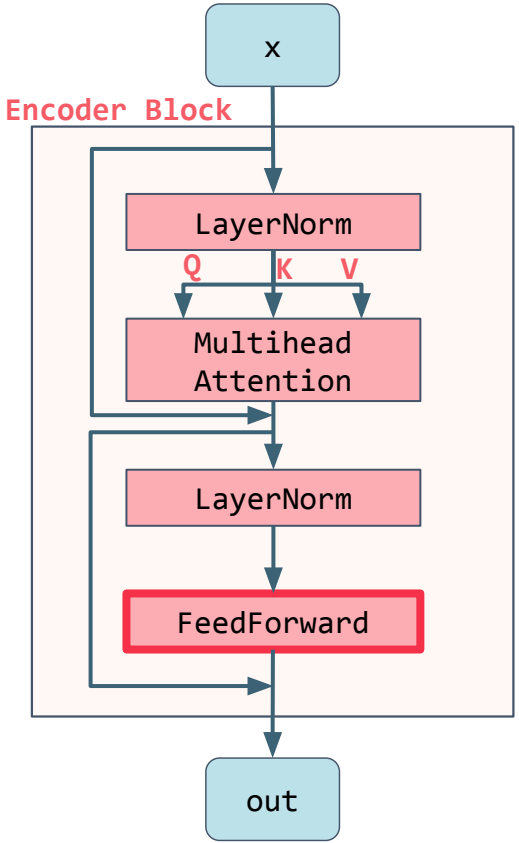

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

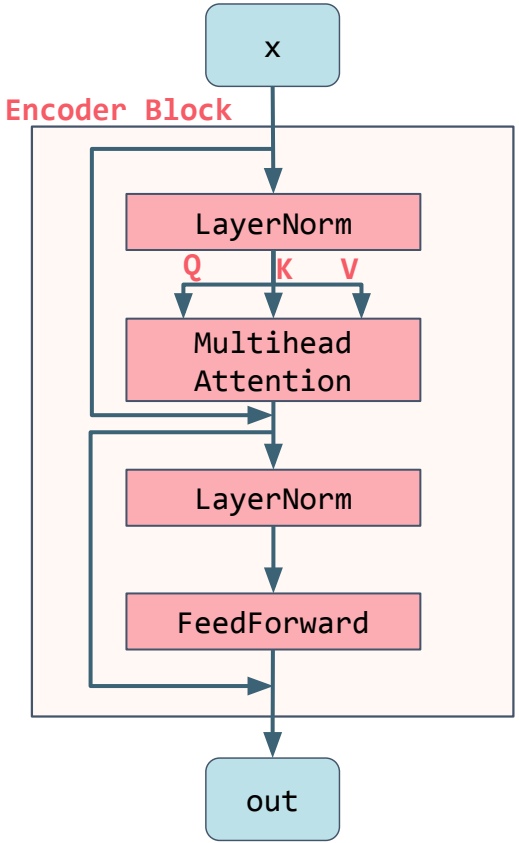

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

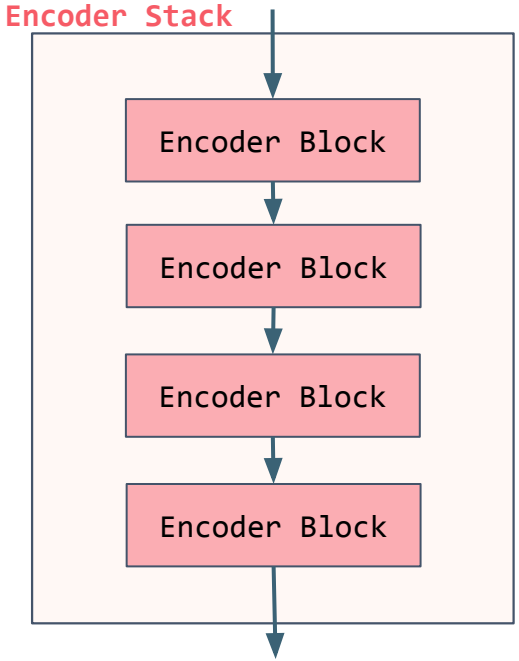

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

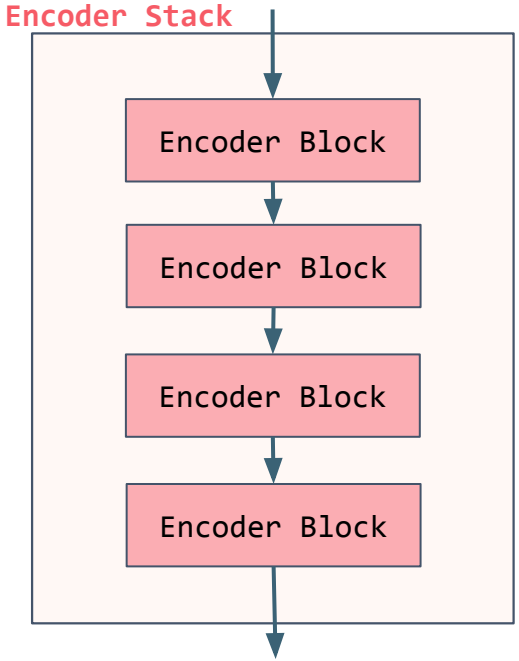

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

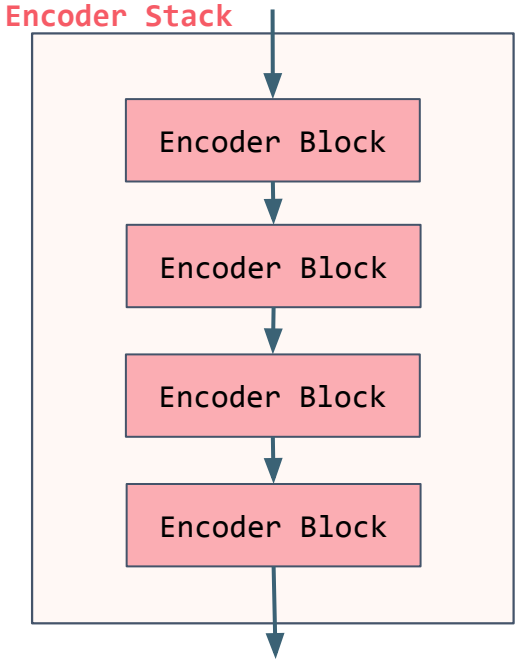

nn.TransformerEncoder

```
layer = nn.TransformerEncoderLayer(d_model=768,
                                   nhead=12,
                                   dim_feedforward=3072,
                                   dropout=0.1,
                                   batch_first=True)


stack = nn.TransformerEncoder(layer, num_layers=4)
```

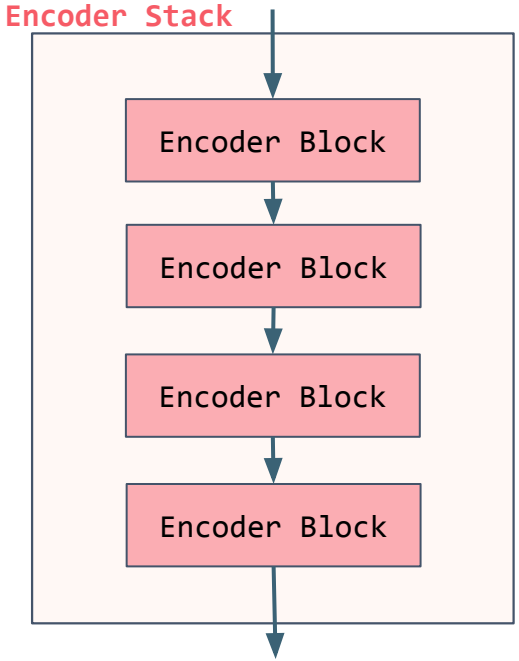

Sentiment Analysis

```
class Sentiment(nn.Module):
        . . .
    def forward(self, x):
        # Input embeddings
        x = self.embedding(x)
        pos_encoding = self.positional_encoding(x)
        x = x + pos_encoding
        x = self.dropout(x)

        for encoder_layer in self.encoder_layers:
            x = encoder_layer(x)

        # Simple pooling
        x = x.mean(dim=1)
        # 2-class linear layer nn.Linear(dim, 2)
        output = self.classifier(x)
        return output
```

Sentiment Analysis

```
class Sentiment(nn.Module):
        . . .
    def forward(self, x):
        # Input embeddings
        x = self.embedding(x)
        pos_encoding = self.positional_encoding(x)
        x = x + pos_encoding
        x = self.dropout(x)

        for encoder_layer in self.encoder_layers:
            x = encoder_layer(x)

        # Simple pooling
        x = x.mean(dim=1)
        # 2-class linear layer nn.Linear(dim, 2)
        output = self.classifier(x)
        return output
```

Sentiment Analysis

```
class Sentiment(nn.Module):
        . . .
    def forward(self, x):
        # Input embeddings
        x = self.embedding(x)
        pos_encoding = self.positional_encoding(x)
        x = x + pos_encoding
        x = self.dropout(x)

        for encoder_layer in self.encoder_layers:
            x = encoder_layer(x)

        # Simple pooling
        x = x.mean(dim=1)
        # 2-class linear layer nn.Linear(dim, 2)
        output = self.classifier(x)
        return output
```

"I really loved this movie — the acting was fantastic!"

POSITIVE

Transformer Architectures

- **Encoders:**
  - Understanding text
  - Classification
  - Named entity recognition
  - Question answering

DeepLearning.AI

Decoders

Specialized Approaches to
Natural Language Processing in Pytorch

Encoders

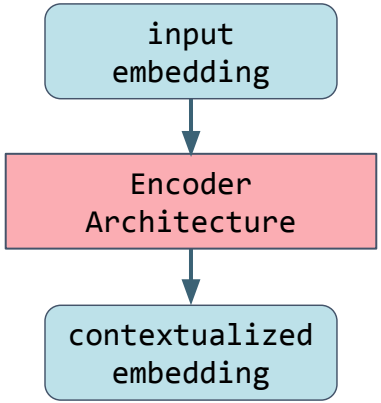
input
embedding
Encoder
Architecture
contextualized
embedding

Encoders

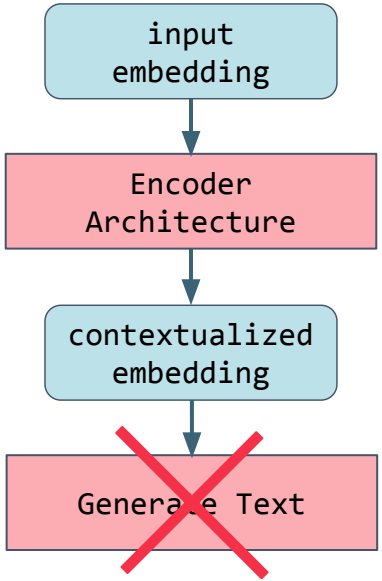
input
embedding
Encoder
Architecture
contextualized
embedding
Generate Text

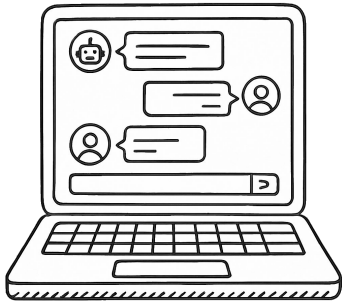

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()




    def forward(self, x):
```

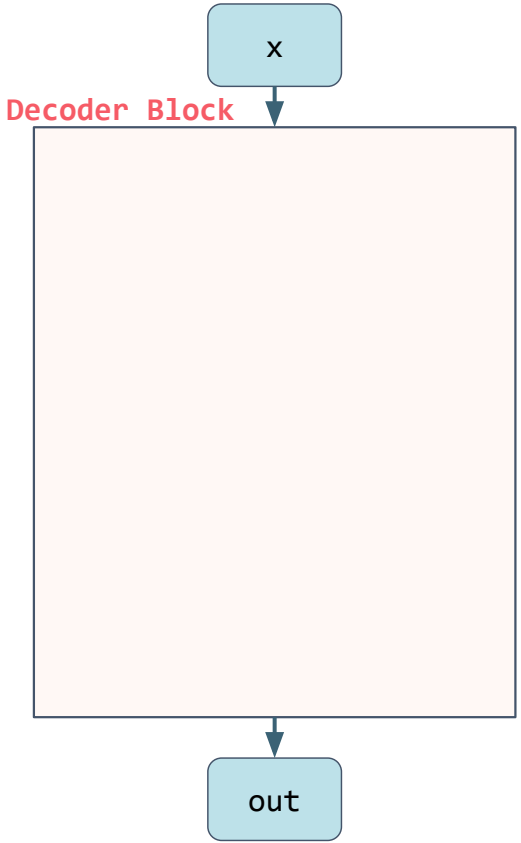

Predicting Tokens

The cat chased the _____

Predicting Tokens

Predicting Tokens

The cat chased **the** _____

Predicting Tokens

Predicting Tokens

Predicting Tokens

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()




    def forward(self, x):
```

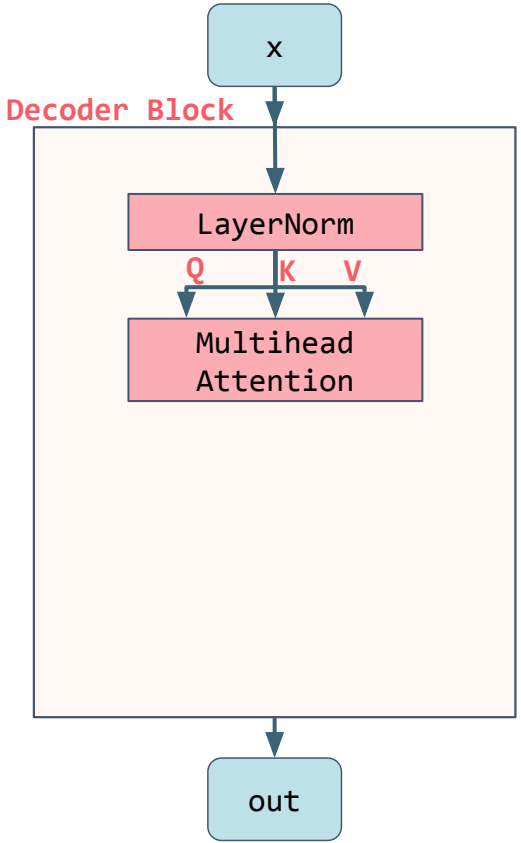

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)


    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
```

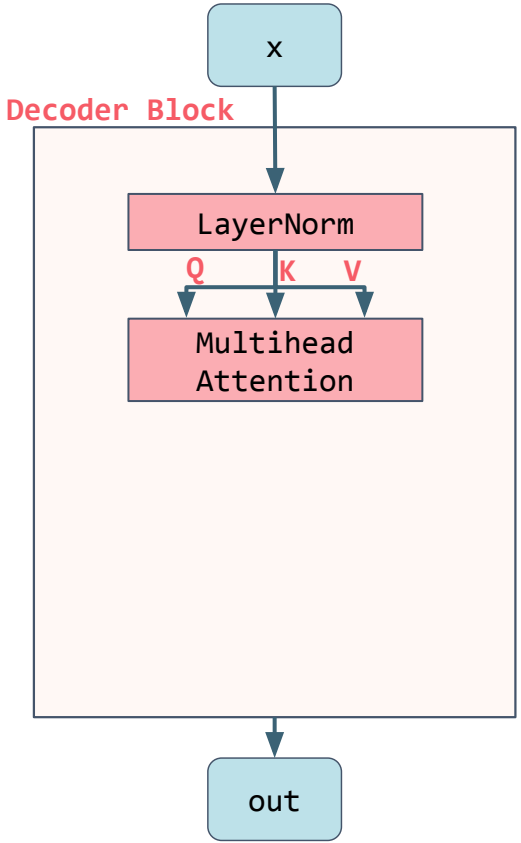

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
```

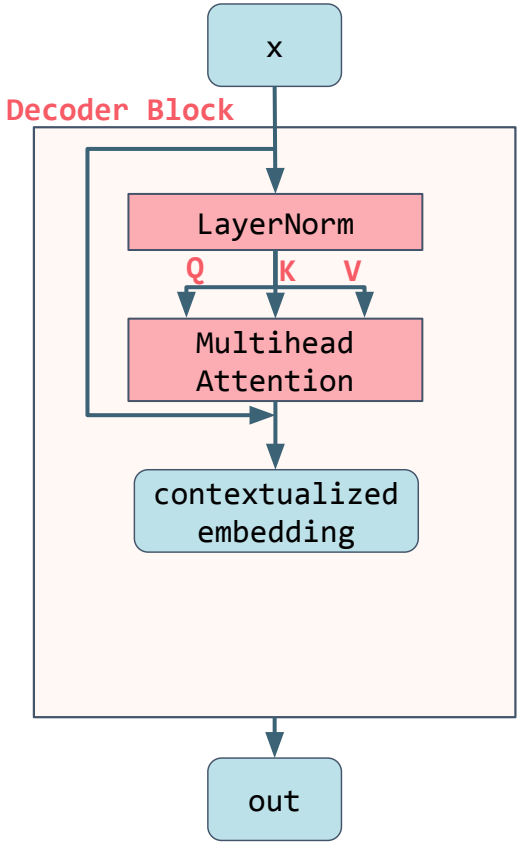

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
```

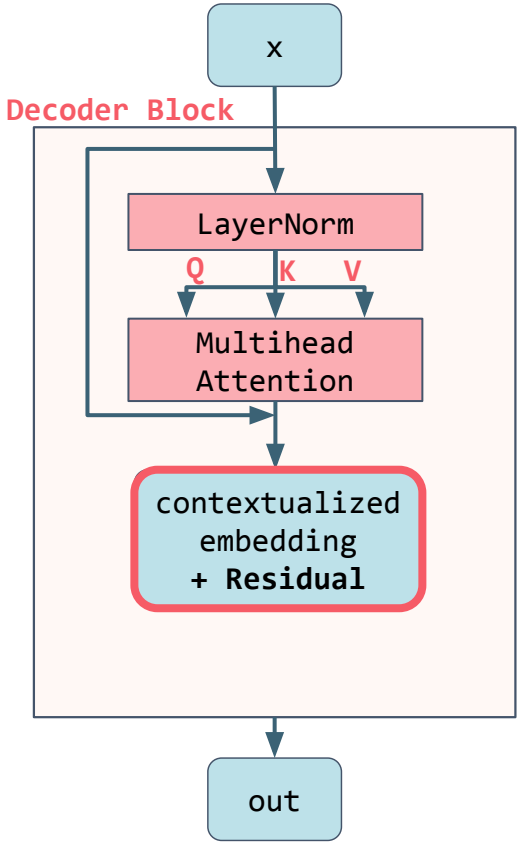

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)



    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
```

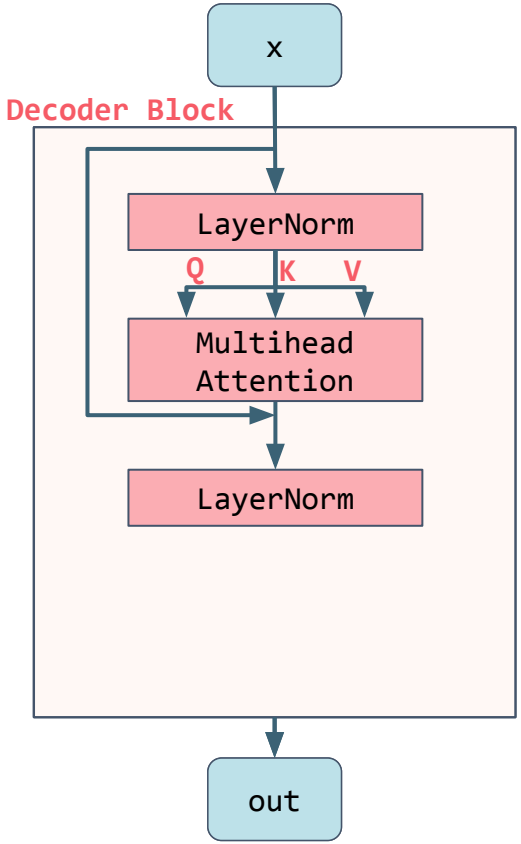

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
```

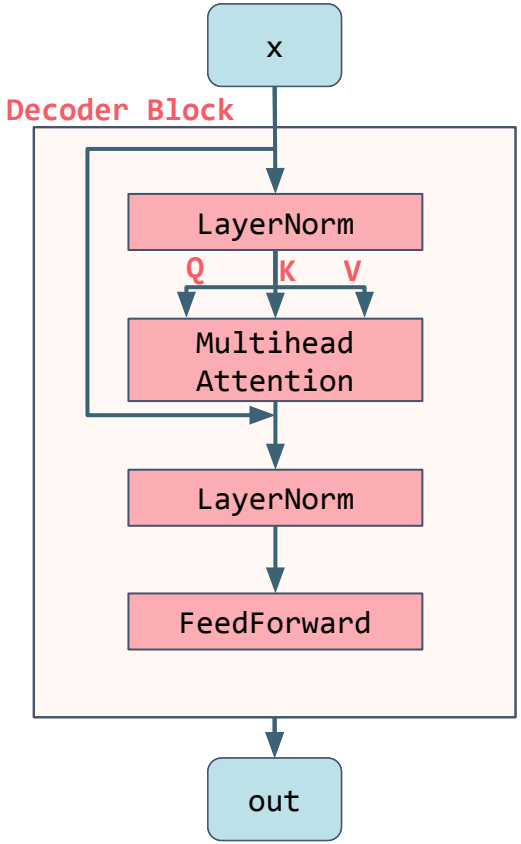

Decoder Block

```
class DecoderBlock(nn.Module):
    def __init__(self, d_model=4, nhead=1, ffn_mult=4):
        super().__init__()
        self.ln1 = nn.LayerNorm(d_model)
        self.mha = nn.MultiheadAttention(. . .)
        self.ln2 = nn.LayerNorm(d_model)
        hidden = ffn_mult * d_model
        self.ffn = nn.Sequential(. . .)

    def forward(self, x):
        x_norm = self.ln1(x)
        attn_out, _ = self.mha(x_norm, x_norm, x_norm)
        x = x + attn_out
        ffn_in = self.ln2(x)
        ffn_out = self.ffn(ffn_in)
        x = x + ffn_out
        return x
```

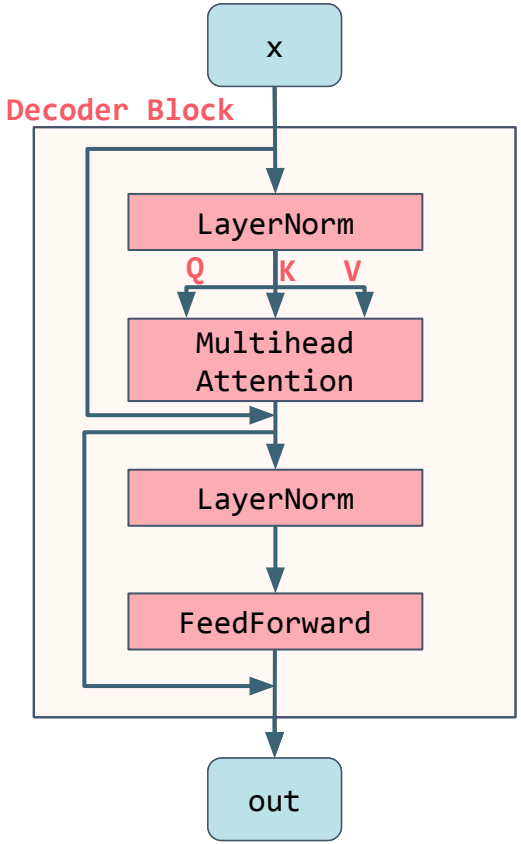

Predicting Tokens

Predicting Tokens

The cat chased **the** ___?___

Predicting Tokens

Predicting Tokens

Predicting Tokens

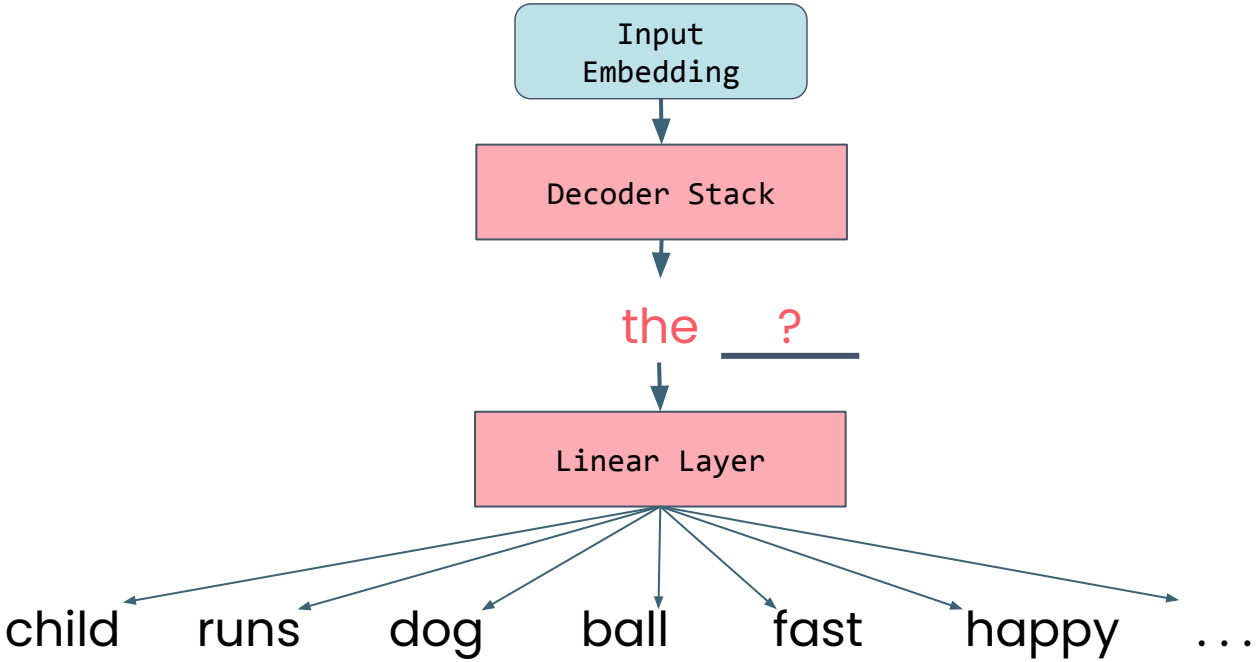

Predicting Tokens

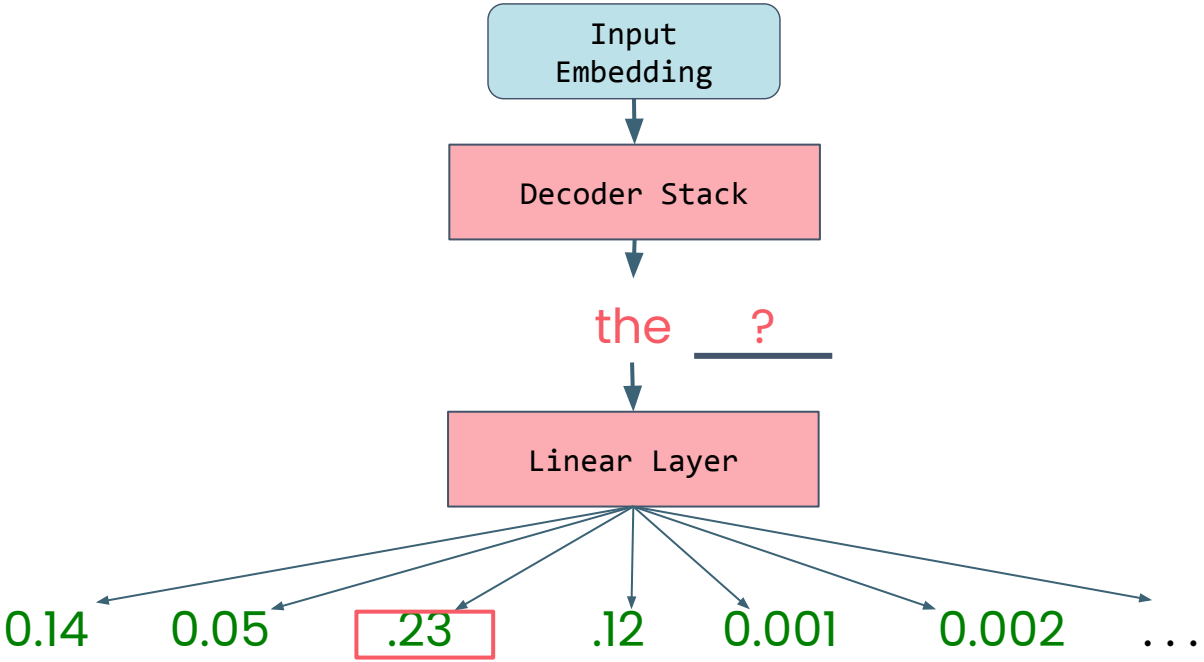

Decoder Block

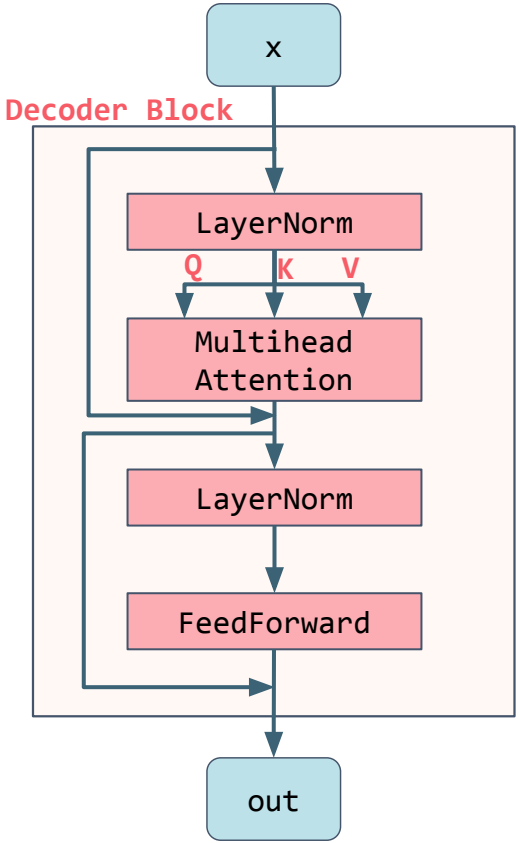

How to Train a Decoder

Input	Targets
The cat chased the ____	**dog**
When the lights went ____	**out**
I would have ____	**left**
She saw a ____	**stranger**

How to Train a Decoder

The cat chased the dog

Input	Targets
The ____	**cat**
The cat ____	**chased**
The cat chased ____	**the**
The cat chased the ____	**dog**

Causal Masks

```
mha = nn.MultiheadAttention(. . .)
```

Causal Masks

```
mha = nn.MultiheadAttention(. . .)


def forward(self, x):

        # Custom masks
        out, _ = mha(attn_mask=custom_mask, . . .)
```

Causal Masks

```
mha = nn.MultiheadAttention(. . .)


def forward(self, x):

        # Custom masks
        out, _ = mha(attn_mask=custom_mask, . . .)

        # is_causal
        out, _ = mha(is_causal=True, . . .)
```

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

		the	cat	chased	the	dog
	the	1.21	-0.259	-0.762	-0.995	0.054
“the	cat	-0.216	0.604	-1.650	-1.061	1.876
	chased	1.462	-0.410	-1.009	-0.990	0.822
	the	1.506	-1.180	0.659	-1.810	0.521
	dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

		the	cat	chased	the	dog
	the	1.21	-0.259	-0.762	-0.995	0.054
“the	cat	-0.216	0.604	-1.650	-1.061	1.876
“the cat	chased	1.462	-0.410	-1.009	-0.990	0.822
	the	1.506	-1.180	0.659	-1.810	0.521
	dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Attention Scores

		the	cat	chased	the	dog
	the	1.21				
“the	cat	-0.216	0.604			
“the cat	chased	1.462	-0.410	-1.009		
“the cat chased	the	1.506	-1.180	0.659	-1.810	
“the cat chased the	dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

```
def make_causal_mask(sz: int):
    mask = torch.full((sz, sz), float("-inf"))
    mask = torch.triu(mask, diagonal=1)
    return mask
```

0.	-inf	-inf	-inf	-inf
0.	0.	-inf	-inf	-inf
0.	0.	0.	-inf	-inf
0.	0.	0.	0.	-inf
0.	0.	0.	0.	0.

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-0.259	-0.762	-0.995	0.054
cat	-0.216	0.604	-1.650	-1.061	1.876
chased	1.462	-0.410	-1.009	-0.990	0.822
the	1.506	-1.180	0.659	-1.810	0.521
dog	0.995	0.521	0.627	-0.511	0.315

+

0.	-inf	-inf	-inf	-inf
0.	0.	-inf	-inf	-inf
0.	0.	0.	-inf	-inf
0.	0.	0.	0.	-inf
0.	0.	0.	0.	0.

Causal Masks

Attention Scores

	the	cat	chased	the	dog
the	1.21	-inf	-inf	-inf	-inf
cat	-0.216	0.604	-inf	-inf	-inf
chased	1.462	-0.410	-1.009	-inf	-inf
the	1.506	-1.180	0.659	-1.810	-inf
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Apply softmax

	the	cat	chased	the	dog
the	1.21	-inf	-inf	-inf	-inf
cat	-0.216	0.604	-inf	-inf	-inf
chased	1.462	-0.410	-1.009	-inf	-inf
the	1.506	-1.180	0.659	-1.810	-inf
dog	0.995	0.521	0.627	-0.511	0.315

Causal Masks

Apply softmax

	the	cat	chased	the	dog
the	1.	0.	0.	0.	0.
cat	0.3058	0.6942	0.	0.	0.
chased	0.8075	0.1242	0.0682	0.	0.
the	0.6523	0.0445	0.2796	0.0237	0.
dog	0.3286	0.2046	0.2274	0.0729	0.1665

Causal Masks

Attention Weights

	the	cat	chased	the	dog
the	1.	0.	0.	0.	0.
cat	0.3058	0.6942	0.	0.	0.
chased	0.8075	0.1242	0.0682	0.	0.
the	0.6523	0.0445	0.2796	0.0237	0.
dog	0.3286	0.2046	0.2274	0.0729	0.1665

Predicting Tokens

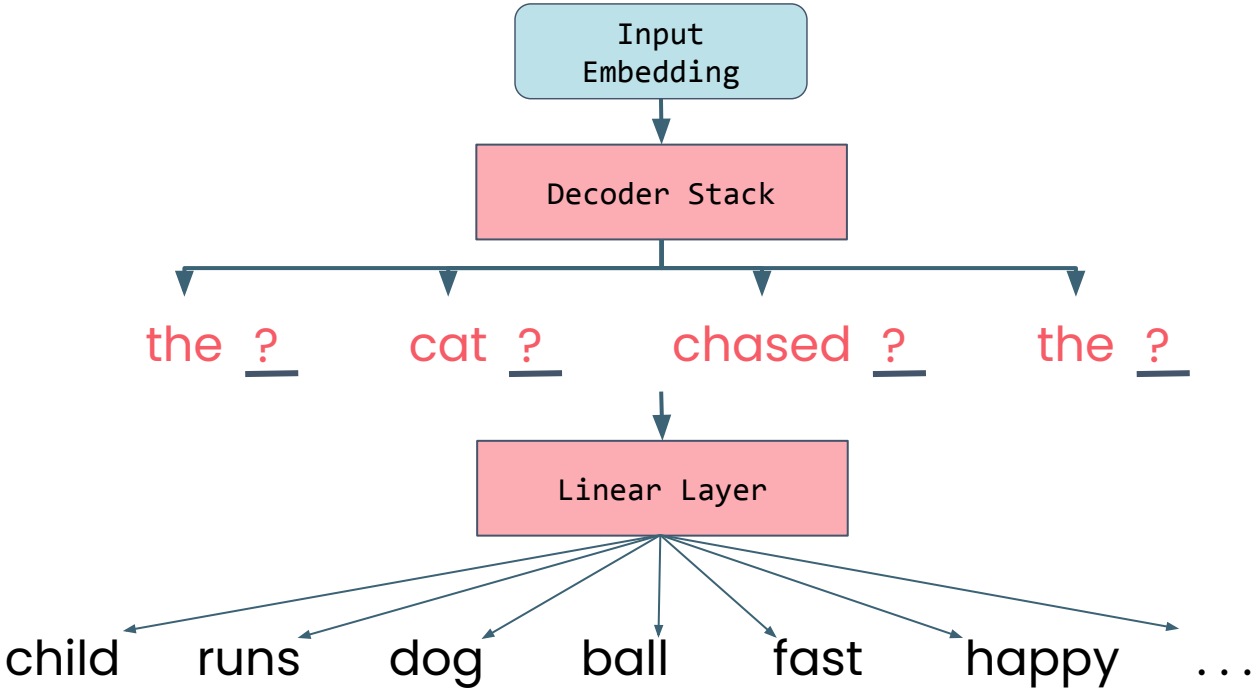

Training Data

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']] # Convert to token IDs
```

Training Data

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']] # Convert to token IDs


# Inputs
inputs = token_sequences[:, :-1] # Inputs as strings ['the', 'cat', 'chased', 'the']
```

Training Data

```
token_sequences = [['the', 'cat', 'chased', 'the', 'dog']] # Convert to token IDs


# Inputs
inputs = token_sequences[:, :-1] # Inputs as strings ['the', 'cat', 'chased', 'the']


# Targets
targets = token_sequences[:, 1:] # Targets as strings ['cat', 'chased', 'the', 'dog']
```

Training Data

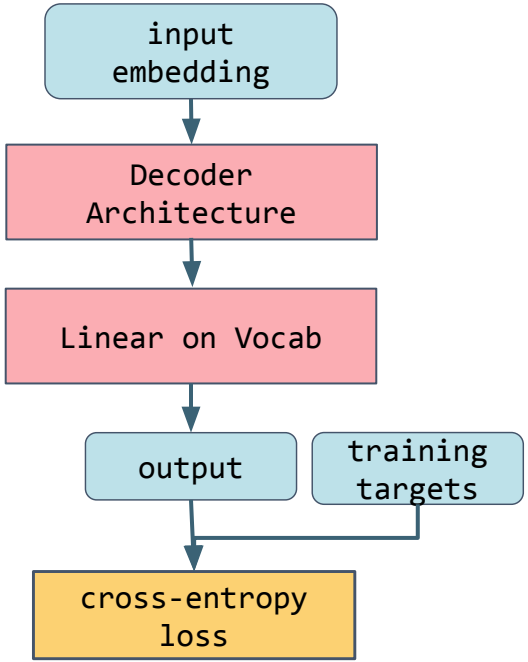

Training Data

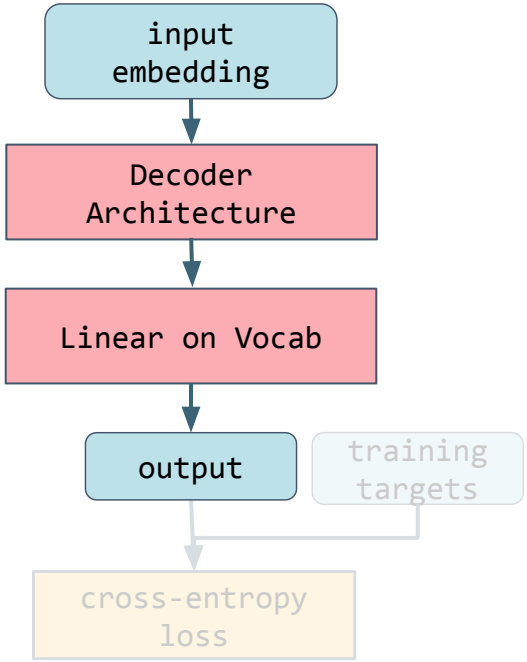

Training Data

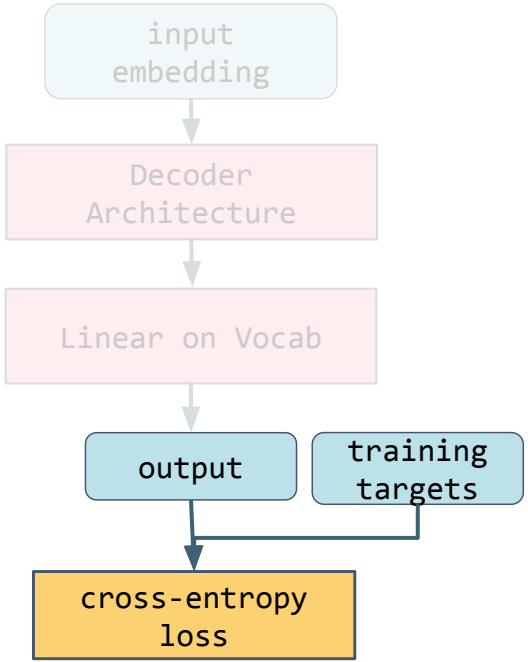

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder
```

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder

nn.TransformerDecoderLayer
nn.TransformerDecoder
```

For encoder decoder models

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder

nn.TransformerDecoderLayer
nn.TransformerDecoder
```

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder
```

\+ causal mask

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder


out = encoder_layer (is_causal=True, . . .)
out = encoder_layer (src_mask=custom_mask, . . .)

out = encoder_stack (is_causal=True, . . .)
out = encoder_stack (mask=custom_mask, . . .)
```

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder


out = encoder_layer (is_causal=True, . . .)
out = encoder_layer (src_mask=custom_mask, . . .)

out = encoder_stack (is_causal=True, . . .)
out = encoder_stack (mask=custom_mask, . . .)
```

Decoder Only Architectures

```
import torch.nn as nn

nn.TransformerEncoderLayer
nn.TransformerEncoder


out = encoder_layer (is_causal=True, . . .)
out = encoder_layer (src_mask=custom_mask, . . .)

out = encoder_stack (is_causal=True, . . .)
out = encoder_stack (mask=custom_mask, . . .)
```

DeepLearning.AI

Encoder - Decoder

Specialized Approaches to
Natural Language Processing in Pytorch

Transformer Architectures

- **Encoders:**
  - Analyze using rich contextual representations
  - Does **not** generate new tokens

- **Decoders:**
  - Can generate new text from old
  - But only by continuing from whatever it was given

Decoders as Chatbots

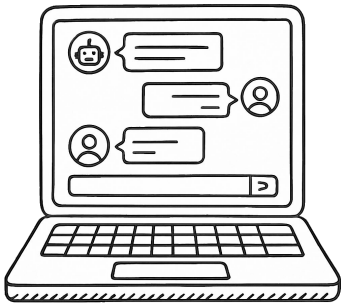

<user> </user>

<assistant> </assistant>

<s> </s>

<user> *did the cat chase the dog?* </user><assistant> . . .

Decoders as Chatbots

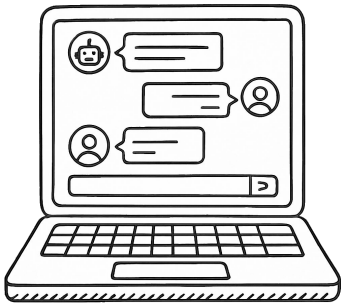

<user> </user>

<assistant> </assistant>

<s> </s>

<user> *did the cat chase the dog?* </user><assistant> *it depends on who you ask* </assistant>

Translation

🇮🇹

The cat chased the dog . . . *Il gatto ha inseguito il cane*

<user> *Translate “the cat chased the dog” into Italian*</user><assistant> . . .

Input

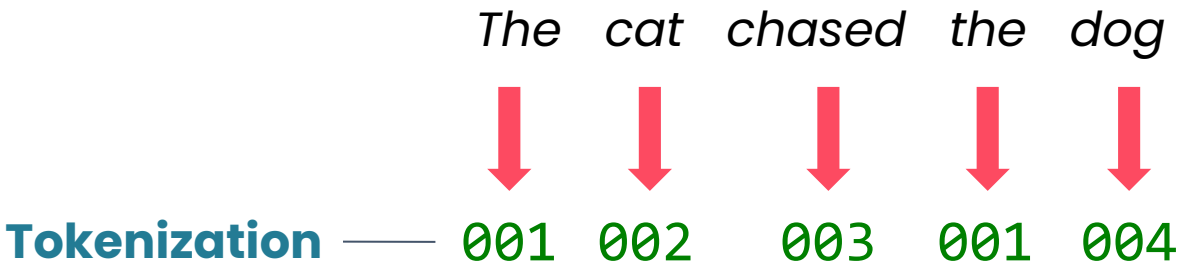

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()




    def forward(self, src):
```

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()




    def forward(self, src):
```

Encoder

src

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()




    def forward(self, src):
        padding_mask = create_padding_mask(src)
```

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        self.token_emb = nn.Embedding(. . .)
        self.pos_enc = PositionalEncoding(. . .)



    def forward(self, src):
        padding_mask = create_padding_mask(src)
```

src

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        self.token_emb = nn.Embedding(. . .)
        self.pos_enc = PositionalEncoding(. . .)



    def forward(self, src):
       padding_mask = create_padding_mask(src)
       src = self.token_emb(src) + self.pos_enc(src)
```

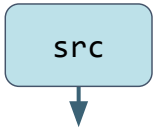

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        . . .




    def forward(self, src):
        . . .
```

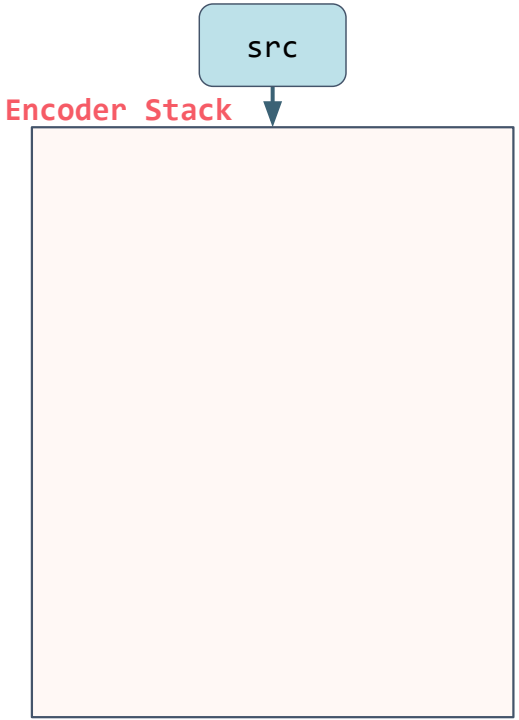

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        . . .
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)



    def forward(self, src):
        . . .
```

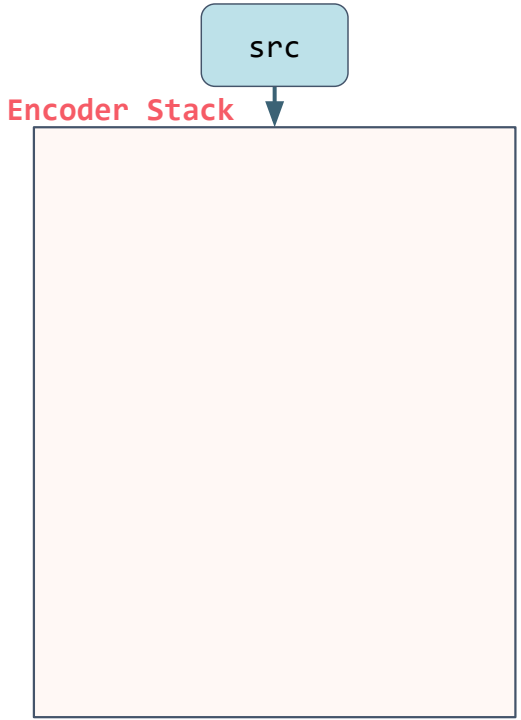

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        . . .
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)

        self.transformer_encoder = nn.TransformerEncoder(
            enc_layer,
            num_layers=num_layers)
    def forward(self, src):
        . . .
```

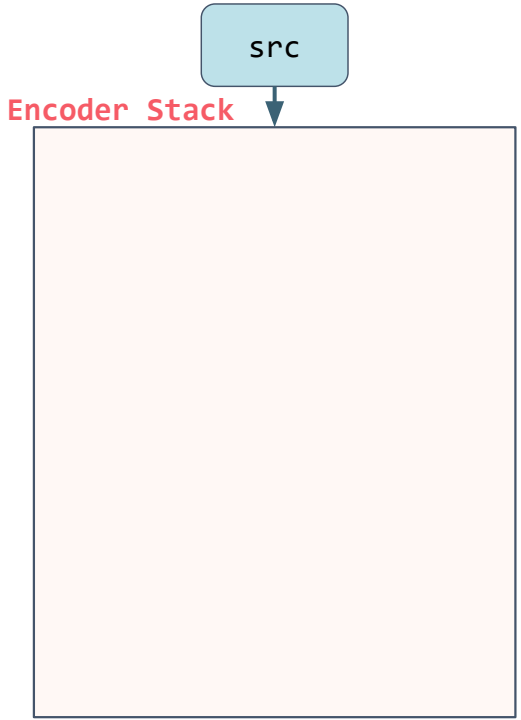

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        . . .
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)

        self.transformer_encoder = nn.TransformerEncoder(
            enc_layer,
            num_layers=num_layers)
    def forward(self, src):
        . . .
```

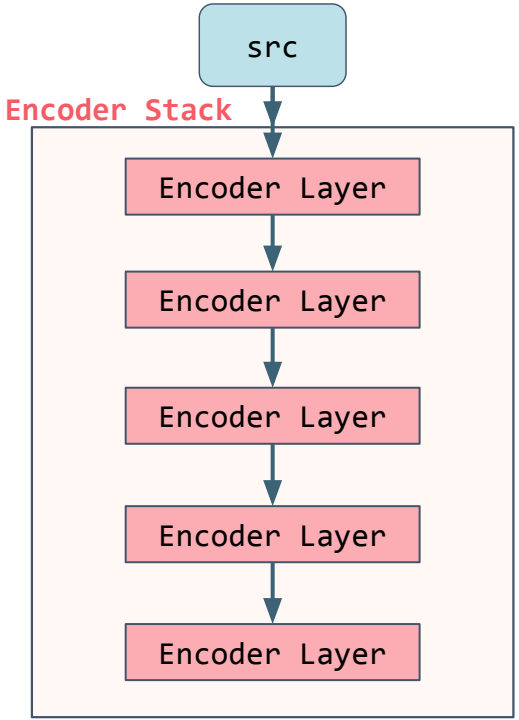

Encoder

```
class Encoder(nn.Module):
    def __init__():
        super().__init__()
        . . .
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)

        self.transformer_encoder = nn.TransformerEncoder(
            enc_layer,
            num_layers=num_layers)
    def forward(self, src):
        . . .
        memory = self.transformer_encoder(src,
            src_key_padding_mask=padding_mask)
        return memory
```

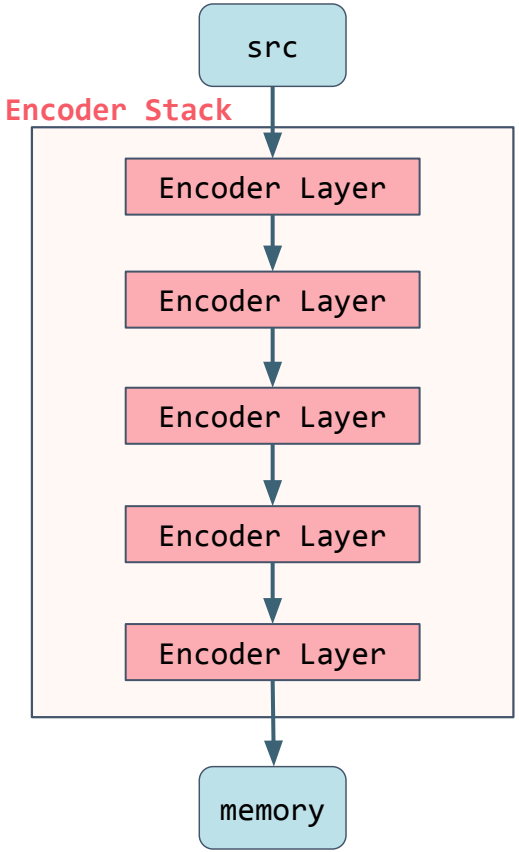

Translation

The cat chased the dog

Il gatto ha inseguito il cane

Translation

The cat chased the dog

‹s›

Decoder

```
class Decoder(nn.Module):
    def __init__():
        super().__init__()
        self.token_emb = nn.Embedding(. . .)
        self.pos_enc = PositionalEncoding(. . .)




    def forward(self, tgt, memory, memory_padding_mask):
       src_key_padding_mask = create_pad_mask(tgt)
       src = self.src_tok(tgt) + self.src_pos(tgt)
```

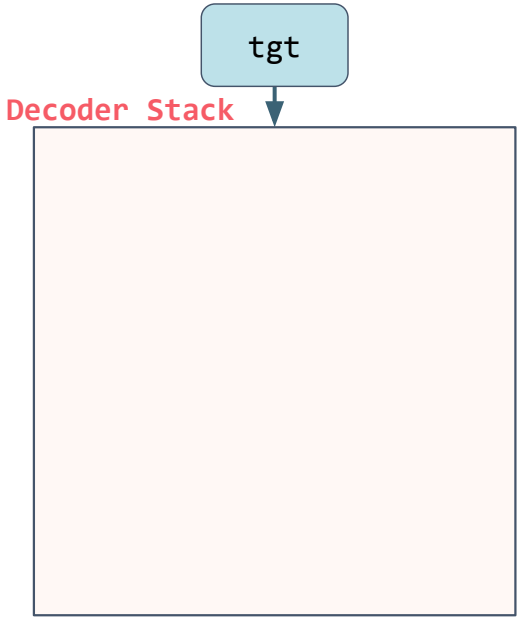

Decoder

```
class Decoder(nn.Module):
    def __init__():
        super().__init__()
        self.token_emb = nn.Embedding(. . .)
        self.pos_enc = PositionalEncoding(. . .)



    def forward(self, tgt, memory, memory_padding_mask):
        src_key_padding_mask = create_pad_mask(tgt)
        src = self.src_tok(tgt) + self.src_pos(tgt)
```

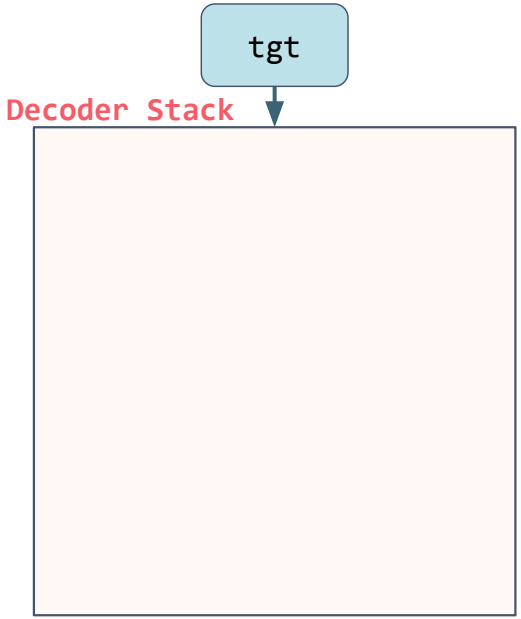

Decoder

```
class Decoder(nn.Module):
    def __init__():
        super().__init__()
        self.token_emb = nn.Embedding(. . .)
        self.pos_enc = PositionalEncoding(. . .)



    def forward(self, tgt, memory, memory_padding_mask):
        src_key_padding_mask = create_pad_mask(tgt)
        src = self.src_tok(tgt) + self.src_pos(tgt)
```

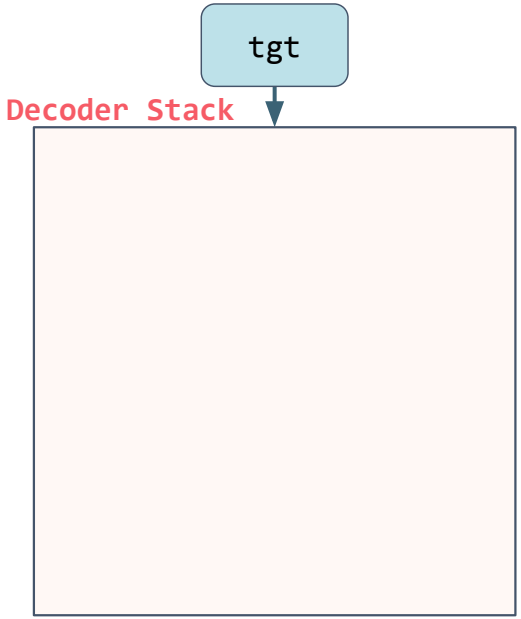

nn.DecoderLayer

```
import torch.nn as nn

nn.TransformerDecoderLayer
nn.TransformerDecoder
```

For encoder decoder models

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .




    def forward(self, tgt, memory, memory_padding_mask):
        . . .
```

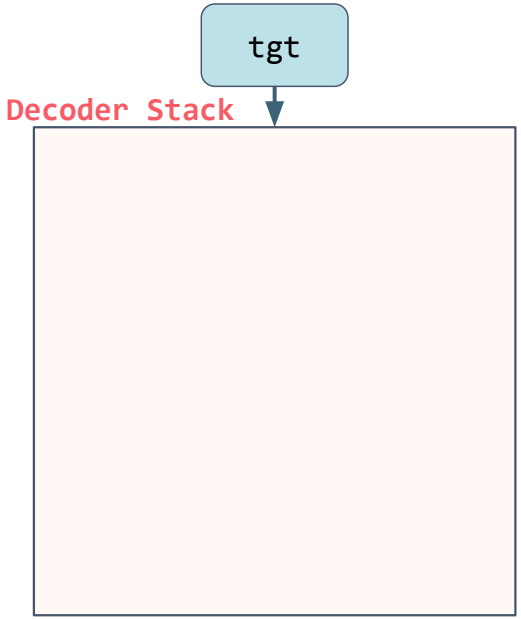

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
        decoder_layer = nn.TransformerDecoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
```

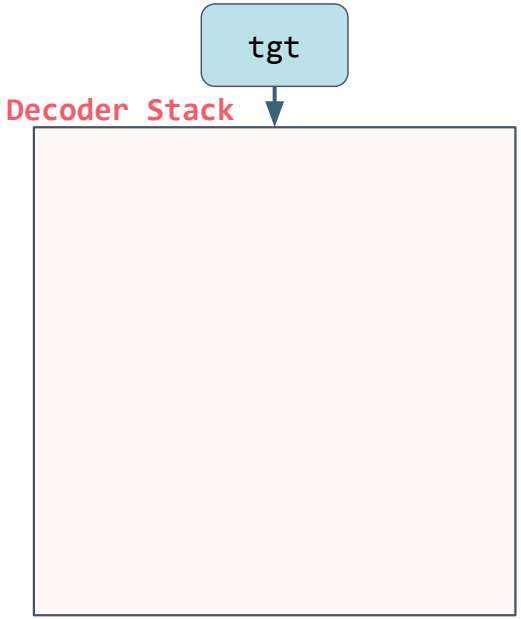

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
        decoder_layer = nn.TransformerDecoderLayer(
            d_model=d_model, nhead=nhead,
            dim_feedforward=dim_feedforward,
            dropout=dropout, batch_first=True)

        self.transformer_decoder = nn.TransformerDecoder(
            decoder_layer,
            num_layers=num_layers)

    def forward(self, tgt, memory, memory_padding_mask):
        . . .
```

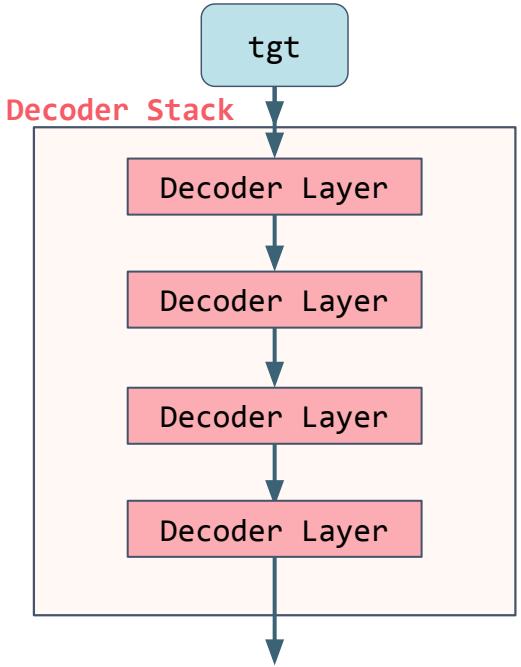

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
```

```
def forward(self, tgt, memory, memory_padding_mask):
    . . .
    decoded= self.transformer_decoder(
        tgt,
        memory,
        tgt_mask=tgt_causal_mask
        memory_mask=None
        tgt_key_padding_mask=tgt_padding_mask
        memory_key_padding_mask=memory_padding_mask)
```

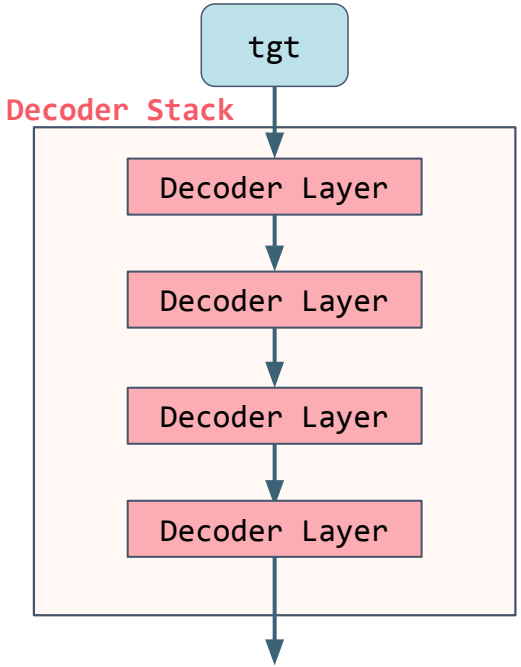

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(
            tgt,
            memory,
            tgt_mask=tgt_causal_mask
            memory_mask=None
            tgt_key_padding_mask=tgt_padding_mask
            memory_key_padding_mask=memory_padding_mask)
```

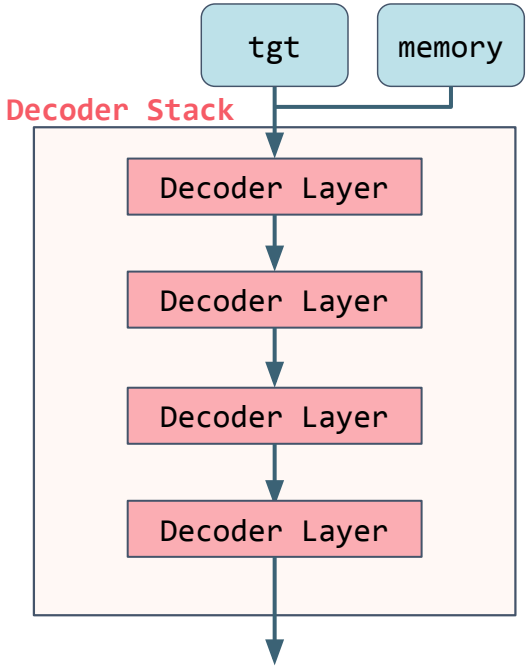

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(
            tgt,
            memory,
            tgt_mask=tgt_causal_mask
            memory_mask=None
            tgt_key_padding_mask=tgt_padding_mask
            memory_key_padding_mask=memory_padding_mask)
```

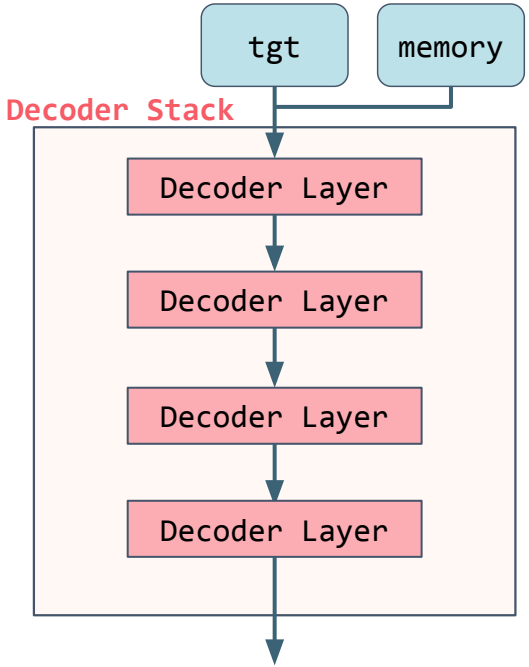

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(
            tgt,
            memory,
            tgt_mask=tgt_causal_mask
            memory_mask=None
            tgt_key_padding_mask=tgt_padding_mask
            memory_key_padding_mask=memory_padding_mask)
```

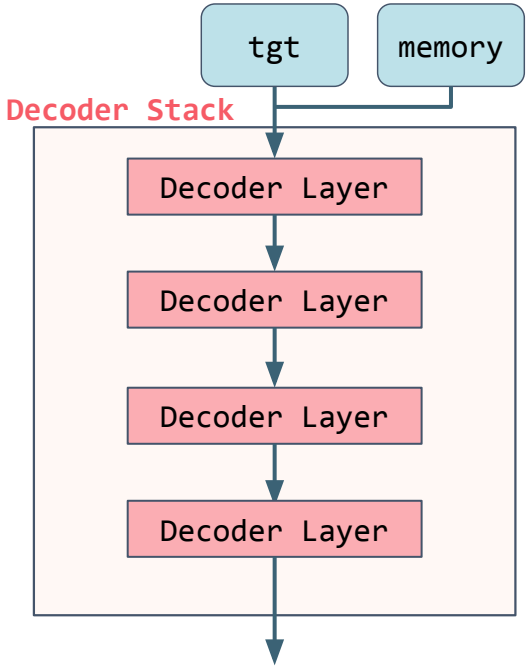

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(
            tgt,
            memory,
            tgt_mask=tgt_causal_mask
            memory_mask=None
            tgt_key_padding_mask=tgt_padding_mask
            memory_key_padding_mask=memory_padding_mask)
```

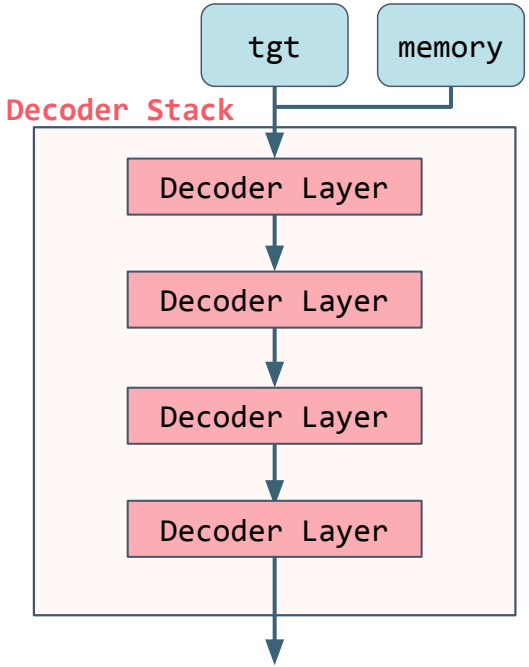

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(
            tgt,
            memory,
            tgt_mask=tgt_causal_mask
            memory_mask=None
            tgt_key_padding_mask=tgt_padding_mask
            memory_key_padding_mask=memory_padding_mask)
```

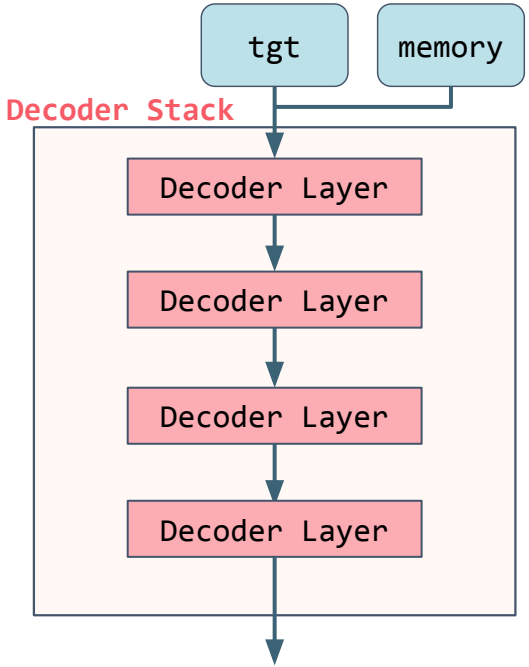

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(. . .)
```

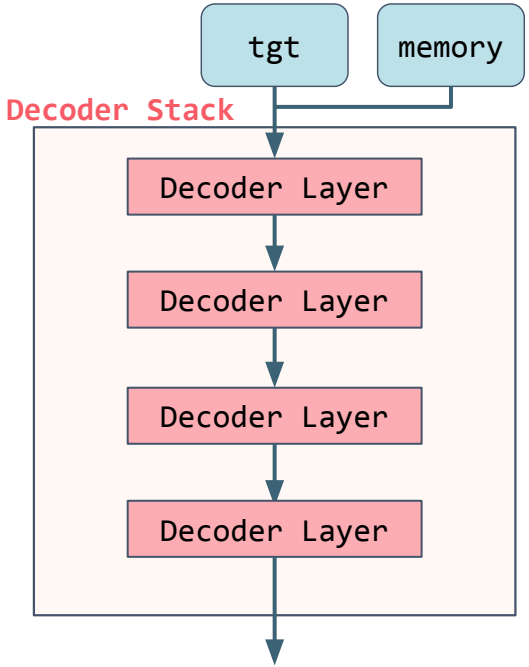

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
        self.out_proj = nn.Linear(d_model, vocab_size)


    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(. . .)
```

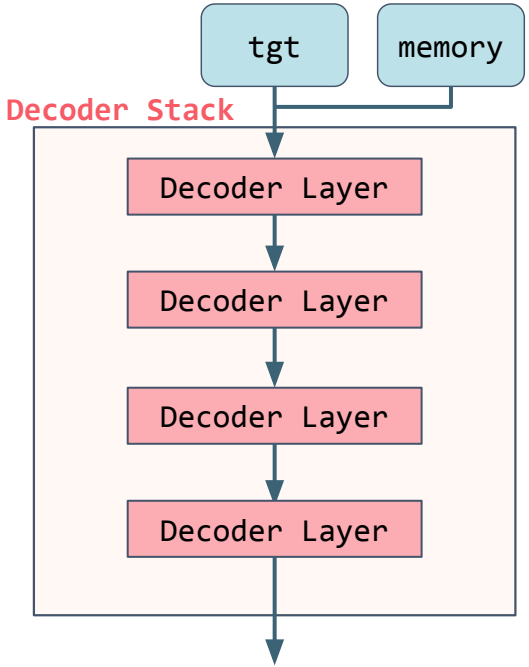

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
        self.out_proj = nn.Linear(d_model, vocab_size)



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(. . .)

        output = self.output_proj(decoded)

        return output
```

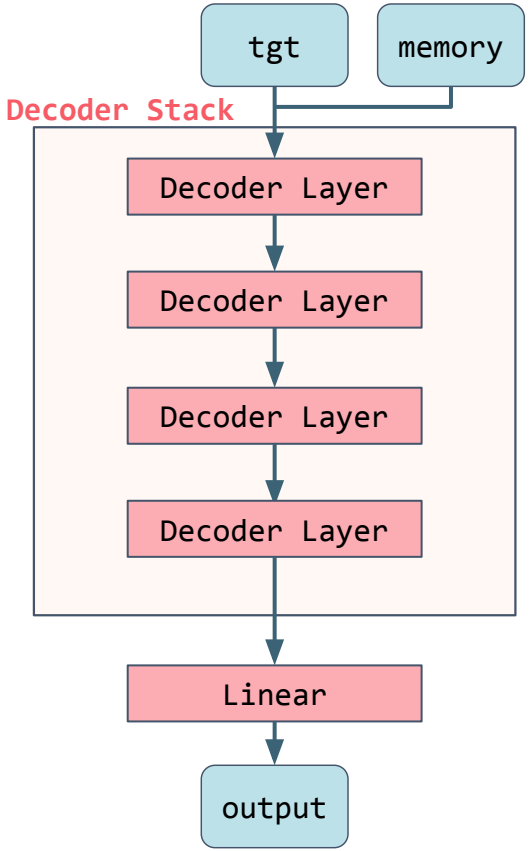

Decoder

```
class Decoder(nn.Module):
    def __init__():
        . . .
        self.out_proj = nn.Linear(d_model, vocab_size)



    def forward(self, tgt, memory, memory_padding_mask):
        . . .
        decoded= self.transformer_decoder(. . .)

        output = self.output_proj(decoded)

        return output
```

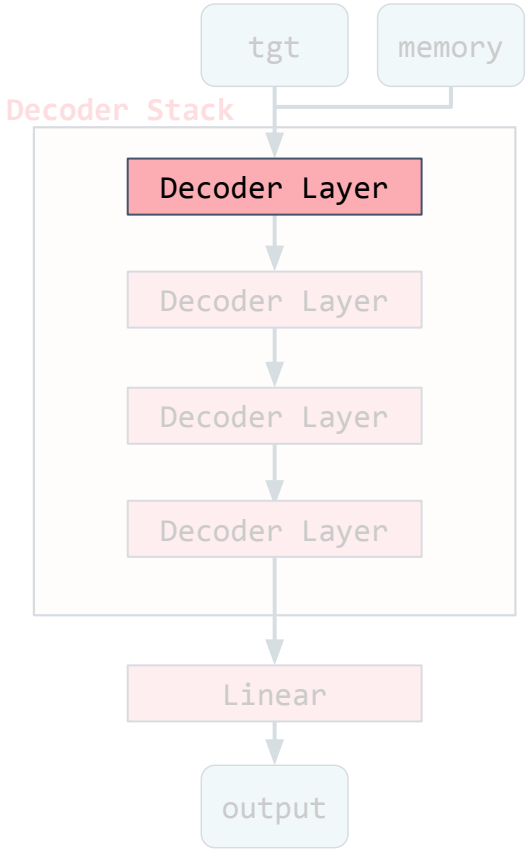

Decoder Layer

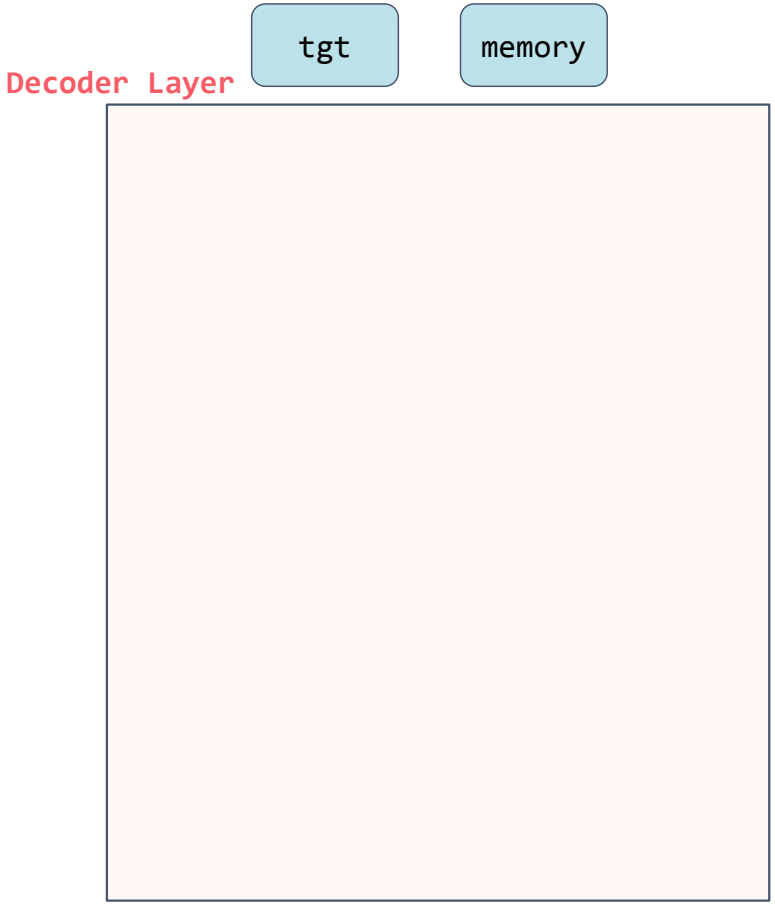

Decoder Layer

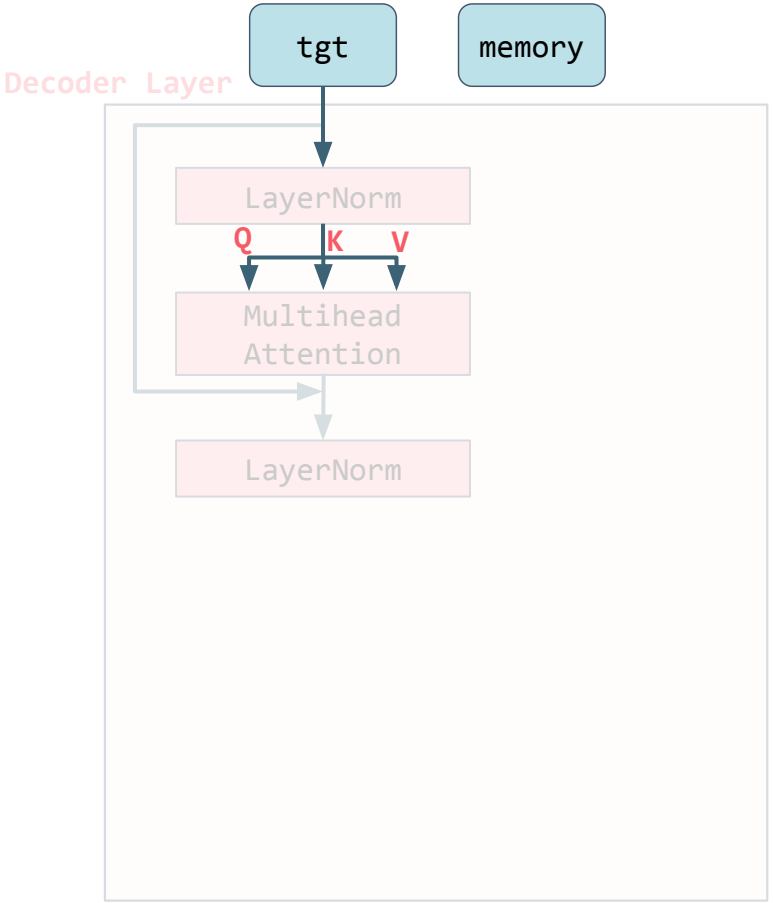

Decoder Layer

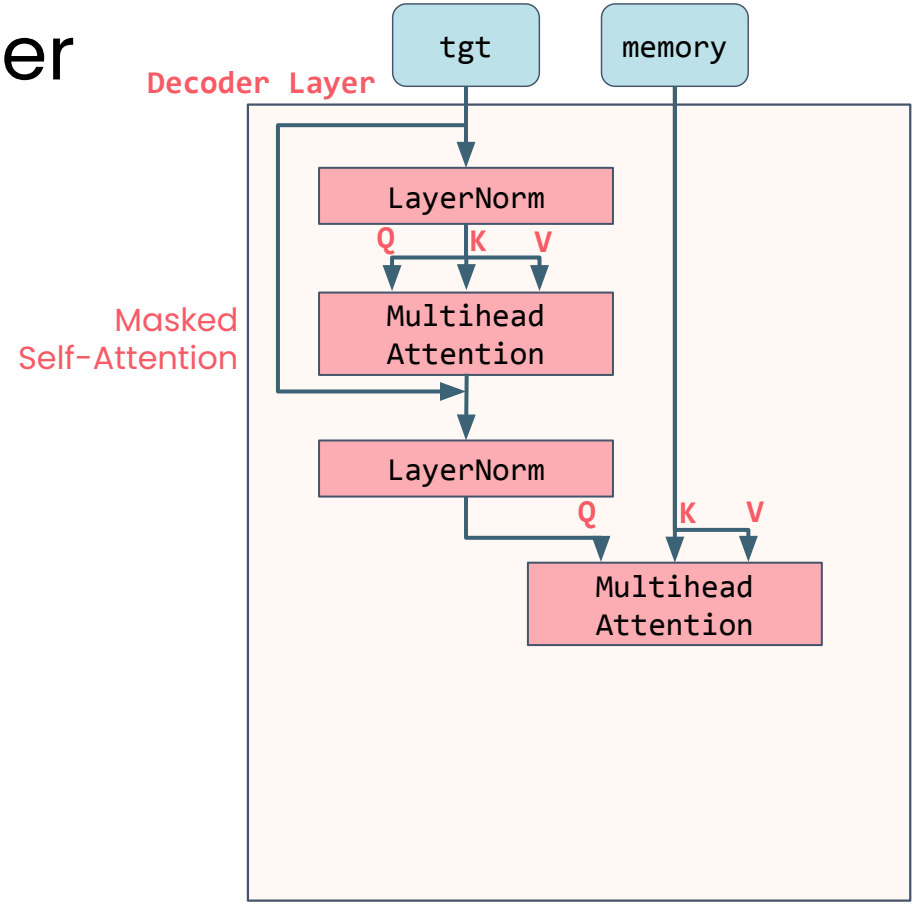

Query
What I'm looking for

Key
What I look like

Value
If you attend to me, here's what I'll give you

Decoder Layer

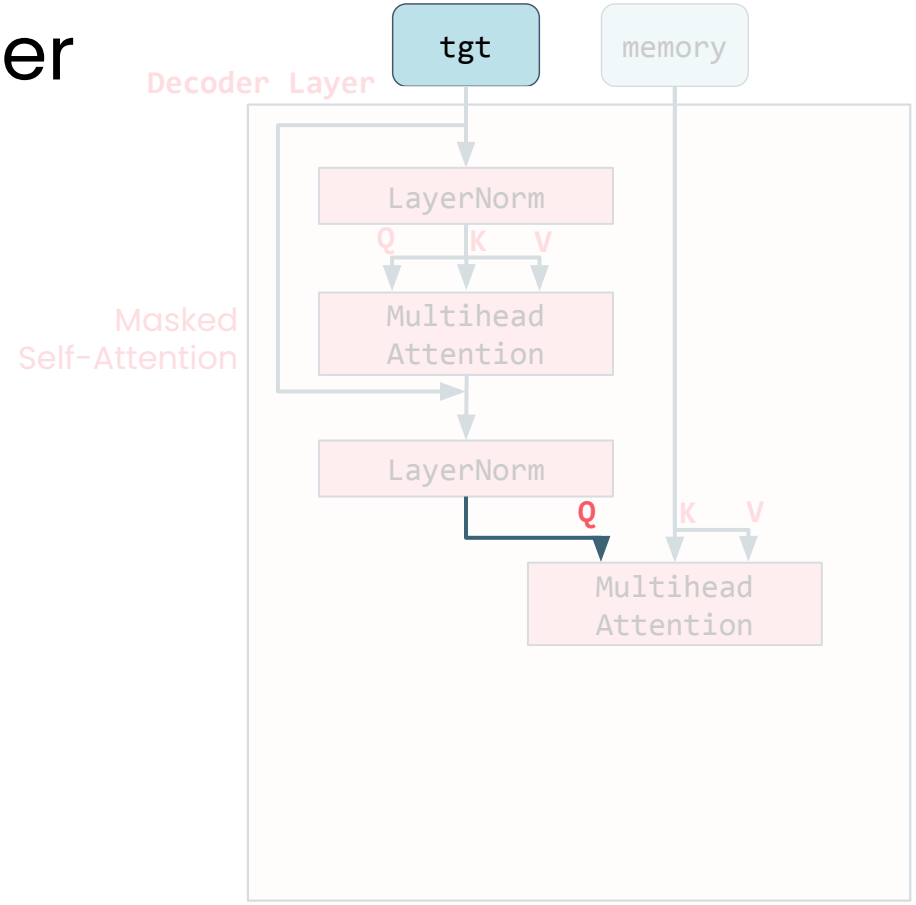

Query
What I'm looking for

Key
What I look like

Value
If you attend to me, here's what I'll give you

Decoder Layer

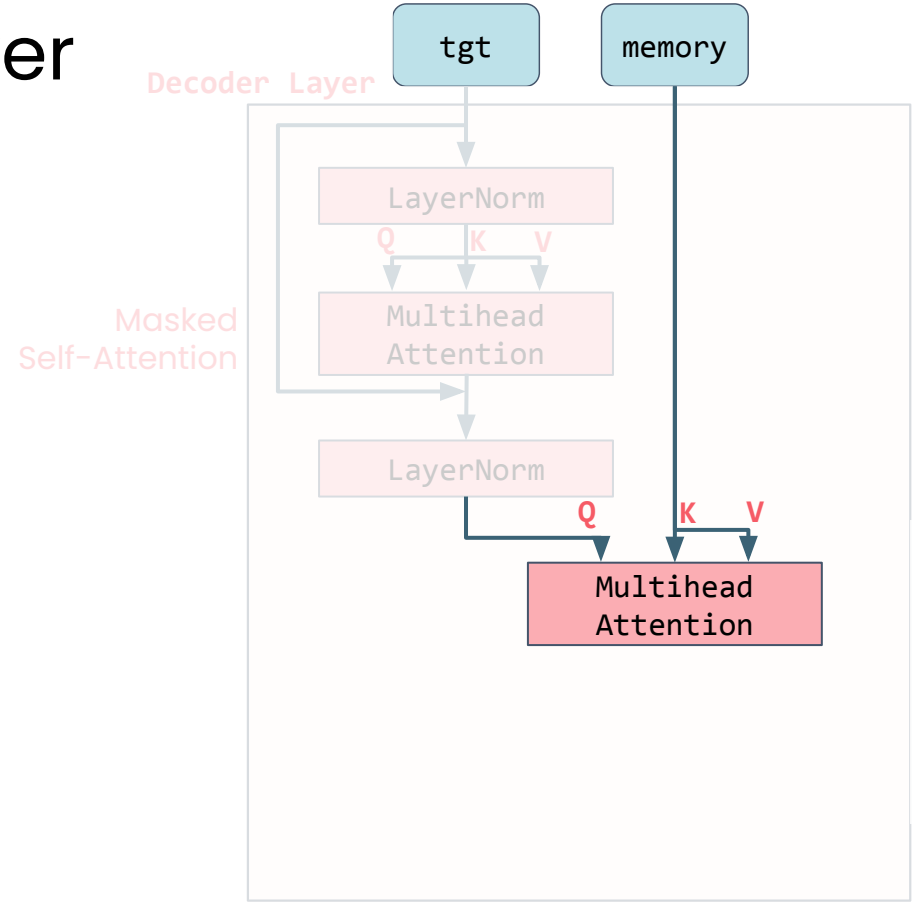

Query
What I'm looking for

Key
What I look like

Value
If you attend to me, here's what I'll give you

Decoder Layer

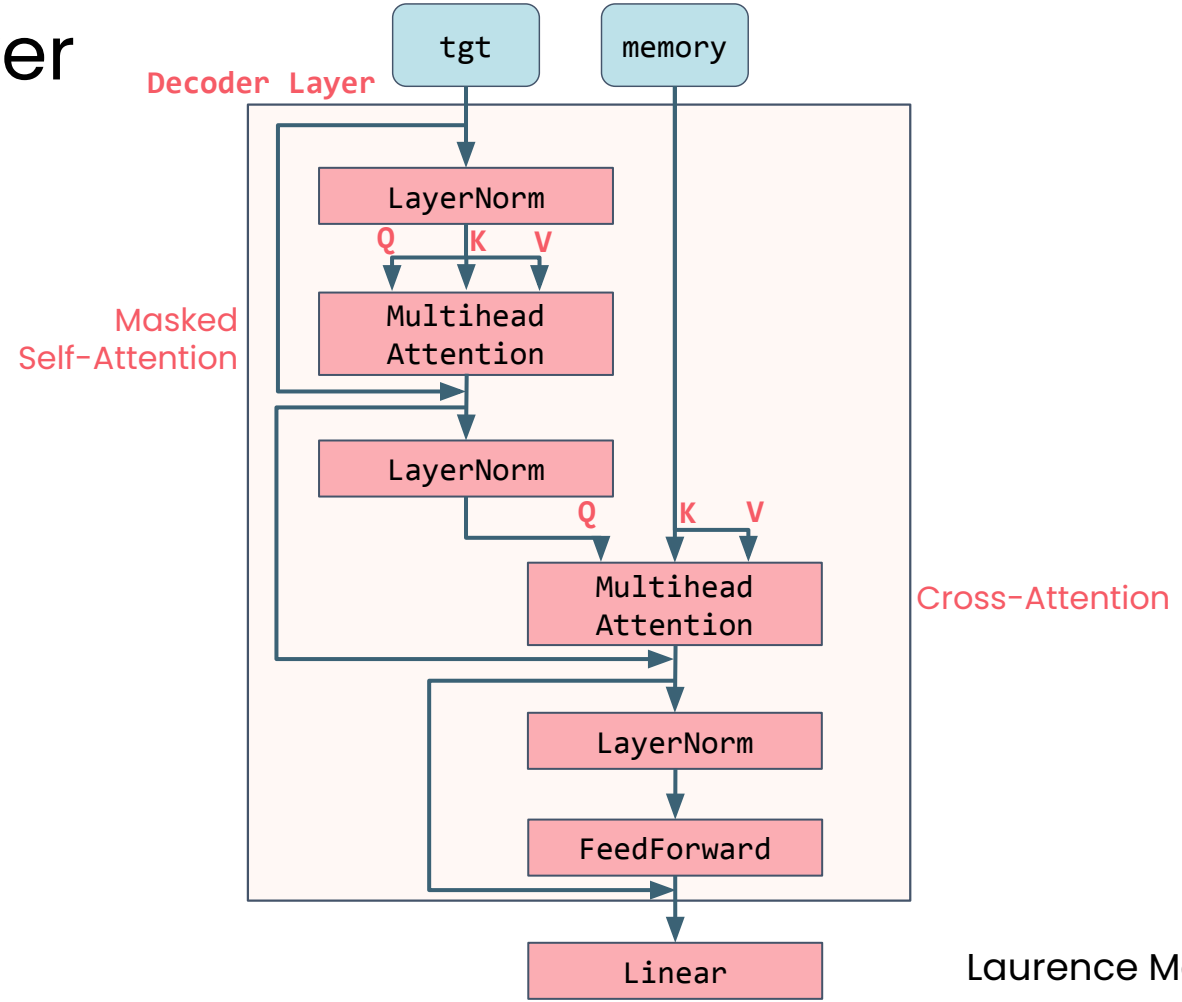

Encoder Decoder

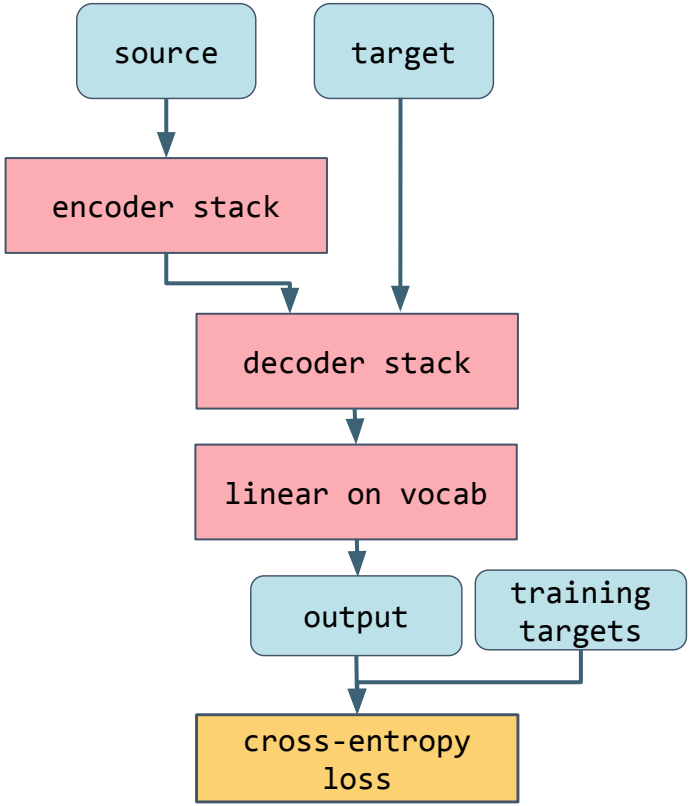

Transformers