



DeepLearning.AI

Model Serialization and Version Control

Preparing Models for Deployment in PyTorch

Deployment

- **Save and Track**
- **Export**
- **Optimize**



Deployment

- **Save and Track**
- Export
- Optimize

Serialization and Version Control





Saving and Tracking




Trained Model

Saving and Tracking





Trained Model

Saving and Tracking



Saving and Tracking



Serialization

```
model = SimpleCNN()  
torch.save(model, 'my_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model, 'my_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model, 'my_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model, 'my_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model.state_dict(), 'my_trained_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model.state_dict(), 'my_trained_model.pt')
```

Serialization

```
model = SimpleCNN()  
torch.save(model.state_dict(), 'my_trained_model.pt')  
  
print(model.state_dict().keys())
```

```
odict_keys(['conv1.weight', 'conv1.bias', 'conv2.weight', 'conv2.bias', 'conv3.weight',  
'conv3.bias', 'fc1.weight', 'fc1.bias', 'fc2.weight', 'fc2.bias'])
```

Serialization

```
model = SimpleCNN()
torch.save(model.state_dict(), 'my_trained_model.pt')

print(model.state_dict().keys())

model.load_state_dict(torch.load('my_trained_model.pt'))
```

```
odict_keys(['conv1.weight', 'conv1.bias', 'conv2.weight', 'conv2.bias', 'conv3.weight',
'conv3.bias', 'fc1.weight', 'fc1.bias', 'fc2.weight', 'fc2.bias'])
```

Serialization

```
model.load_state_dict(torch.load('downloaded_model.pt'))
```

Serialization

```
model.load_state_dict(torch.load('downloaded_model.pt'))
```

Serialization

```
model.load_state_dict(torch.load('downloaded_model.pt'))
```

Can execute arbitrary code!

Serialization

```
model.load_state_dict(torch.load('downloaded_model.pt', weights_only=True))
```


Serialization

```
model.load_state_dict(torch.load('downloaded_model.pt', weights_only=True))
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
        f'checkpoint_epoch_{epoch}.pt')
```

Training...




Epoch 1
...
Epoch 5
Save!

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')
```

Training...



Epoch 1
...
Epoch 5
Save!

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')  
  
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
}  
}
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')  
  
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
}
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')  
  
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
}  
}
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')  
  
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
}  
}
```

Checkpoints

```
if epoch % 5 == 0:  
    torch.save(checkpoint,  
               f'checkpoint_epoch_{epoch}.pt')  
  
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
}  
}
```

Checkpoints

`torch.save`



```
checkpoint_epoch_5.pt  
checkpoint_epoch_10.pt  
checkpoint_epoch_15.pt
```

Checkpoints

`torch.save`



```
Model_final_final.pt  
model_final_ACTUALLY_FINAL.pt  
Model_final_use_this_one.pt
```

Metadata

```
torch.save(model.state_dict(), f'resnet18_acc{accuracy:.1f}_epoch{epoch}.pt')
```

Metadata

```
torch.save(model.state_dict(), f'resnet18_{acc{accuracy:.1f}}_epoch{epoch}.pt')
```

Metadata

```
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
    # Add any other training state you need  
    'model_name': 'resnet18',  
    'dataset_version': 'v2.1',  
    'learning_rate': 0.001,  
    'accuracy': accuracy,  
    'date': datetime.now().isoformat()  
    'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'  
}
```

Metadata

```
checkpoint = {
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'loss': loss,
    # Add any other training state you need
    'model name': 'resnet18',
    ''dataset_version': 'v2.1',
    'learning_rate': 0.001,
    'accuracy': accuracy,
    'date': datetime.now().isoformat()
    'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'
}
```

Metadata

```
checkpoint = {
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'loss': loss,
    # Add any other training state you need
    'model_name': 'resnet18',
    'dataset_version': 'v2.1',
    'learning_rate': 0.001,
    'accuracy': accuracy,
    'date': datetime.now().isoformat()
    'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'
}
```

Metadata

```
checkpoint = {
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'loss': loss,
    # Add any other training state you need
    'model_name': 'resnet18',
    'dataset_version': 'v2.1',
    'learning_rate': 0.001,
    'accuracy': accuracy,
    'date': datetime.now().isoformat(),
    'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'
}
```

Metadata

```
checkpoint = {
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'loss': loss,
    # Add any other training state you need
    'model_name': 'resnet18',
    'dataset_version': 'v2.1',
    'learning_rate': 0.001,
    'accuracy': accuracy,
    'date': datetime.now().isoformat()
}
'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'
```

Metadata

```
checkpoint = {  
    'epoch': epoch,  
    'model_state_dict': model.state_dict(),  
    'optimizer_state_dict': optimizer.state_dict(),  
    'loss': loss,  
    # Add any other training state you need  
    'model_name': 'resnet18',  
    'dataset_version': 'v2.1',  
    'learning_rate': 0.001,  
    'accuracy': accuracy,  
    'date': datetime.now().isoformat()  
    'notes': 'Training paused - accuracy still ~64%. Next: try lower LR or add dropout'  
}
```

mlflow

mlflow

mlflow

torch.save



```
Model_final_final.pt  
model_final_ACTUALLY_FINAL.pt  
Model_final_use_this_one.pt  
chpt_v2b_final_maybe.pt  
run17_overfit_szn.pt  
tuesday_night_fix.pt  
best_so_far_i_think.pt  
dont_delete_this_one.pt  
...
```



mlflow

Time created State: Active Datasets Sort: Run Name

Columns Group by

metrics.rmse < 1 and params.model = "tree"

New run



Time created State: Active Datasets Sort: Run Name

Time created State: Active Datasets Sort: Run Name

Columns Group by

	Metrics				Parameters			
	Run Name	Created	Duration	test_accuracy	train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr...	✓ 1	3.3s	0.826666666...	0.85546875	256	12	0.0001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.1s	0.938666666...	0.9296875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.94	0.96875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.954666666...	0.96875	64	8	0.0005
<input type="checkbox"/>	mlp_lr...	✓ 1	5.3s	0.966666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.2s	0.965333333...	1	64	8	0.001



Time created State: Active Datasets Sort: test_accuracy

Metrics.rmse < 1 and params.model = "tree"

New run

Time created

State: Active

Datasets

Sort: test_accuracy

Columns

Group by

	Metrics					Parameters		
	Run Name	Created	Duration	test_accuracy	train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr...	✓ 1	5.3s	0.966666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.2s	0.965333333...	1	64	8	0.001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.954666666...	0.96875	64	8	0.0005
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.94	0.96875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.1s	0.938666666...	0.9296875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.3s	0.826666666...	0.85546875	256	12	0.0001



Q params.lr = "0.001"

Time created ▾ State: Active ▾ Datasets ▾ Sort: test_accuracy ▾

Columns ▾ Group by ▾

		Metrics				Parameters		
<input type="checkbox"/>	Run Name	Created	Duration	test_accuracy	train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr... ●	✓ 1	5.3s	0.9666666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.2s	0.9653333333...	1	64	8	0.001
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.0s	0.9546666666...	0.96875	64	8	0.0005
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.0s	0.94	0.96875	128	10	0.0003
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.1s	0.9386666666...	0.9296875	128	10	0.0003
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.3s	0.8266666666...	0.85546875	256	12	0.0001



Q params.lr = "0.001"

Time created ▾ State: Active ▾ Datasets ▾ Sort: test_accuracy ▾

Columns ▾ Group by ▾

		Metrics				Parameters		
<input type="checkbox"/>	Run Name	Created	Duration	test_accuracy	<input type="checkbox"/> train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr... ●	✓ 1	5.3s	0.9666666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr... ●	✓ 1	3.2s	0.9653333333...	1	64	8	0.001

Comparing 5 runs

▼ Parameters

Show diff only

batch_size	64	256	128	128	64
epochs	8	12	10	10	8
lr	0.0005	0.0001	0.0003	0.0003	0.001
seed	5	4	3	2	1
weight_decay	0.0005	0.0001	0.0	0.0001	0.001

▼ Metrics

Show diff only

test_accuracy	0.955	0.827	0.94	0.939	0.965
train_accuracy	0.969	0.855	0.969	0.93	1
val_accuracy	0.953	0.832	0.923	0.932	0.961



image-clf-demo > Runs >

mlp_lr=0.0005_wd=0.0005_bs=64_seed=5

⋮

Overview

Model metrics

System metrics

Traces

Artifacts



Q Search metric charts

✓ Auto-refresh



▼ Model metrics (3)

+ Add chart





≡ mlflow

+ New

Experiments

Models

Prompts

Experiments

Create Compare Delete

Filter experiments by name

<input type="checkbox"/> Name	Tags
<input type="checkbox"/> text-classification-bert-baseline	framework: pytorch task:
<input type="checkbox"/> language-model-finetuning	owner: Laurence task: lan
<input type="checkbox"/> sentiment-analysis-transformers	
<input type="checkbox"/> cnn-architecture-comparison	framework: pytorch



mlflow

Time created State: Active Datasets Sort: Run Name

Columns Group by

	Metrics				Parameters			
	Run Name	Created	Duration	test_accuracy	train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr...	1	3.3s	0.826666666...	0.85546875	256	12	0.0001
<input type="checkbox"/>	mlp_lr...	1	3.1s	0.938666666...	0.9296875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	1	3.0s	0.94	0.96875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	1	3.0s	0.954666666...	0.96875	64	8	0.0005
<input type="checkbox"/>	mlp_lr...	1	5.3s	0.966666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr...	1	3.2s	0.965333333...	1	64	8	0.001



mlflow

Time created State: Active Datasets Sort: Run Name

Columns Group by

	Run Name	Created	Duration	Metrics	Parameters			
				test_accuracy	train_accuracy	batch_size	epochs	lr
<input type="checkbox"/>	mlp_lr...	✓ 1	3.3s	0.8266666666...	0.85546875	256	12	0.0001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.1s	0.9386666666...	0.9296875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.94	0.96875	128	10	0.0003
<input type="checkbox"/>	mlp_lr...	✓ 1	3.0s	0.9546666666...	0.96875	64	8	0.0005
<input type="checkbox"/>	mlp_lr...	✓ 1	5.3s	0.9666666666...	0.984375	64	8	0.001
<input type="checkbox"/>	mlp_lr...	✓ 1	3.2s	0.9653333333...	1	64	8	0.001



DeepLearning.AI

Exporting Models with ONNX

Preparing Models for Deployment in PyTorch

Deployment

- **Save and Track**
- Export
- Optimize


Serialization and Version Control



Deployment


- Save and Track
- **Export**
- Optimize

Export



PyTorch-only


Export



servers



phones



embedded devices




the cloud

Export




export



servers



phones




embedded devices



the cloud


ONNX



servers



phones




embedded devices




the cloud

ONNX



ONNX



Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()  
  
# Example tensor with expected input shape  
dummy_input = torch.randn(1, 3, 224, 224)
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()  
  
# Example tensor with expected input shape  
dummy_input = torch.randn(1, 3, 224, 224)
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()  
  
# Example tensor with expected input shape  
dummy_input = torch.randn(1, 3, 224, 224)  
              |  
             batch size
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()  
  
# Example tensor with expected input shape  
dummy_input = torch.randn(1, 3, 224, 224)  
              |  
              RGB channels
```

Load a Model

```
model = MyModel()  
model.load_state_dict(torch.load('model_weights.pt'))  
model.eval()  
  
# Example tensor with expected input shape  
dummy_input = torch.randn(1, 3, 224, 224)  
              | . |  
              height & width
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,           # PyTorch model
    dummy_input,     # Input tensor
    "model.onnx",    # Output file name
    export_params=True, # Store weights in model file
    opset_version=11, # ONNX version
    do_constant_folding=True, # Optimize constant folding
    input_names=['input'],    # Name inputs
    output_names=['output'],   # Name outputs
    dynamic_axes={          # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```


Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Operator	Since Version
AveragePool	22,19,11,10..
Flatten	24,23,21,13..
MatMul	13,9,1
MaxPool	22,12,11,10..
Reshape	24,23,21,19..
Squeeze	24,23,21,13..

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```



Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

4

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['input'],       # Name inputs
    output_names=['output'],     # Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['images'],      # Name inputs
    output_names=['predictions'],# Name outputs
    dynamic_axes={              # Support variable batch size
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['images'],      # Name inputs
    output_names=['predictions'],# Name outputs
    dynamic_axes={
        'input': {0: 'batch_size'},
        'output': {0: 'batch_size'}
    }
)
```

```
torch.randn(1, 3, 224, 224)
```

Export to ONNX

```
# Export to ONNX format
torch.onnx.export(
    model,                      # PyTorch model
    dummy_input,                # Input tensor
    "model.onnx",               # Output file name
    export_params=True,          # Store weights in model file
    opset_version=11,            # ONNX version
    do_constant_folding=True,   # Optimize constant folding
    input_names=['images'],      # Name inputs
    output_names=['predictions'],# Name outputs
    dynamic_axes={
        'input': [0: 'batch_size'],
        'output': {0: 'batch_size'}
    }
)
```

```
torch.randn(1, 3, 224, 224)
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference

```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX Inference


```
import onnxruntime as ort
import numpy as np

# Load ONNX model
session = ort.InferenceSession("model.onnx")

# Prepare input
input_name = session.get_inputs()[0].name
input_data = np.random.randn(1, 3, 224, 224).astype(np.float32)

# Run inference
outputs = session.run(None, {input_name: input_data})
prediction = outputs[0]
```

ONNX






DeepLearning.AI

Pruning

Preparing Models for Deployment in PyTorch

Real-World Models



92% Accuracy



Fast




Lightweight



Efficient

ONNX



Model Compression




Fast



Lightweight



Efficient



Compression Techniques

- **Pruning**
- **Quantization**


Compression Techniques

- **Pruning**
- **Quantization**

Compression Techniques


- **Pruning:** Not every part of a model is important
- **Quantization**

Compression Techniques




Millions of Parameters

Compression Techniques




Millions of Parameters

Unstructured Pruning




More Parameters than Needed

Unstructured Pruning




More Parameters than Needed

Unstructured Pruning



More Parameters than Needed

Unstructured Pruning



More Parameters than Needed

Unstructured Pruning

```
model.parameters()
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)  
  
prune.random_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune

prune.l1_unstructured(model.fc1, name="weight", amount=0.3)

prune.random_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune

prune.l1_unstructured(model.fc1, name="bias", amount=0.3)

prune.random_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)  
  
prune.random_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune

prune.l1_unstructured(model.fc1, name="weight", amount=100)

prune.random_unstructured(model.fc1, name="weight", amount=0.3)
```

Unstructured Pruning

```
print("fc1 weights:\n", model.fc1.weight)

tensor([[ -0.2658,   0.0693,   0.0689,   0.3122],
       [  0.4445,   0.1917,   0.1849,   0.1577],
       [  0.0758,   0.3863,   0.1648,   0.3379],
       . . .]
```

Unstructured Pruning

```
print("fc1 weights:\n", model.fc1.weight)

tensor([[ -0.2658,   0.0,     0.0,     0.3122], same size
       [ 0.4445,   0.1917,   0.1849,   0.0      ],
       [ 0.0,     0.3863,   0.0,     0.3379], same speed
       . . .]
```

Unstructured Pruning

```
import torch.nn.utils.prune as prune

prune.l1_unstructured(model.fc1, name="weight", amount=100)


prune.random_unstructured(model.fc1, name="weight", amount=0.3)

model.fc1.weight = model.fc1.weight.to_sparse()
```


Store non-zero only

Smaller but not optimized


Structured Pruning



Structured Pruning



Structured Pruning




Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```

Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```


nn.Linear(2, 4)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```


nn.Linear(2, 4)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```


nn.Linear(2, 4)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```


nn.Linear(2, 4)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)
```


nn.Linear(2, 2)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=1)
```


nn.Linear(2, 2)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=1)
```


nn.Linear(2, 2)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=1)
```


nn.Linear(1, 2)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=1)  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```


```
nn.Linear(1, 2)
```



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=1)  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```


```
nn.Linear(1, 2)
```



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```


nn.Linear(1, 2)



Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.fc1, name="weight", amount=0.5, n=2, dim=0)  
  
prune.l1_unstructured(model.fc1, name="weight", amount=0.3)
```

```
nn.Linear(1, 2)
```



L1 vs. L2 Structured Pruning

weights:

Neuron A

0.82
-0.15
0.08
-0.04

Neuron B

0.45
0.38
-0.32
0.26

L1 Structured Pruning

`torch.sum(torch.abs(weights))`

A: 1.09 < B: 1.41

L2 Structured Pruning

`torch.sqrt(torch.sum(weights ** 2))`

A: 0.84 > B: 0.72

Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.conv1, name="weight", amount=0.5, n=2, dim=0)
```

Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.conv1, name="weight", amount=0.5, n=2, dim=0)
```

Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.conv1, name="weight", amount=0.5, n=2, dim=0)
```


Cut filters

Structured Pruning

```
import torch.nn.utils.prune as prune  
  
prune.ln_structured(model.conv1, name="weight", amount=0.5, n=2, dim=1)
```


Cut inputs to filters

Structured Pruning




Structured Pruning


Fewer parameters
Fewer feature maps



Larger impact on
accuracy




Pruning




Metric	Before pruning	After pruning
Accuracy	95 %	93 %
Inference time	100 ms	70 ms
File size	50 MB	25 MB


Pruning



Pruning



Pruning






DeepLearning.AI


Static and Dynamic Quantization

Preparing Models for Deployment in PyTorch

Pruning




Pruning



32-bit Weights

Model Weight

0.0283746




32-bit floating point value

0 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 1 | 1 | 1 | 0

During Training

Model Weight


0.0283746



During Training

Model Weight


0.0283744



During Training

Model Weight


0.0283732



During Training

Model Weight

0.0283668



During Training

Model Weight

0.0283668

32-bit floating point value

0	0	1	1	1	1	0	0	1	1	1	0	1	0	0	0	0	1	1	1	0	0	0	0	1	0	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Training Complete

Model Weight

0.0283668



32-bit floating point value

0 | 0 | 1 | 1 | 1 | 1 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 1 | 1 | 1 | 1 | 0

Still need this much precision?

Training Complete

Model Weight

0.0283668

32-bit floating point value


0	0	1	1	1	1	0	0	1	1	1	0	1	0	0	0	0	1	1	1	0	0	0	0	1	0	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Training Complete

Model Weight

26

32-bit floating point value



Noise Tolerance

Model Weight

0.0283668

Class Predictions

Cat: 0.82

Dog: 0.17

Bird: 0.01

Noise Tolerance

Model Weight

0.0280000

Class Predictions

Cat: 0.82

Dog: 0.17

Bird: 0.01

Noise Tolerance

Model Weight

0.0280000

Class Predictions

Cat: 0.81
Dog: 0.18
Bird: 0.01

Quantization

Model Weight

0.0283668

32-bit float

0	0	1	1	1	1	0	0	1	1	1	0	1	0	0	0	0	1	1	1	0	0	0	0	1	0	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Quantization

Model Weight

26

32-bit float

0 0 1 1 1 1 0 0 1 1 1 0 1 0 0 0 0 1 1 1 0 0 0 0 1 0 1 1 1 1 1 0

8-bit int
↓

0 0 0 1 1 0 1 0

Quantization

```
class SimpleModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Linear(4, 1)

    def forward(self, x):
        x = self.layer1(x)
        return x
```

Quantization

```
class SimpleModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Linear(4, 1)

    def forward(self, x):
        x = self.layer1(x)
        return x
```

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])
```

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])
```

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])
```

```
bias = model.layer1.bias.data
```

bias not quantized

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])
```

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])  
  
scale = weights.abs().max()/127 # abs(-1.0412)/127
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])  
  
scale = weights.abs().max()/127 # abs(-1.0412)/127  
print(scale)  
> tensor(0.0082)
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data  
print(weights)  
> tensor([0.12537, -0.3749, 0.7526, -1.0412])  
  
scale = weights.abs().max()/127 # abs(-1.0412)/127  
print(scale)  
> tensor(0.0082)  
  
int8_weights = torch.round(weights/scale).to(torch.int8)
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data
print(weights)
> tensor([0.12537, -0.3749, 0.7526, -1.0412])

scale = weights.abs().max()/127 # abs(-1.0412)/127
print(scale)
> tensor(0.0082)

int8_weights = torch.round(weights/scale).to(torch.int8)
print(int8_weights)
> tensor([15, -46, 92, -127], dtype=torch.int8)
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data
print(weights)
> tensor([0.12537, -0.3749, 0.7526, -1.0412])

scale = weights.abs().max()/127 # abs(-1.0412)/127
print(scale)
> tensor(0.0082)

int8_weights = torch.round(weights/scale).to(torch.int8)
print(int8_weights)
> tensor([15, -46, 92, -127], dtype=torch.int8)
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data
print(weights)
> tensor([0.12537, -0.3749, 0.7526, -1.0412])

scale = weights.abs().max()/127 # abs(-1.0412)/127
print(scale)
> tensor(0.0082)

int8_weights = torch.round(weights/scale).to(torch.int8)
print(int8_weights)
> tensor([15, -46, 92, -127], dtype=torch.int8)
```

int8 range: -128 to 127

Quantizing Weights

```
weights = model.layer1.weight.data
print(weights)
> tensor([0.12537, -0.3749, 0.7526, -1.0412])


scale = weights.abs().max()/127 # abs(-1.0412)/127
print(scale)
> tensor(0.0082)

int8_weights = torch.round(weights/scale).to(torch.int8)
print(int8_weights)
> tensor([15, -46, 92, -127], dtype=torch.int8)


dequantized = int8_weights.float * scale
```

int8 range: -128 to 127

Quantization




Quantization



Known Weights:
Scale values can be precomputed

Quantization




How do you quantize numbers that only appear at runtime?

You can't derive scale values ahead of time

Calibration

```
# Sample data
```

```
([[ 0.25, -0.40,  0.80,  0.10],  
 [-0.10,  0.30, -0.50,  0.60],  
 [ 0.75,  0.20, -0.10, -0.30],  
 [-0.45, -0.25,  0.50,  0.40],  
 [ 0.10,  0.90,  0.20, -0.60],  
 [-0.20,  0.15, -0.70,  0.50],  
 . . .
```




Static Quantization

```
class QuantizedCNN(nn.Module):
    def __init__(self):
        super(QuantizedCNN, self).__init__()

    def forward(self, x):

        return x
```




Static Quantization

```
class QuantizedCNN(nn.Module):
    def __init__(self):
        super(QuantizedCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()

    def forward(self, x):

        return x
```



Static Quantization

```
class QuantizedCNN(nn.Module):
    def __init__(self):
        super(QuantizedCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        ... # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):

        return x
```




Static Quantization

```
class QuantizedCNN(nn.Module):
    def __init__(self):
        super(QuantizedCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        . . . # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):

        return x
```




Static Quantization

```
class QuantizedCNN(nn.Module):
    def __init__(self):
        super(QuantizedCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        . . . # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):
        x = self.quant(x)
        . . . # Model Layers
        x = self.dequant(x)
        return x
```



Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()
```

Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()
```

Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()
```

Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')
```

Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')
```


Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')  
  
torch.quantization.prepare(quantized_static_model, inplace=True)
```

Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')  
  
torch.quantization.prepare(quantized_static_model, inplace=True)  
  
calibrate(quantized_static_model, testloader)
```


Observers



Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')  
  
torch.quantization.prepare(quantized_static_model, inplace=True)  
  
calibrate(quantized_static_model, testloader)  
  
torch.quantization.convert(quantized_static_model, inplace=True)
```

After Convert



Preparing the Model

```
quantized_static_model = QuantizedCNN()  
  
quantized_static_model.load_state_dict(baseline_model.state_dict())  
quantized_static_model.eval()  
  
# Set as 'x86' in the lab - see lab for details  
quantized_static_model.qconfig = torch.quantization.get_default_qconfig('fbgemm')  
  
torch.quantization.prepare(quantized_static_model, inplace=True)  
  
calibrate(quantized_static_model, testloader)  
  
torch.quantization.convert(quantized_static_model, inplace=True)
```

Dynamic Quantization

```
baseline_model.eval()
```

Dynamic Quantization

```
baseline_model.eval()

quantized_model = torch.quantization.quantize_dynamic(
    baseline_model,
    {nn.Linear},
    dtype=torch.qint8
)
```

Dynamic Quantization

```
baseline_model.eval()

quantized_model = torch.quantization.quantize_dynamic(
    baseline_model,
    {nn.Linear},
    dtype=torch.qint8
)
```

Dynamic Quantization

```
baseline_model.eval()  
  
No calibration step  
quantized_model = torch.quantization.quantize_dynamic(  
    baseline_model,  
    {nn.Linear},  
    dtype=torch.qint8  
)
```

Dynamic vs. Static Quantization

- **Dynamic doesn't support all layers** (e.g. conv2d)
- **Static and Dynamic only runs on CPUs**

Dynamic vs. Static Quantization

Dynamic:

- Minimal code
- No calibration
- nn.Linear heavy models

Static:

- More setup
- Bigger speed and memory savings
- Works with convolutions (CNNs)



DeepLearning.AI

Quantization Aware Training

Preparing Models for Deployment in PyTorch

Quantization

Model Weight

0.0283668

32-bit floating point value

0	0	1	1	1	1	0	0	1	1	1	0	1	0	0	0	0	1	1	1	0	0	0	0	1	0	1	1	1	1	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---


Quantization

Model Weight


26

32-bit floating point value


0 0 1 1 1 1 0 0 1 1 1 0 1 0 0 0 0 1 1 1 0 0 0 0 1 0 1 1 1 1 1 0




Quantization




Quantization



Quantization Aware Training (QAT)



Quantization Aware Training (QAT)




QAT

```
class QATCNN(nn.Module):
    def __init__(self):
        super(QATCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        . . . # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):
        x = self.quant(x)
        . . . # Model Layers
        x = self.dequant(x)
        return x
```




QAT

```
class QATCNN(nn.Module):
    def __init__(self):
        super(QATCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        . . . # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):
        x = self.quant(x)
        . . . # Model Layers
        x = self.dequant(x)
        return x
```




QAT

```
class QATCNN(nn.Module):
    def __init__(self):
        super(QATCNN, self).__init__()

        self.quant = torch.quantization.QuantStub()
        ... # Model Layers
        self.dequant = torch.quantization.DeQuantStub()

    def forward(self, x):
        x = self.quant(x)
        ... # Model Layers
        x = self.dequant(x)
        return x
```




QAT

```
# Prepare the model for Quantization-Aware Training.  
qat_model = torch.quantization.prepare_qat(qat_model)
```


QAT

```
# Prepare the model for Quantization-Aware Training.  
qat_model = torch.quantization.prepare_qat(qat_model)
```


QAT




QAT




QAT




QAT




QAT



QAT



QAT



Fake Quantization

Activation: **0.125**

Scale: **0.041**

Fake Quantize { Quantize: **round(0.125 / 0.041)**

Fake Quantization

Activation: **0.125**

Scale: **0.041**

Fake Quantize {
Quantize: **3**
Dequantize: **3*0.041**


Fake Quantization

Activation: **0.125**


Scale: **0.041**

Fake Quantize {
Quantize: **3**
Dequantize: **0.123**


QAT on Layers




QAT on Layers




QAT on Layers



QAT on Layers




QAT on Layers




QAT on Layers

```
class QuantizedCNN(nn.Module):  
    ...  
  
    def fuse_model(self):  
        torch.quantization.fuse_modules(self,  
            ['conv1', 'bn1', 'relu1'], inplace=True)
```




QAT on Layers

```
class QuantizedCNN(nn.Module):  
    ...  
  
    def fuse_model(self):  
        torch.quantization.fuse_modules(self,  
            ['conv1', 'bn1', 'relu1'], inplace=True)
```



QAT on Layers

```
class QuantizedCNN(nn.Module):  
    ...  
  
    def fuse_model(self):  
        torch.quantization.fuse_modules(self,  
            ['conv1', 'bn1', 'relu1'], inplace=True)
```



QAT Setup

- **Include stubs**
- **Fuse layers**
- **Load pretrained weights**
- **Prepare and finetune the model**
- **Convert final model**



DeepLearning.AI

PyTorch for Deep Learning

Conclusion